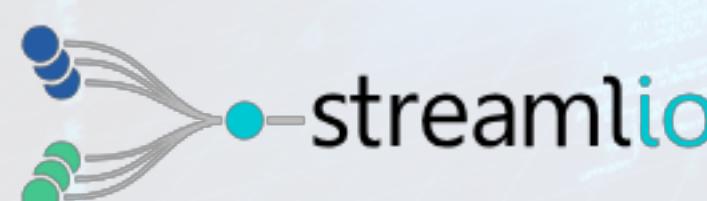


# UNIFYING MESSAGING, QUEUING, STREAMING & COMPUTE WITH APACHE PULSAR

KARTHIK RAMASAMY

CO-FOUNDER AND CEO





# Connected World





# Ubiquity of Real-Time Data Streams & Events

# EVENT/STREAM DATA PROCESSING

- ◆ Events are analyzed and processed as they arrive
- ◆ Decisions are timely, contextual and based on fresh data
- ◆ Decision latency is eliminated
- ◆ Data in motion



# EVENT/STREAM PROCESSING PATTERNS

MONITORING

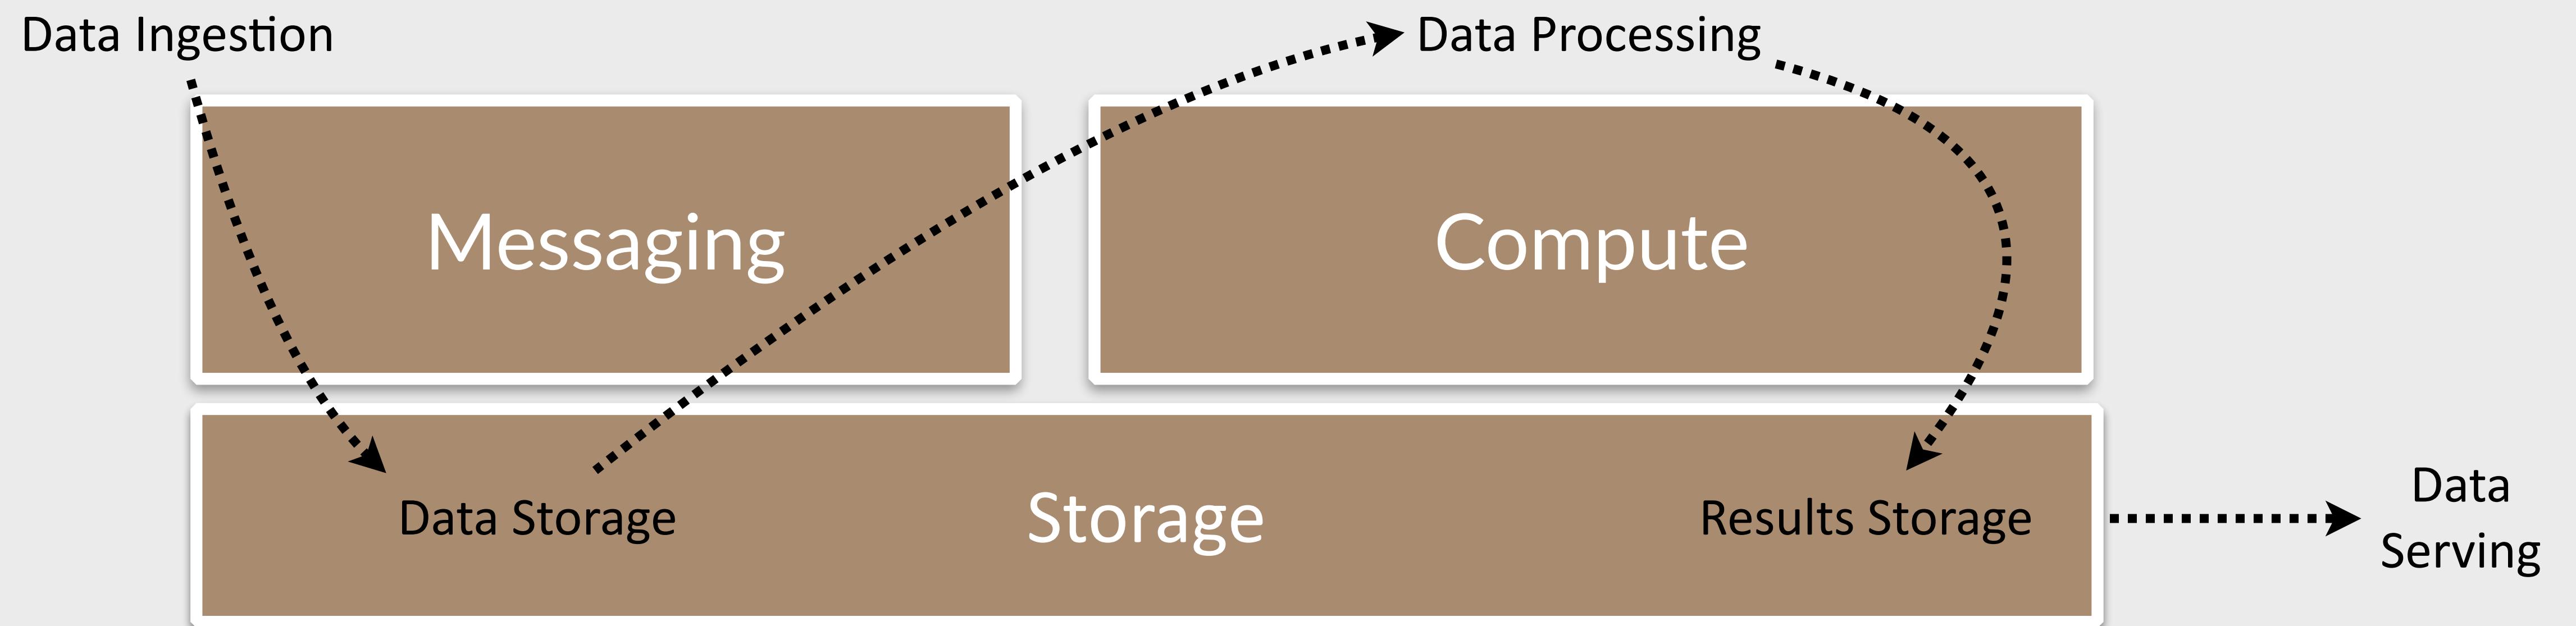
MICROSERVICES

WORKFLOWS

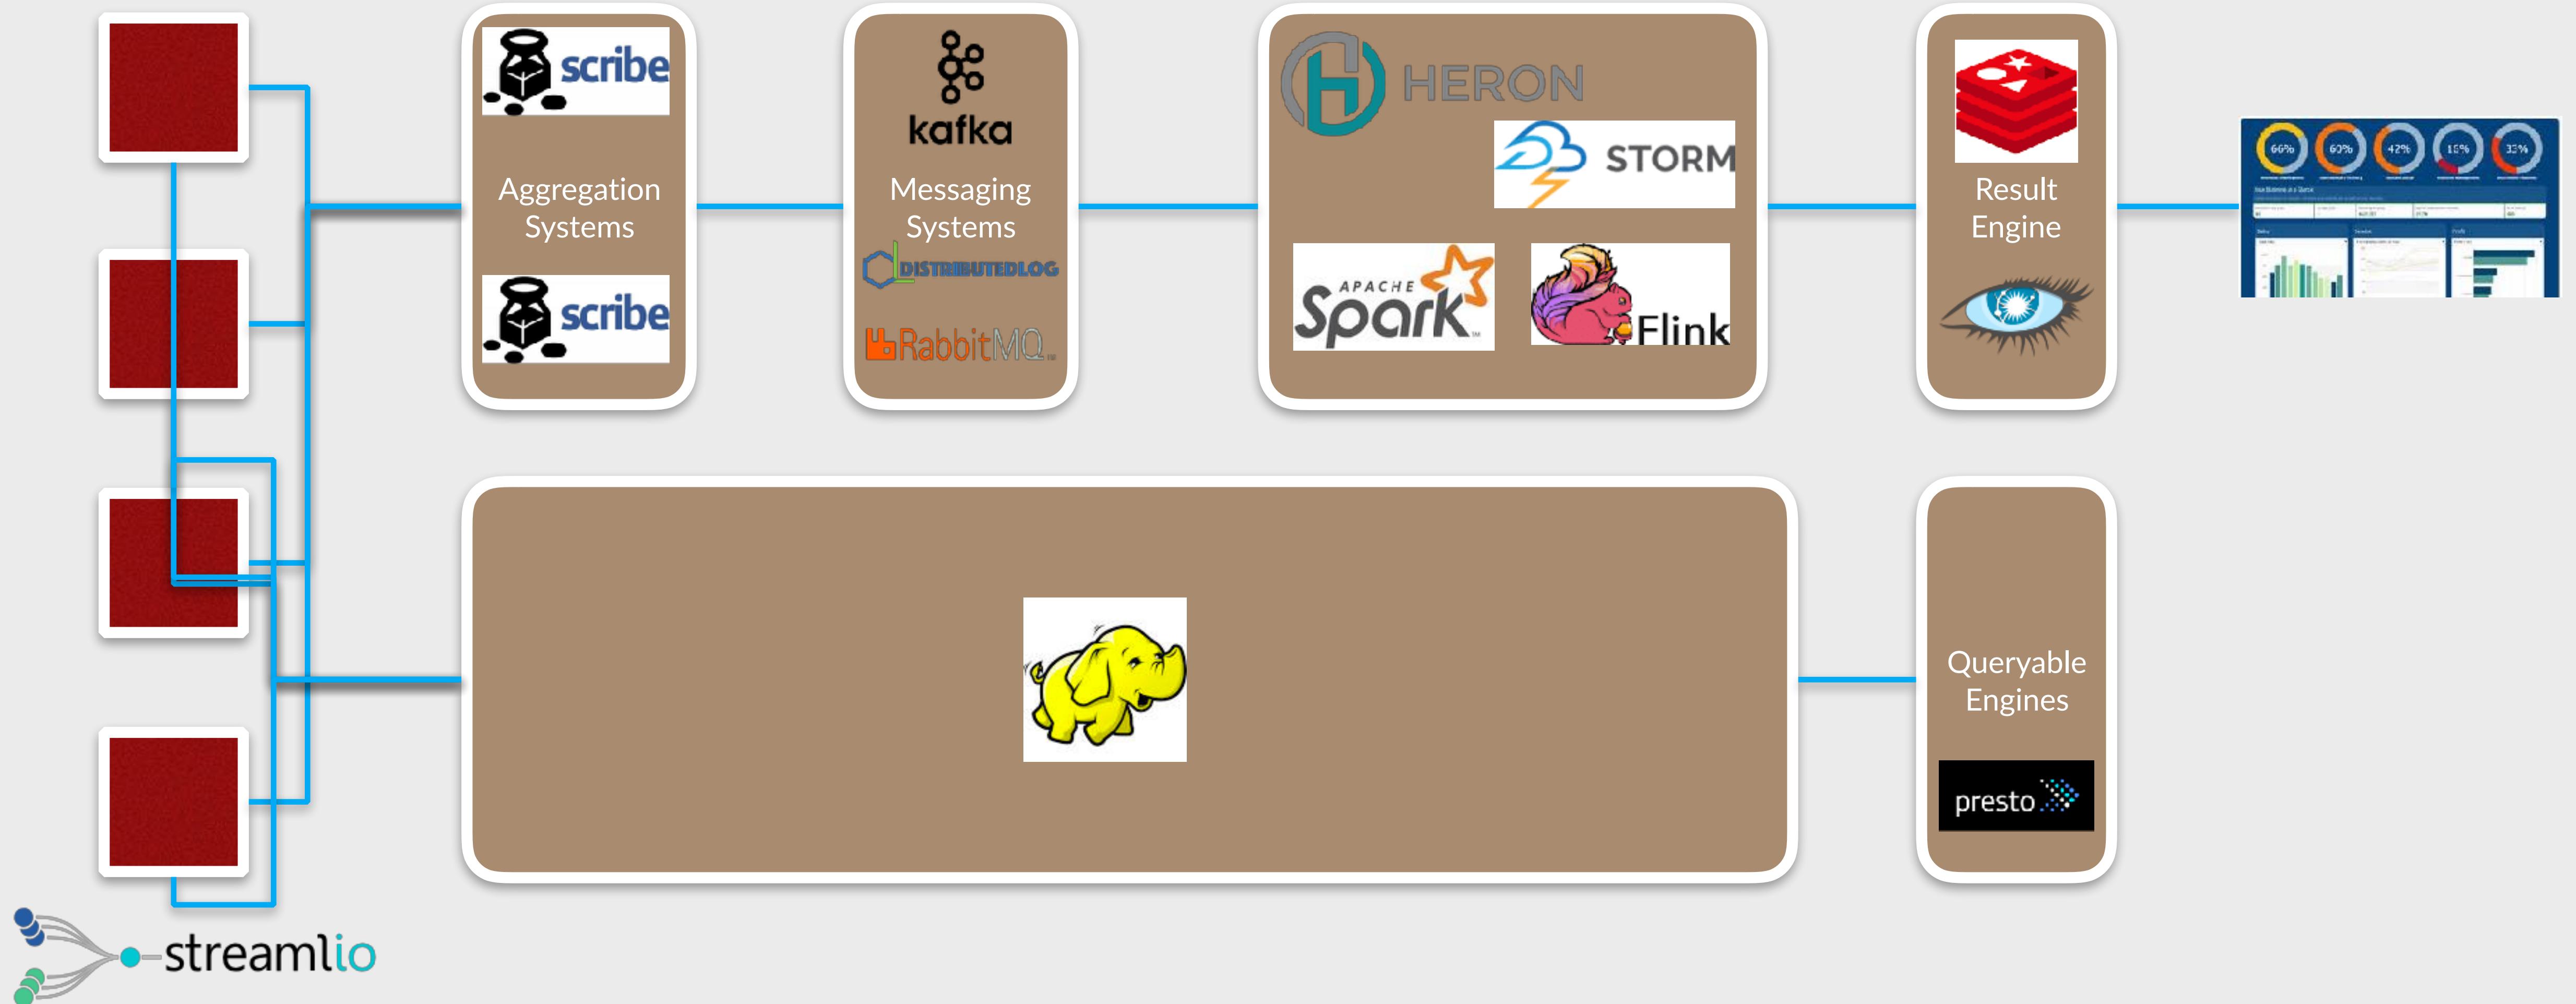
MODEL INFERENCE

ANALYTICS

# STREAM PROCESSING PATTERN



# ELEMENTS OF EVENT/STREAM PROCESSING



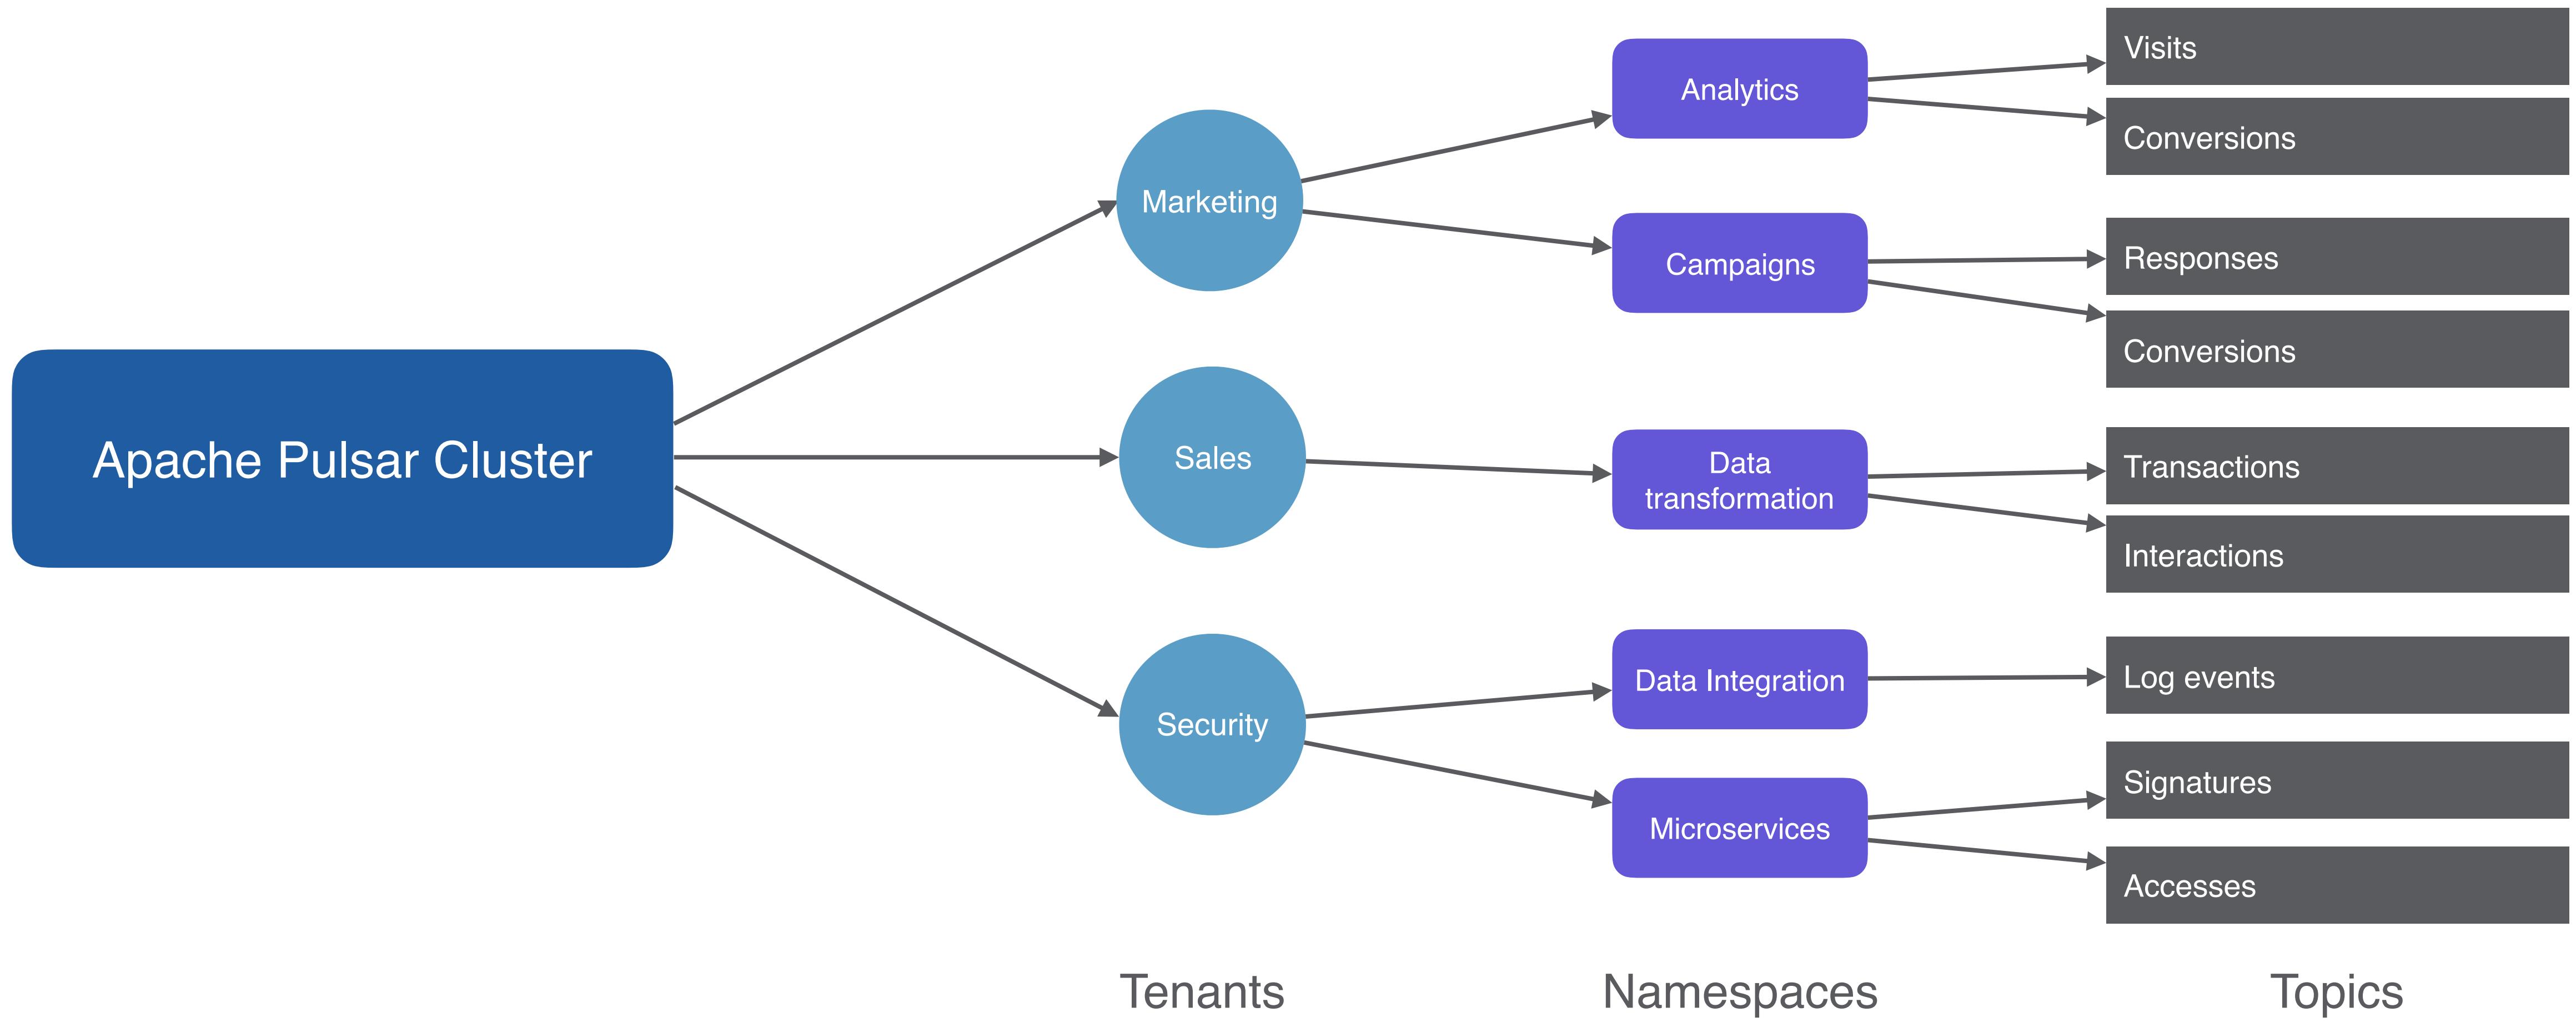
# APACHE PULSAR

**Flexible Messaging + Streaming System**  
backed by a durable log storage

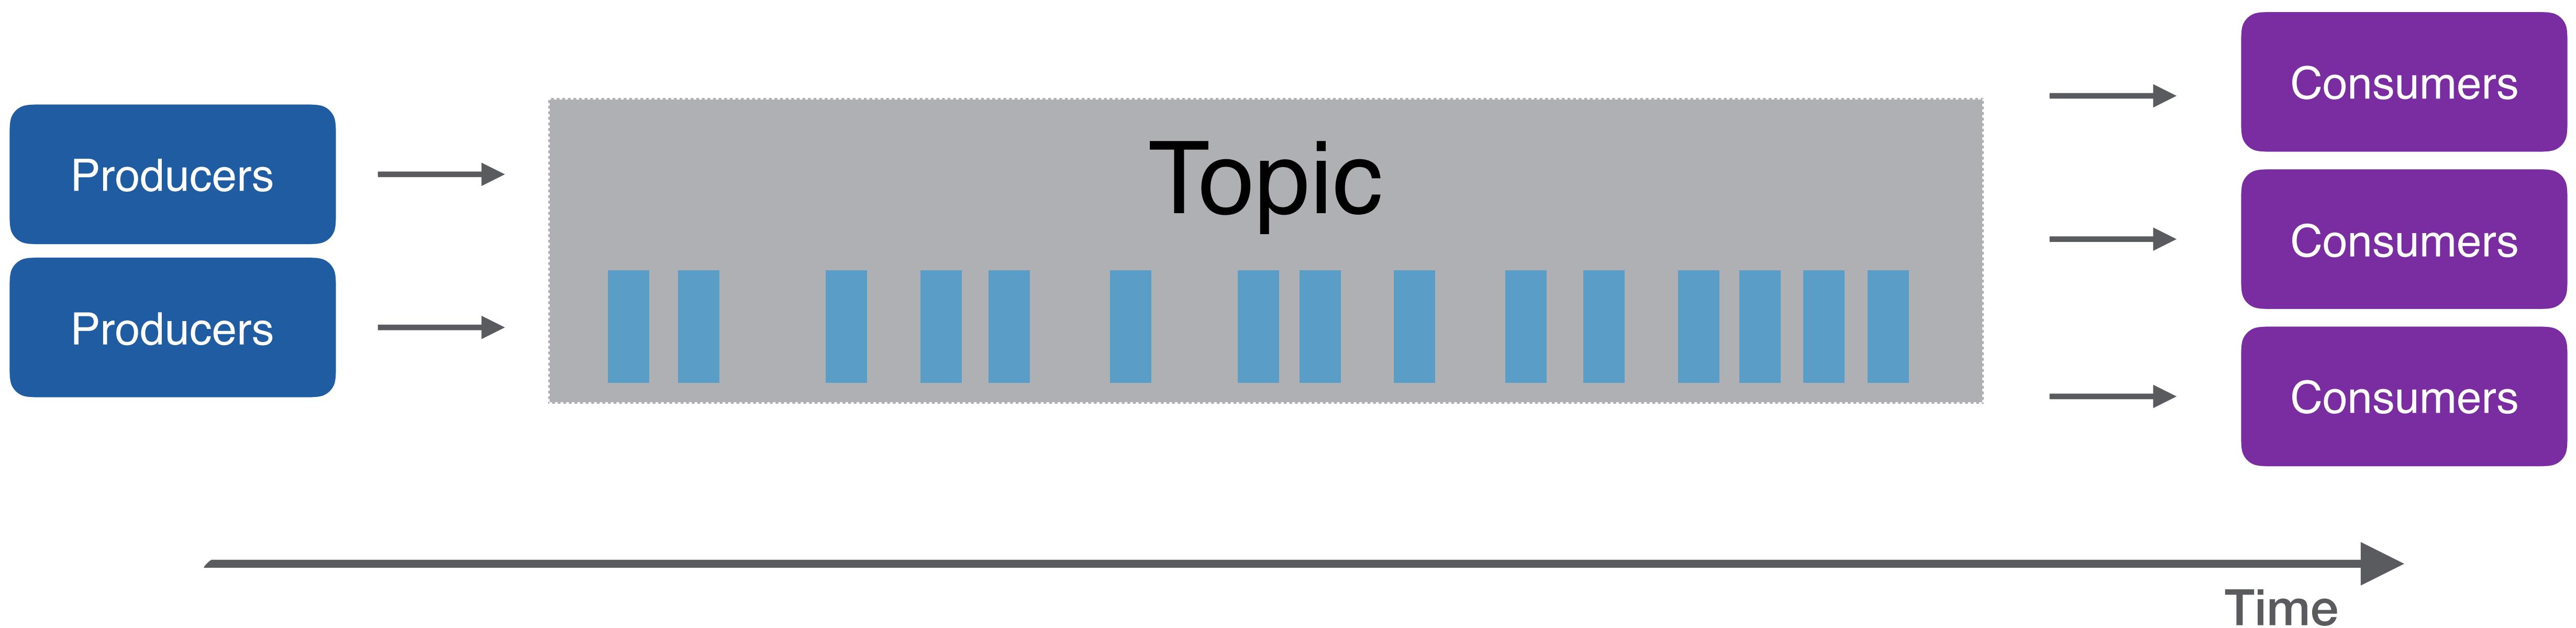


# Key Concepts

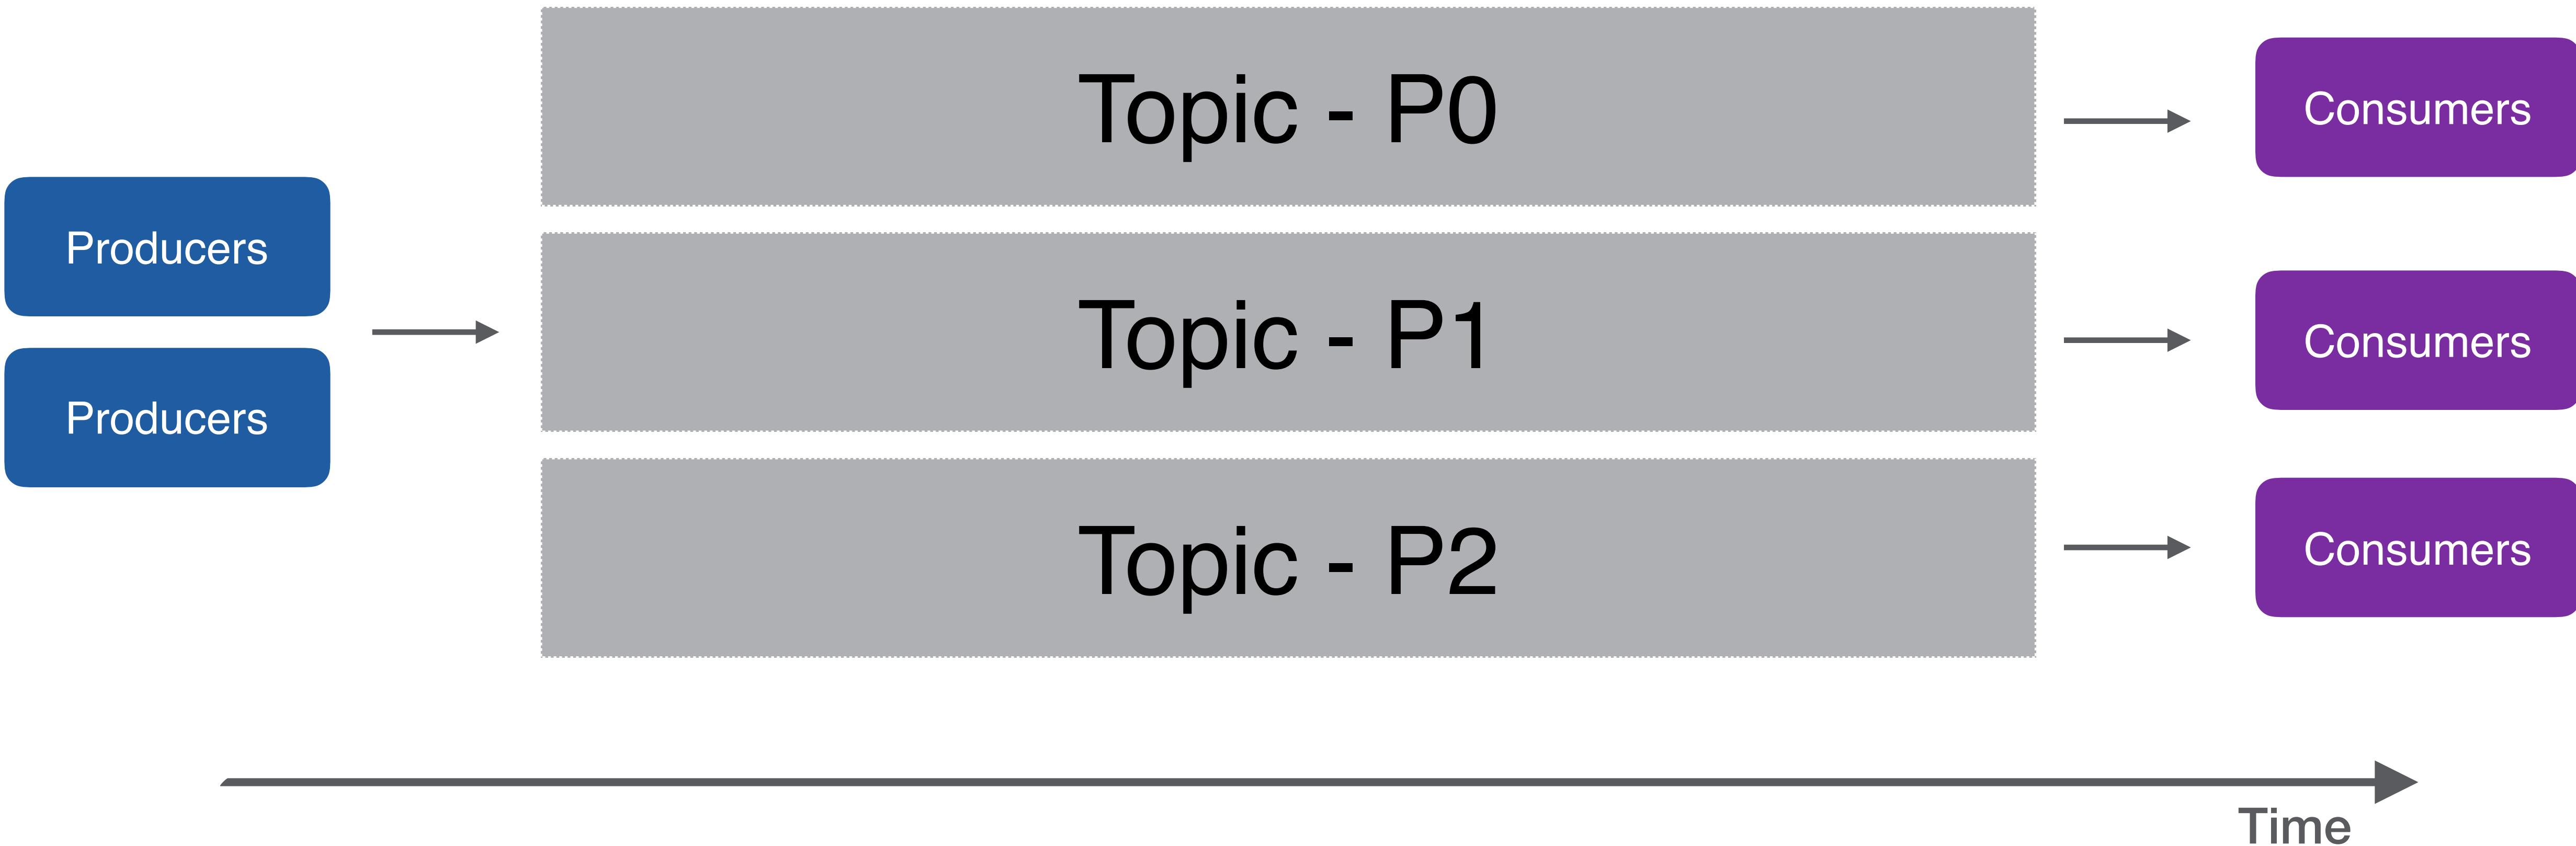
# Core concepts: Tenants, namespaces, topics



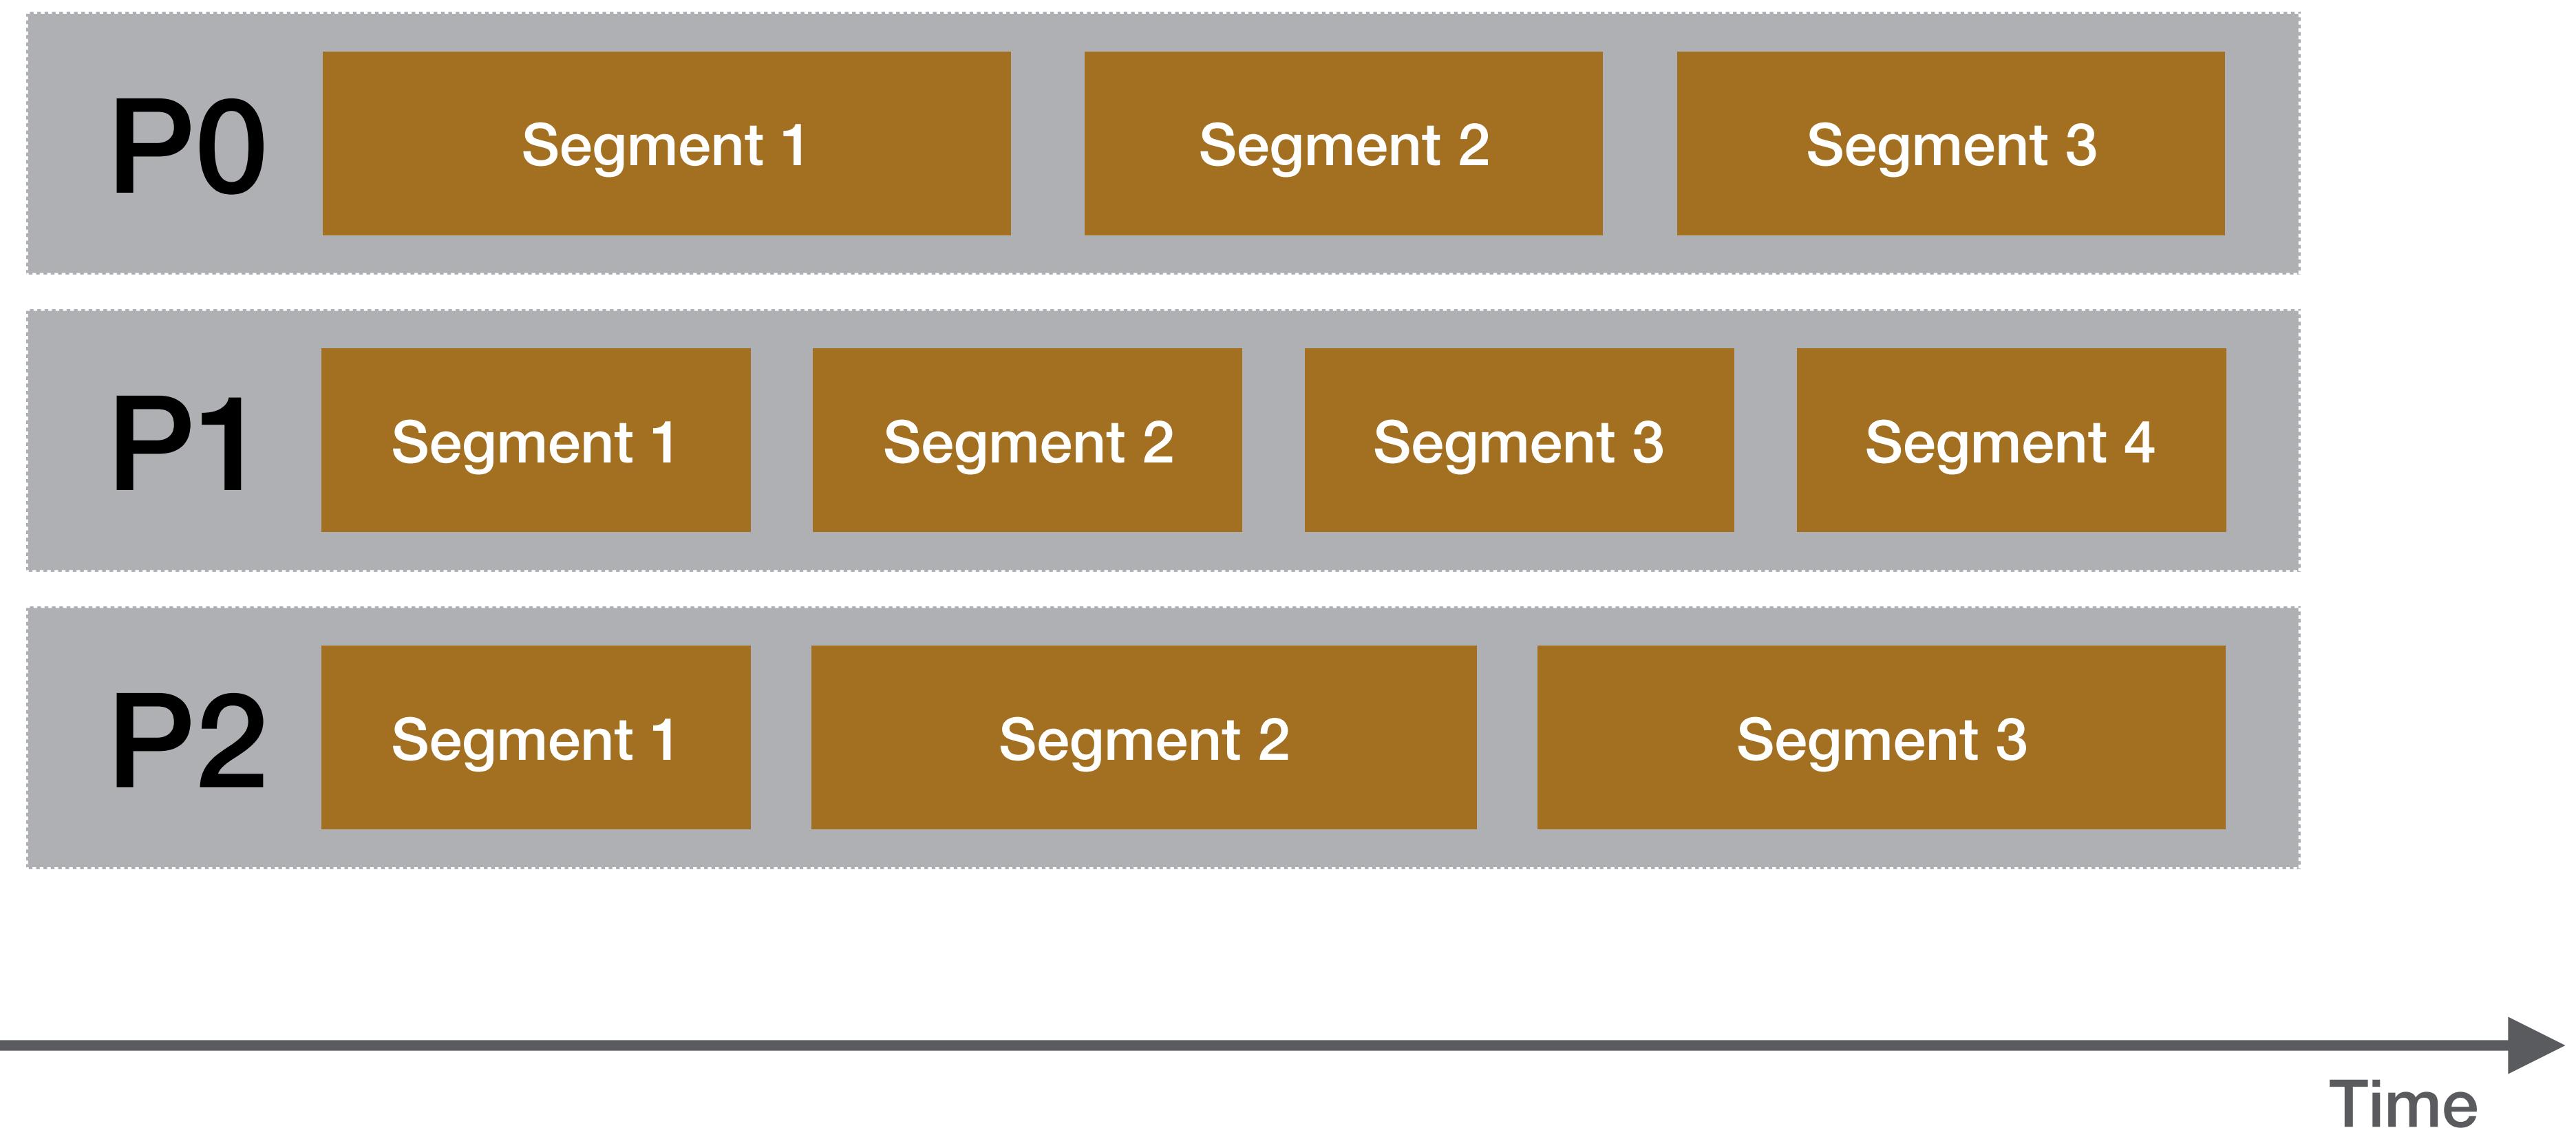
# Topics



# Topic partitions

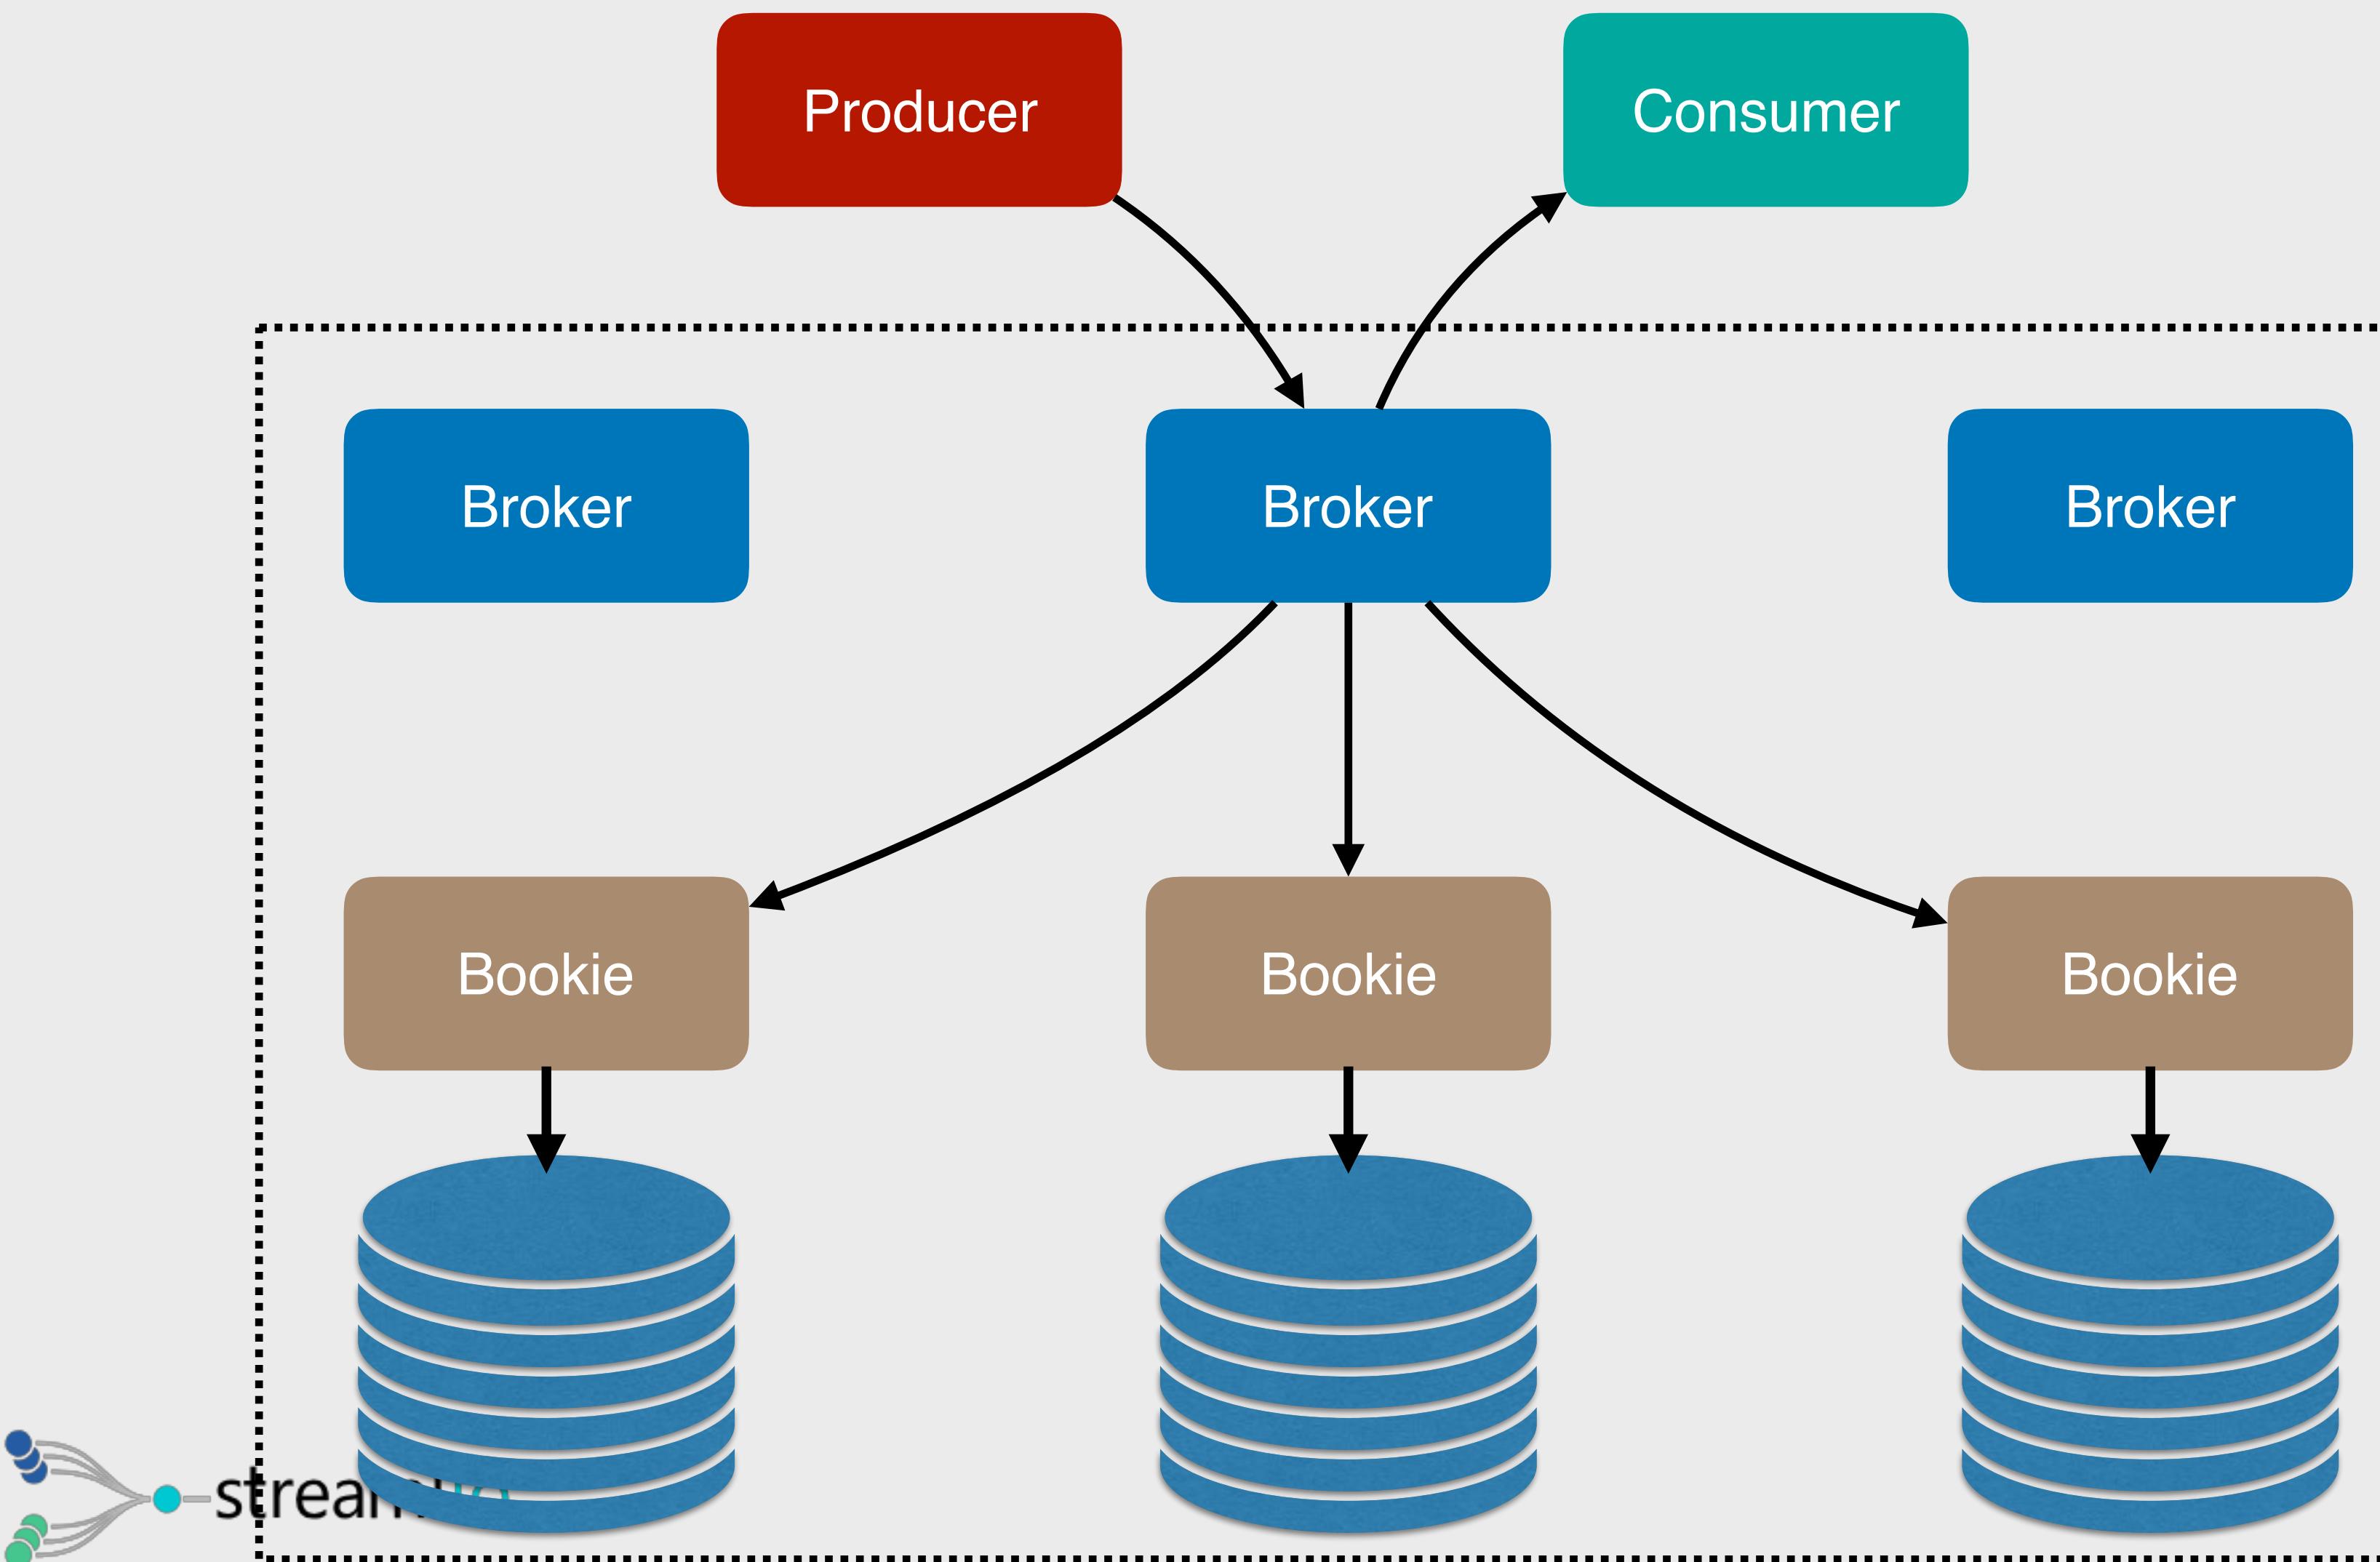


# Segments



# Architecture

# APACHE PULSAR



## SERVING

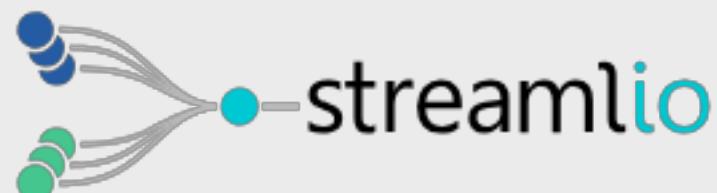
Brokers can be added independently  
Traffic can be shifted quickly across brokers

## STORAGE

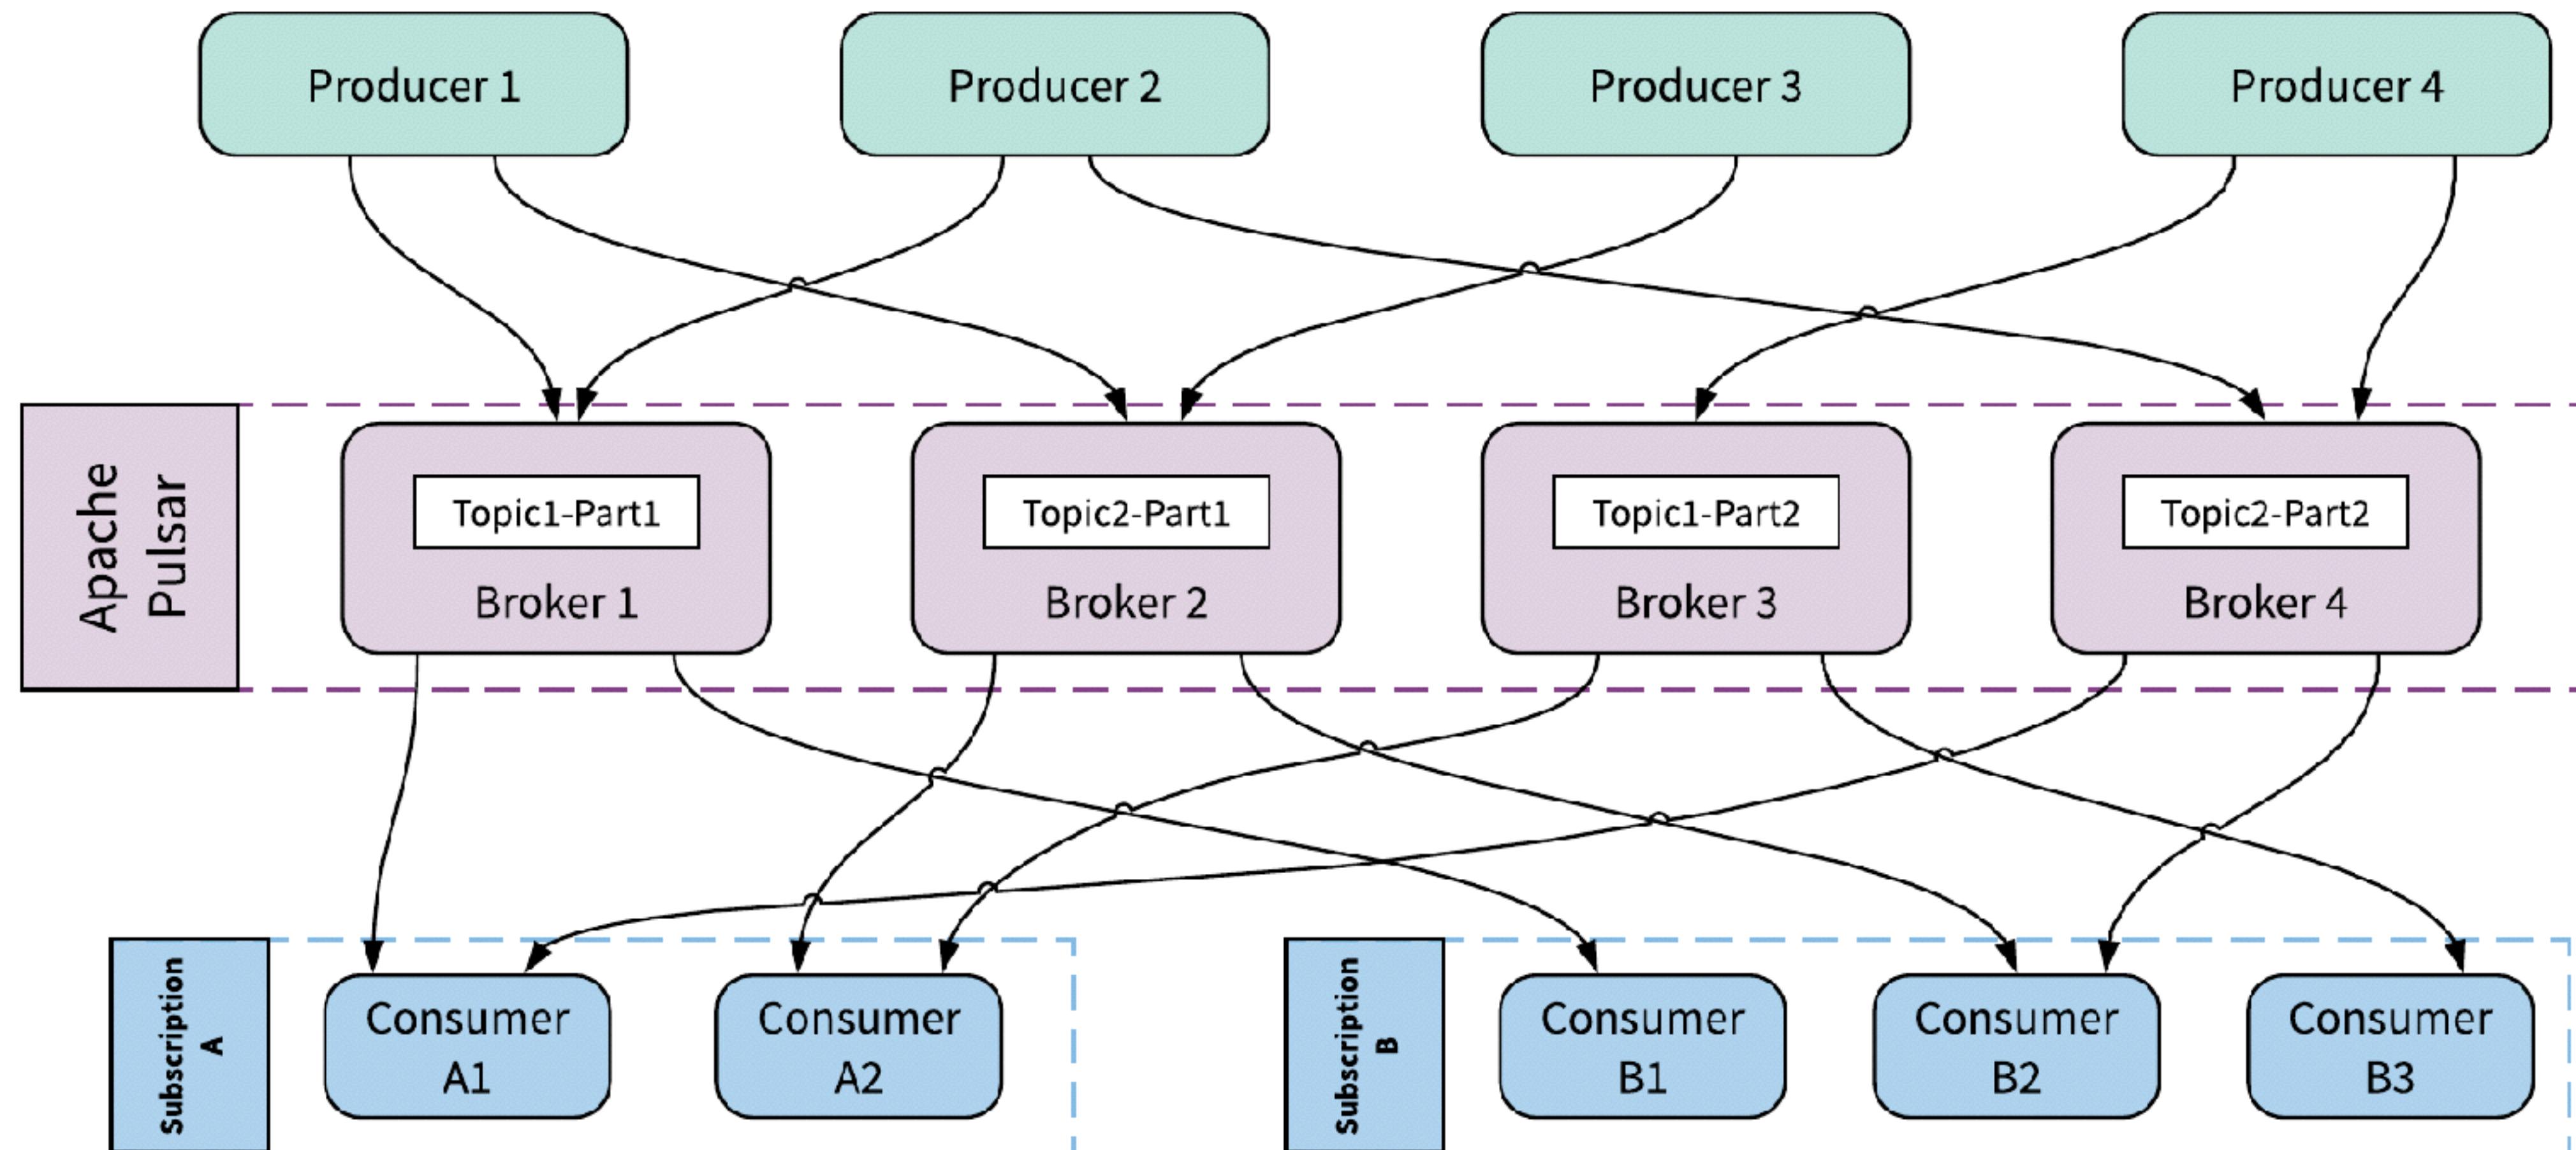
Bookies can be added independently  
New bookies will ramp up traffic quickly

# APACHE PULSAR - BROKER

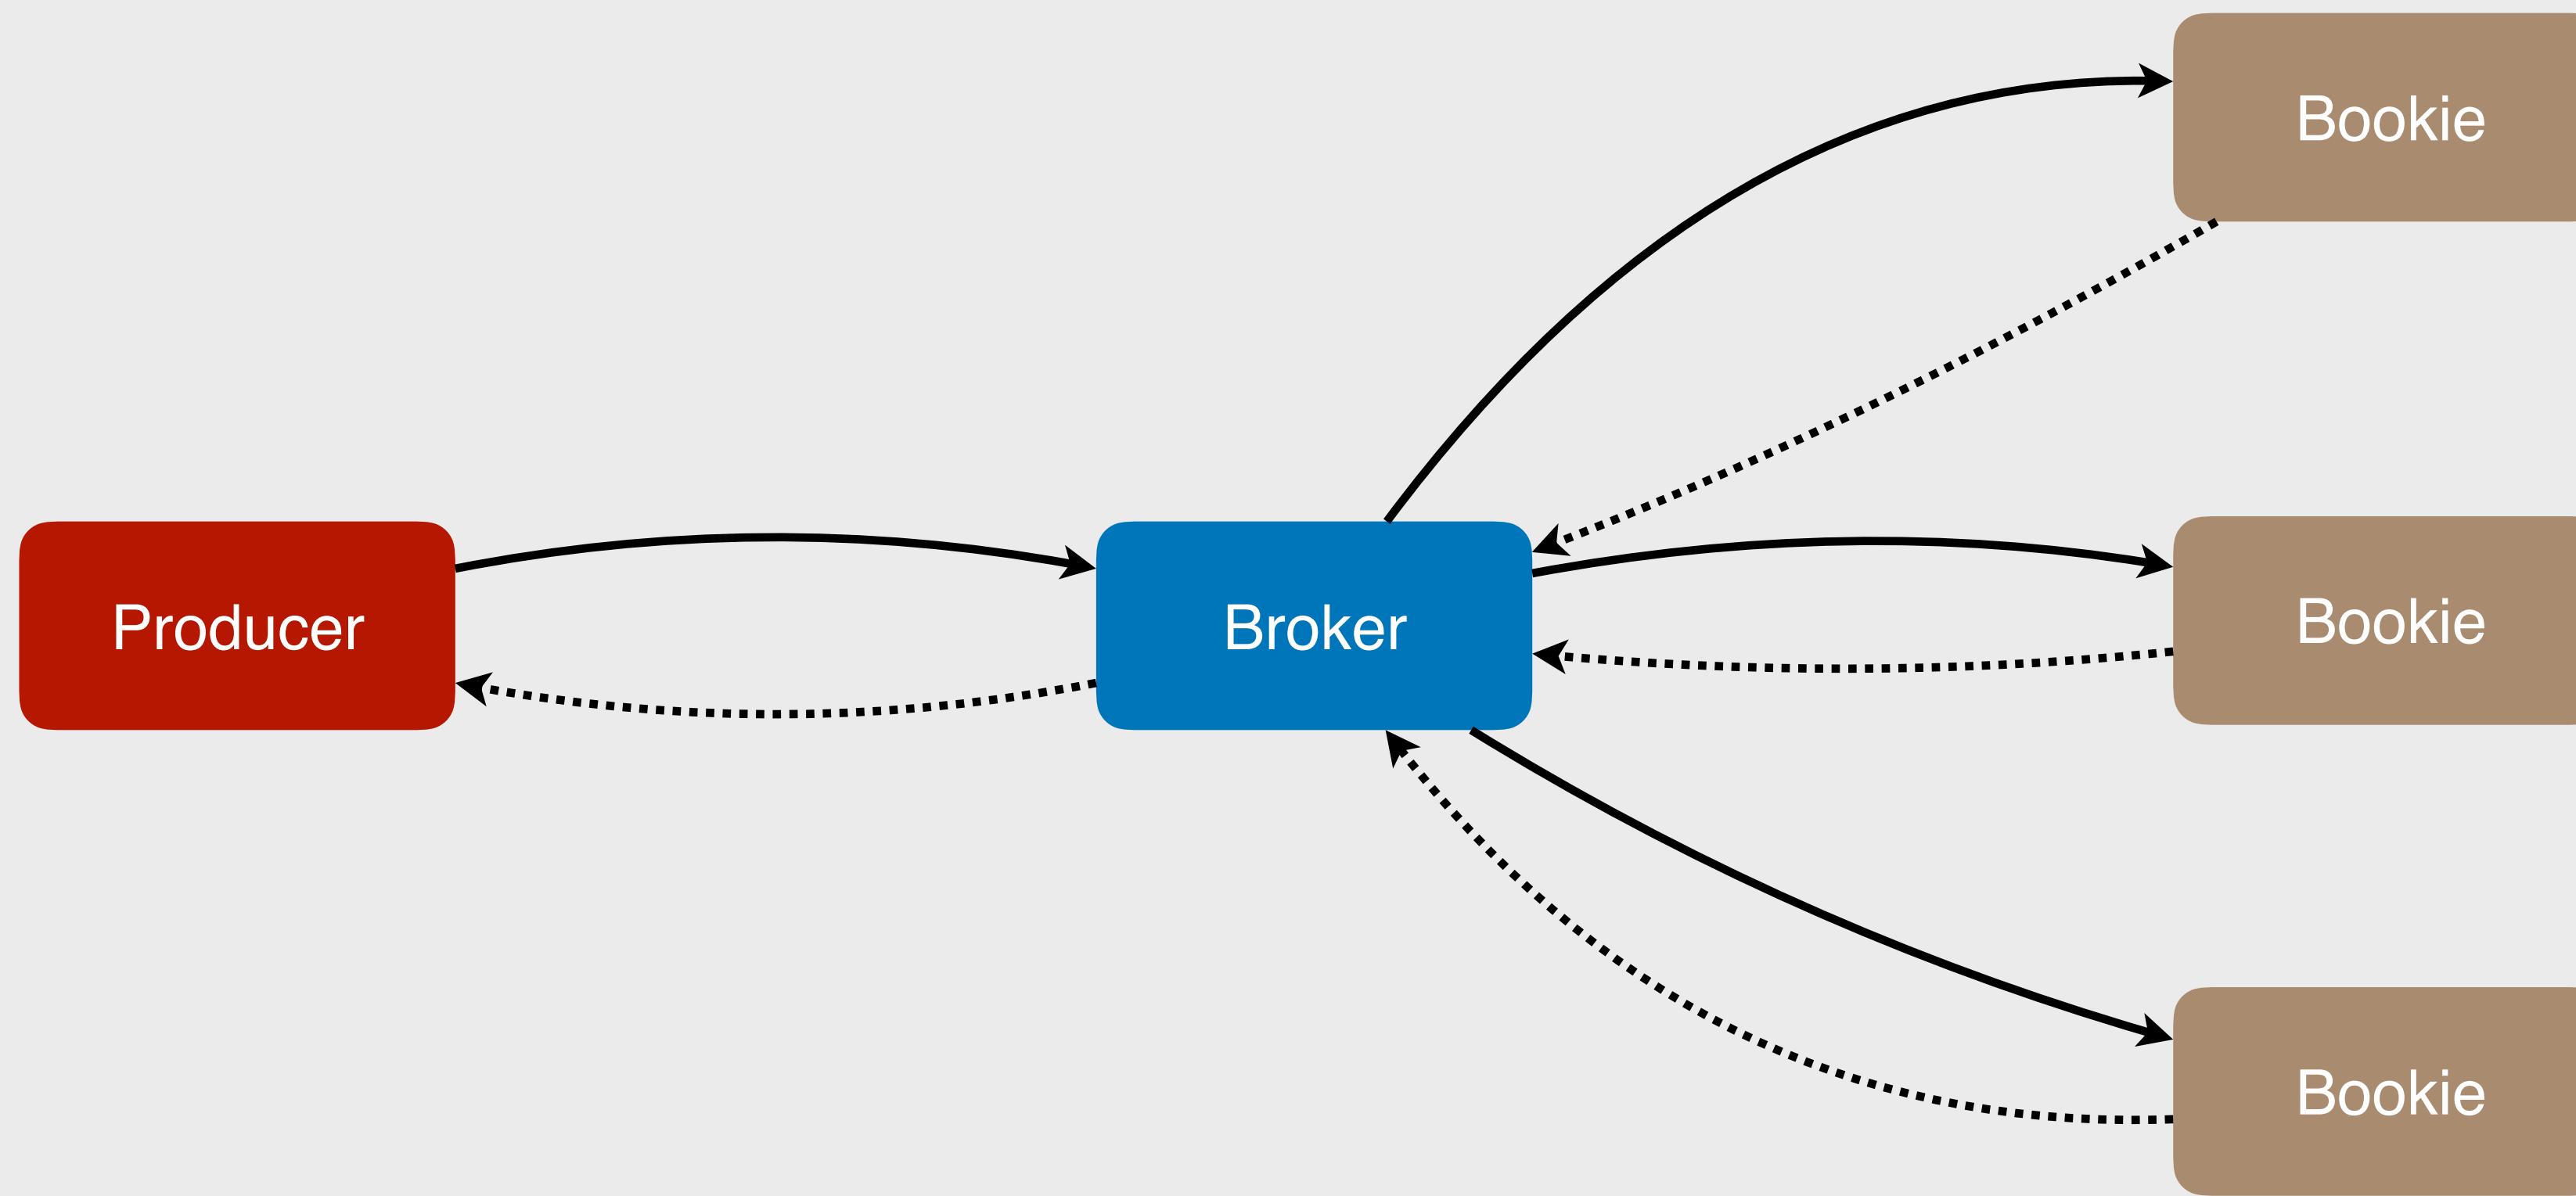
- ◆ Broker is the only point of interaction for clients (producers and consumers)
- ◆ Brokers acquire ownership of group of topics and “serve” them
- ◆ Broker has no durable state
- ◆ Provides service discovery mechanism for client to connect to right broker



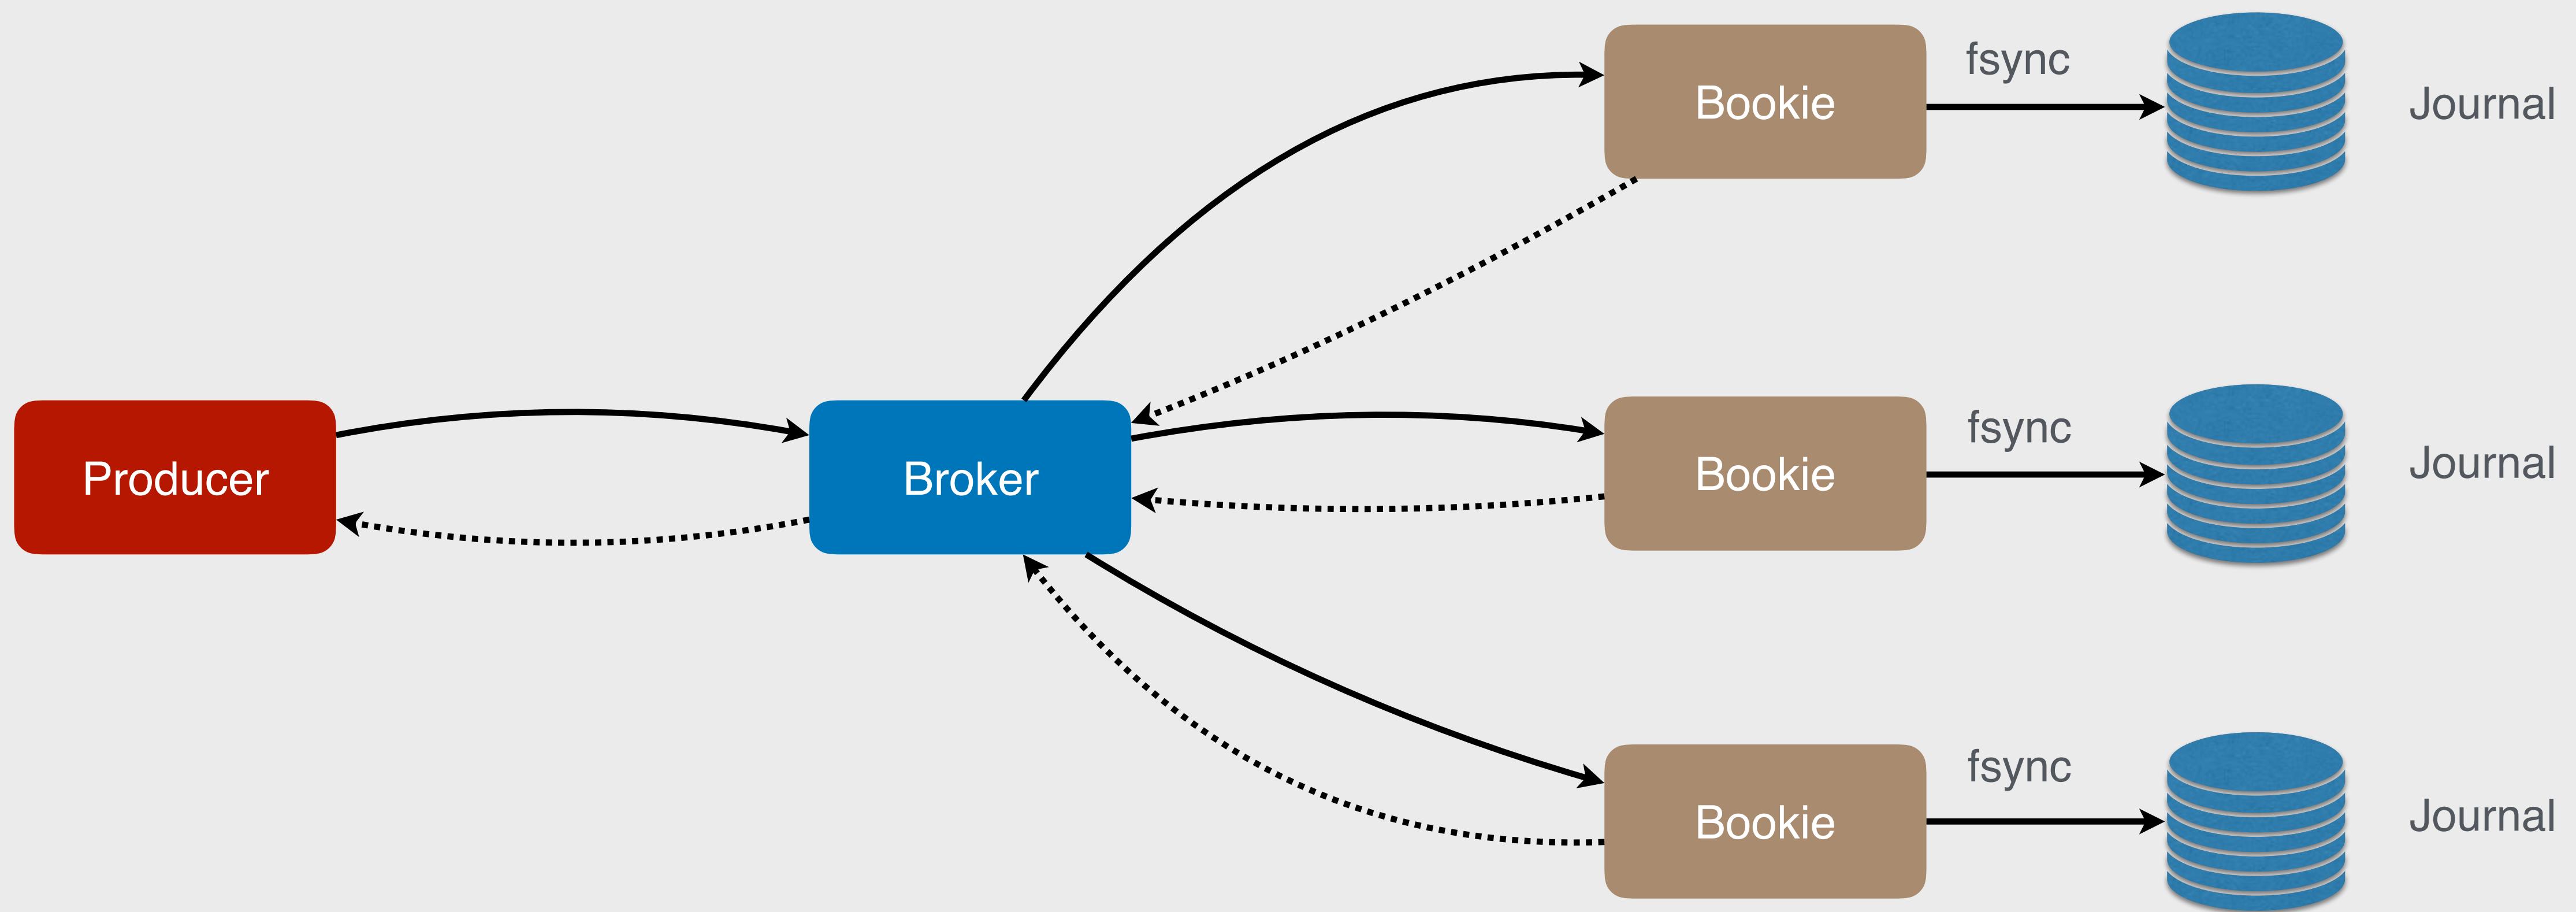
# APACHE PULSAR - BROKER



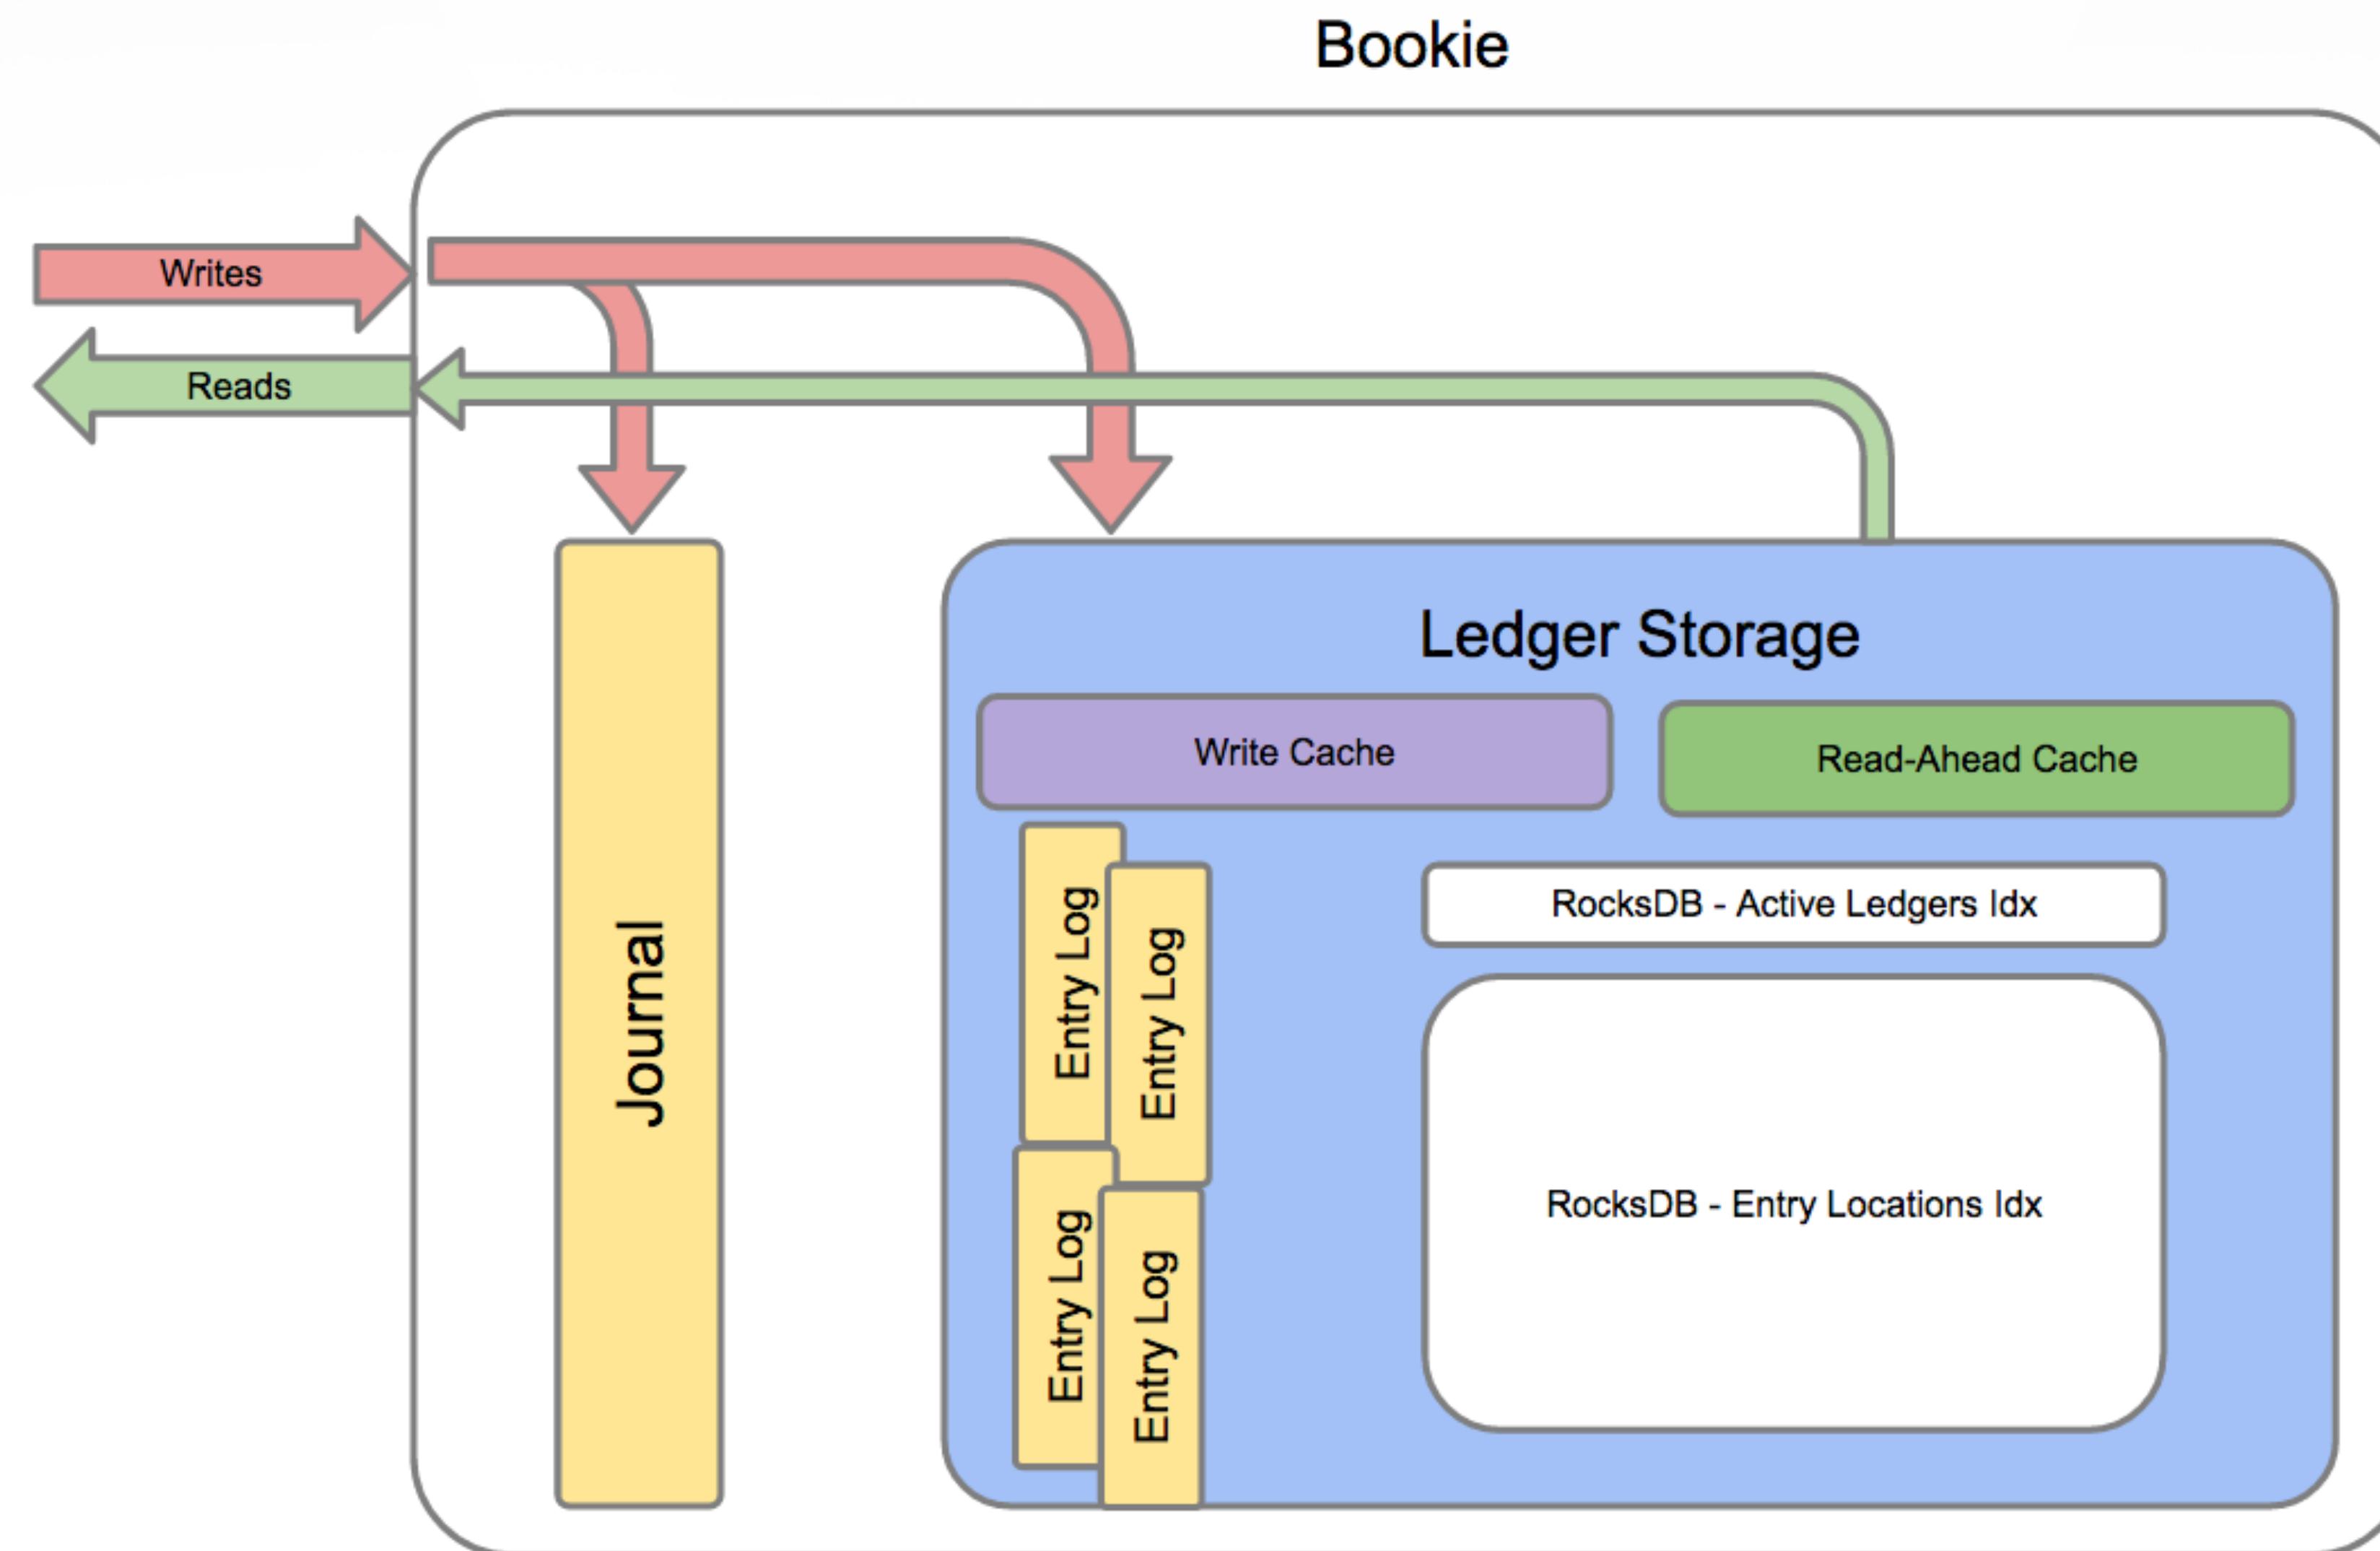
# APACHE PULSAR - CONSISTENCY



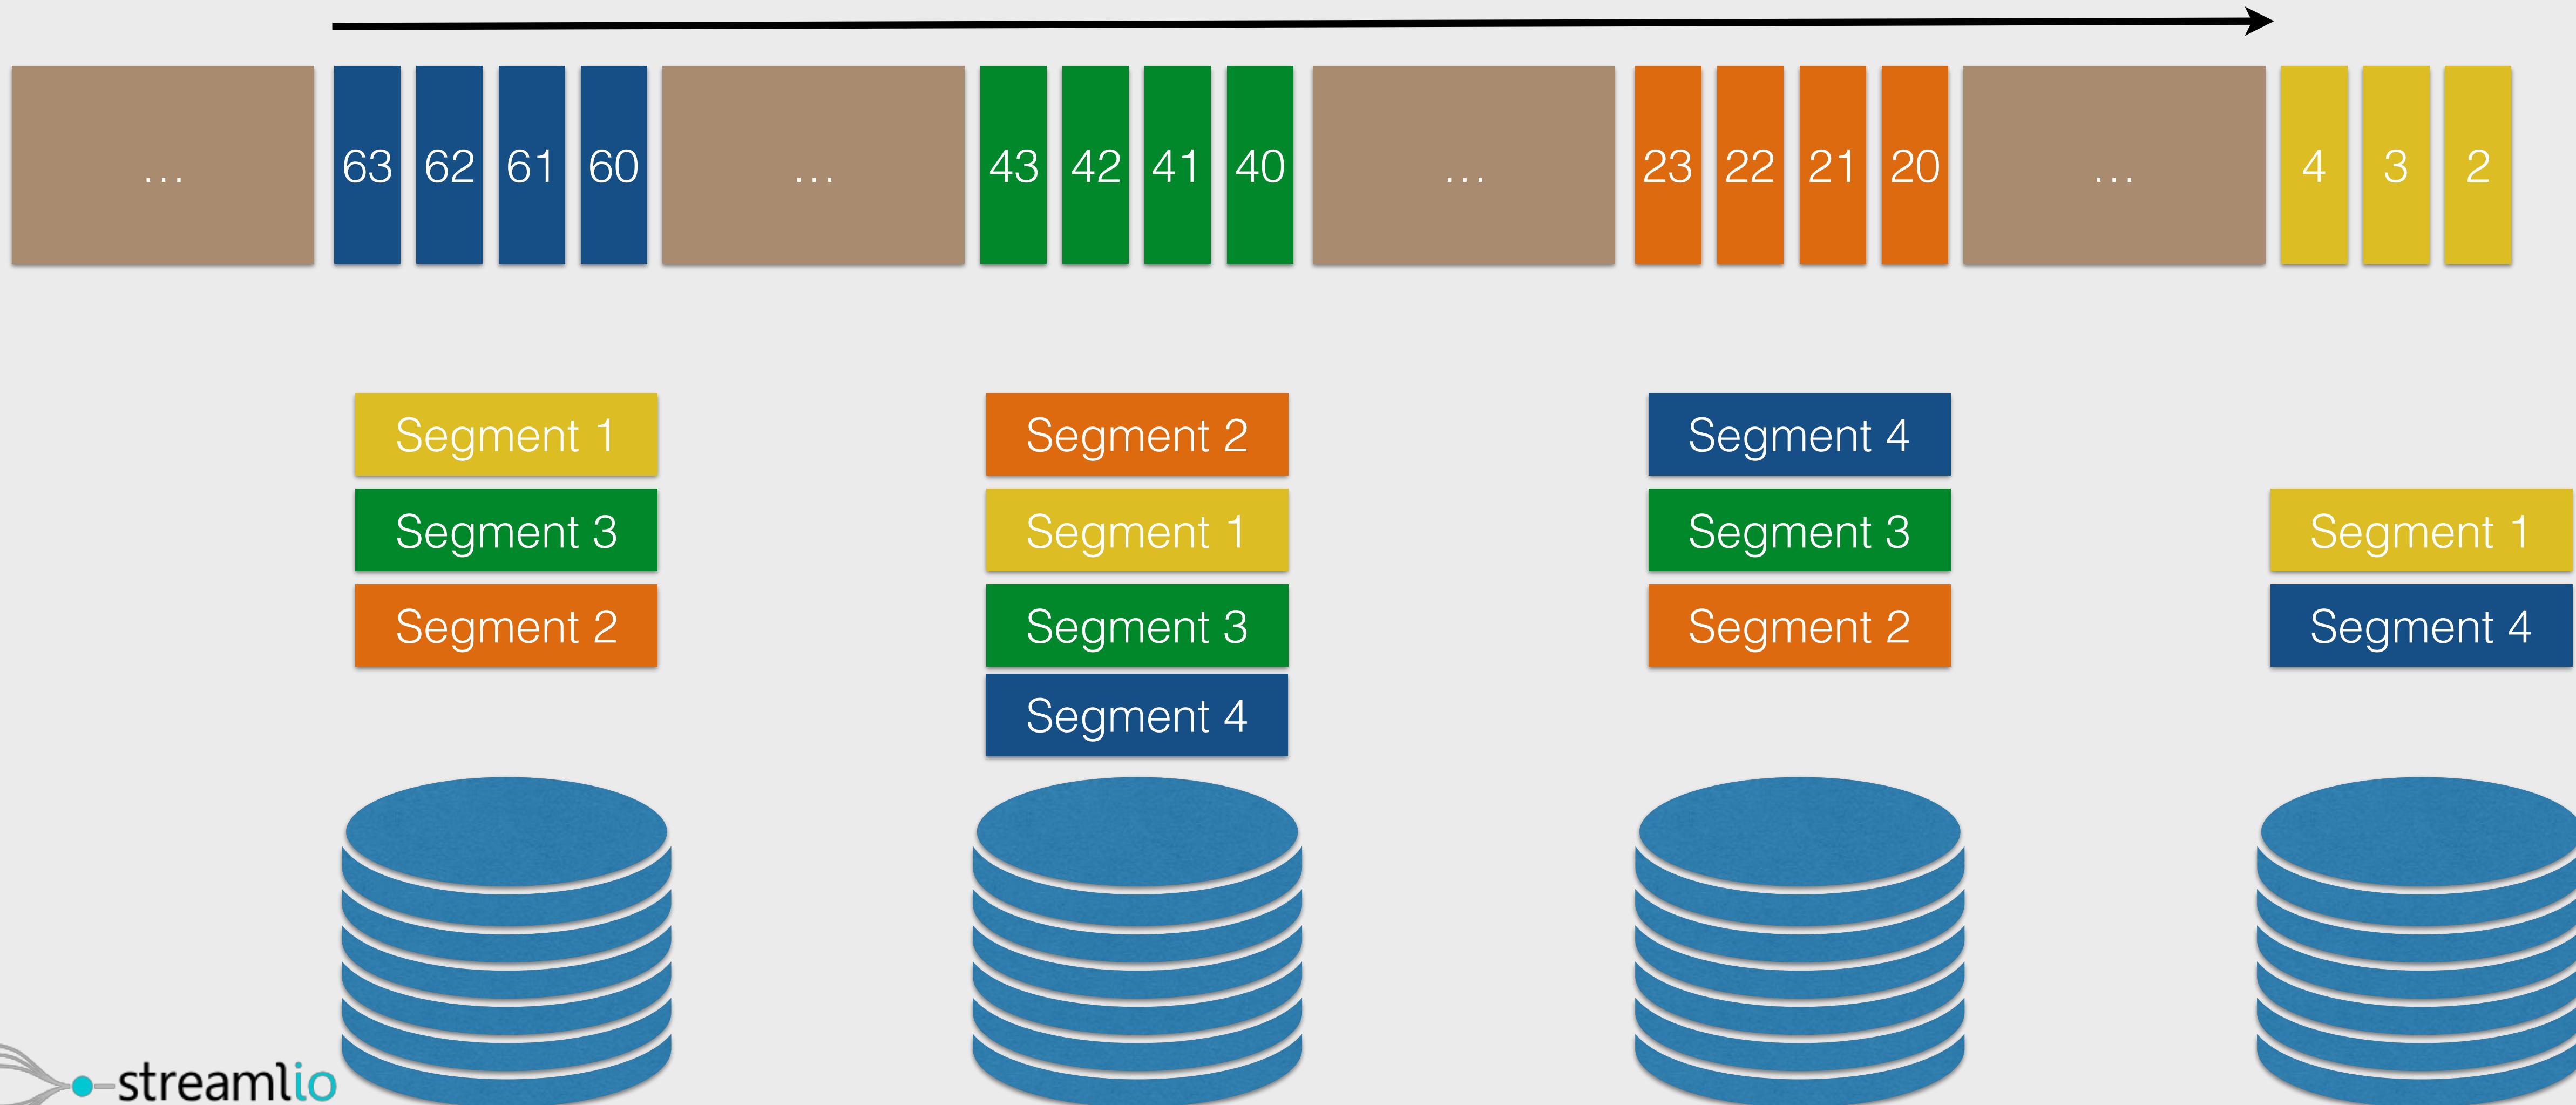
# APACHE PULSAR - DURABILITY (NO DATA LOSS)



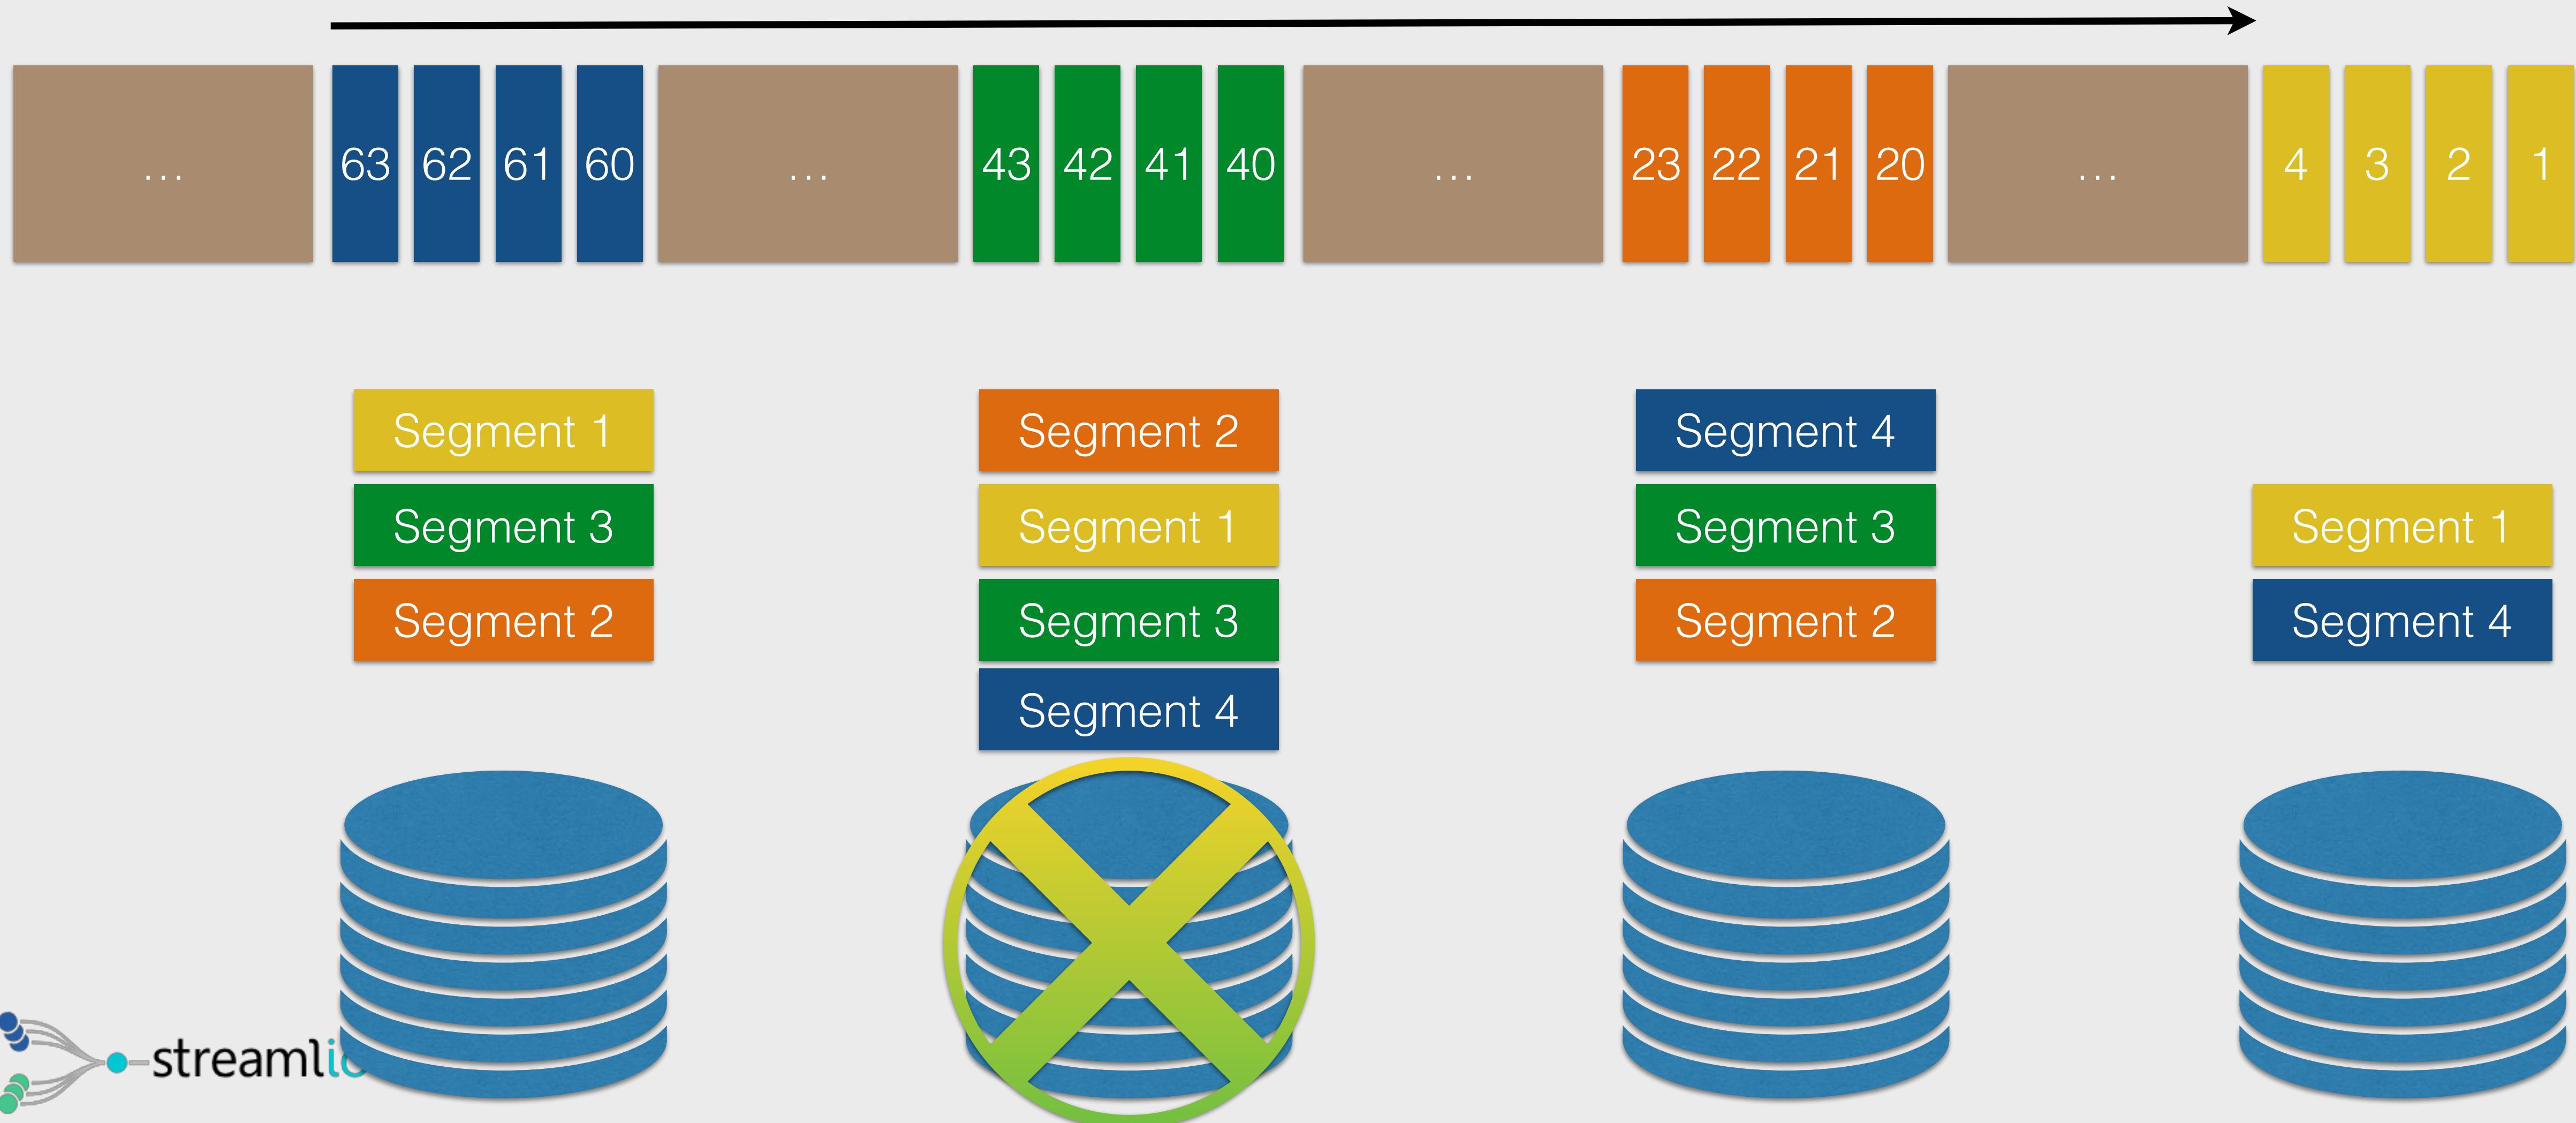
# APACHE PULSAR - ISOLATION



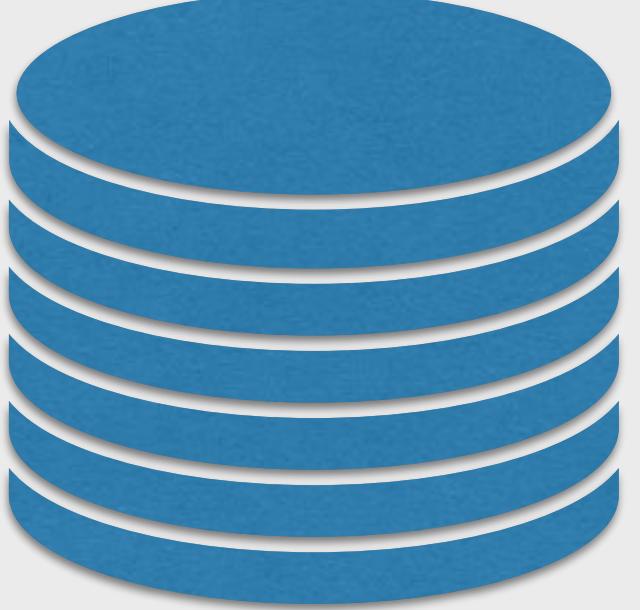
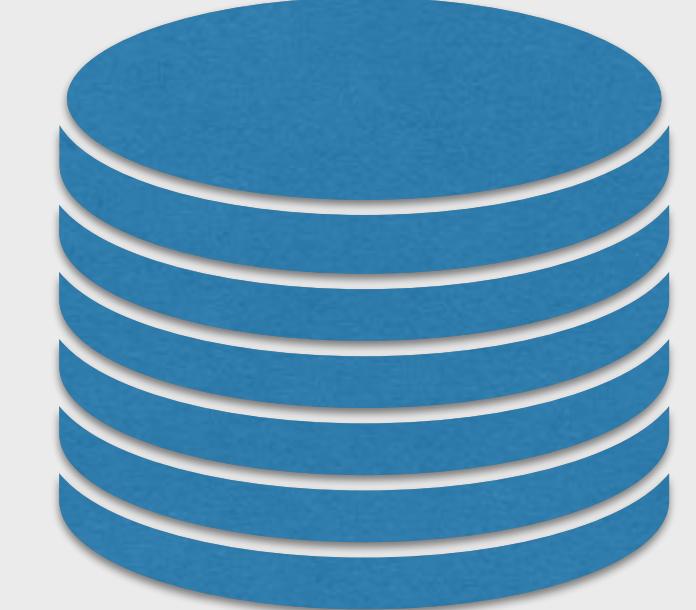
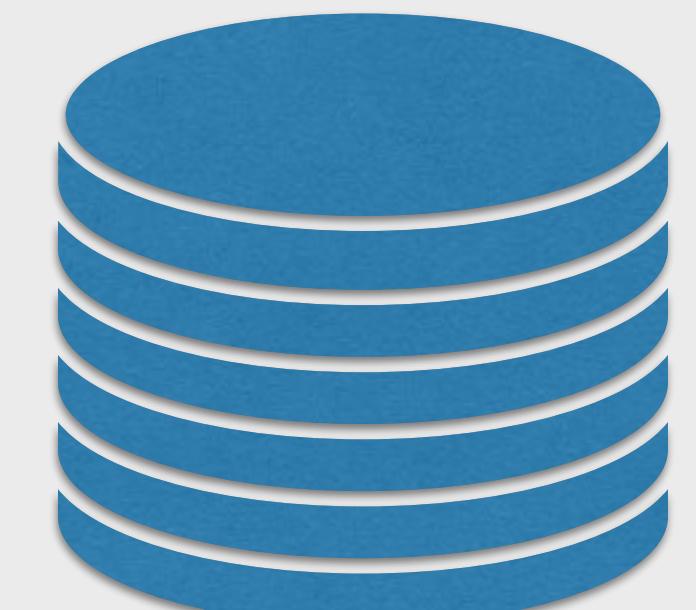
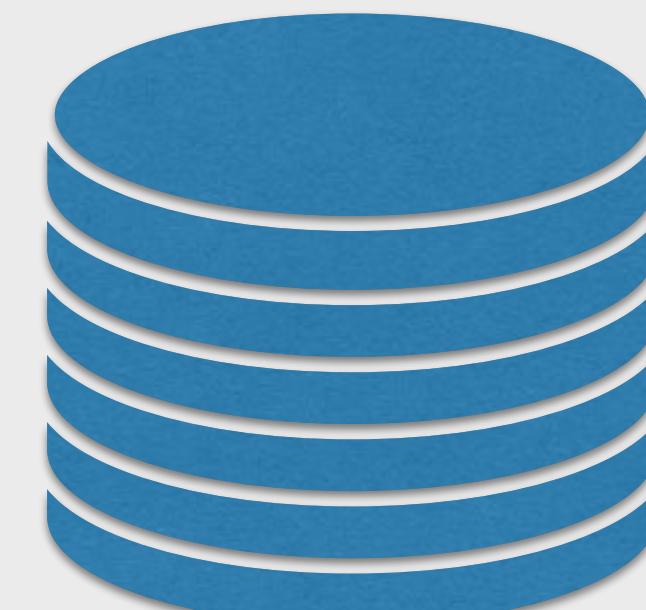
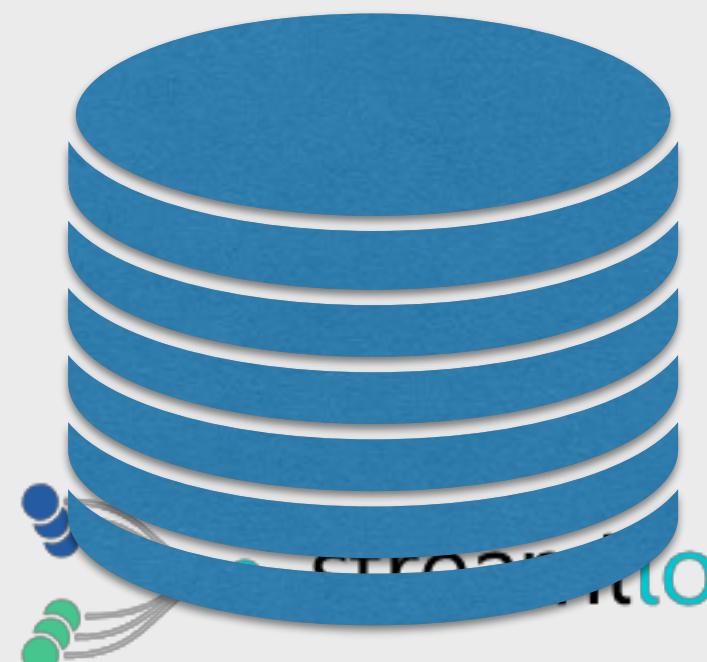
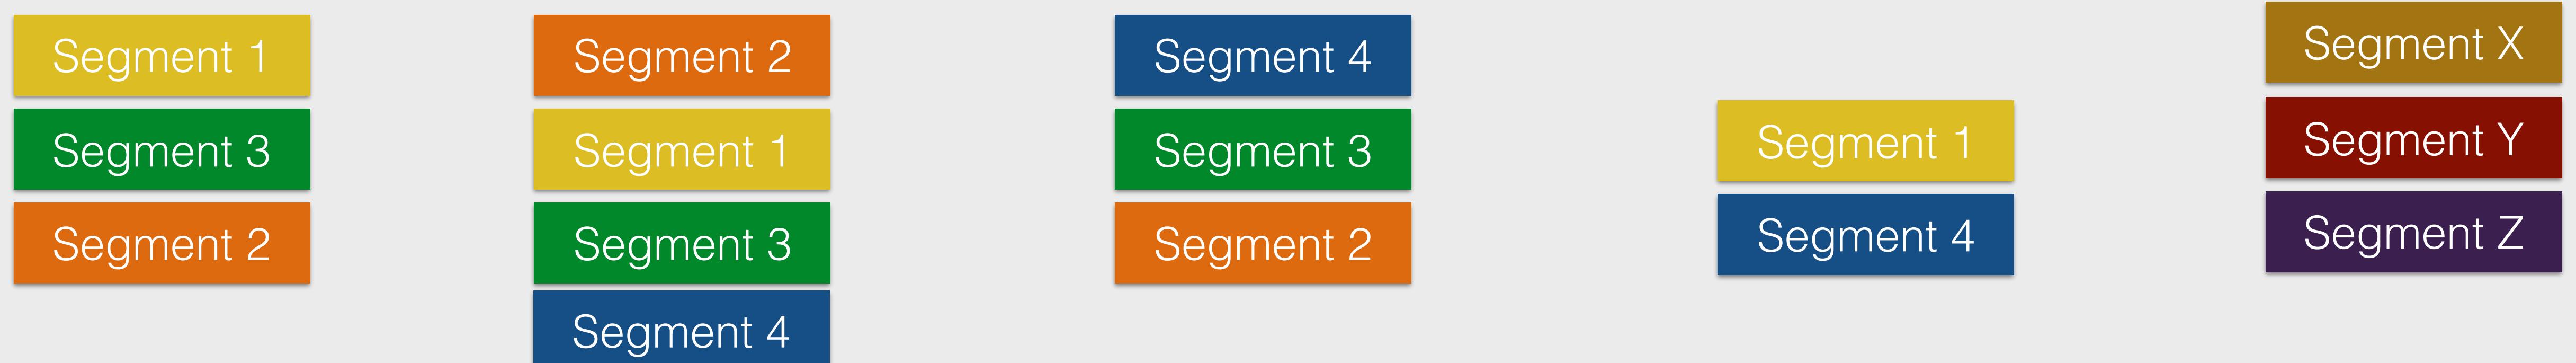
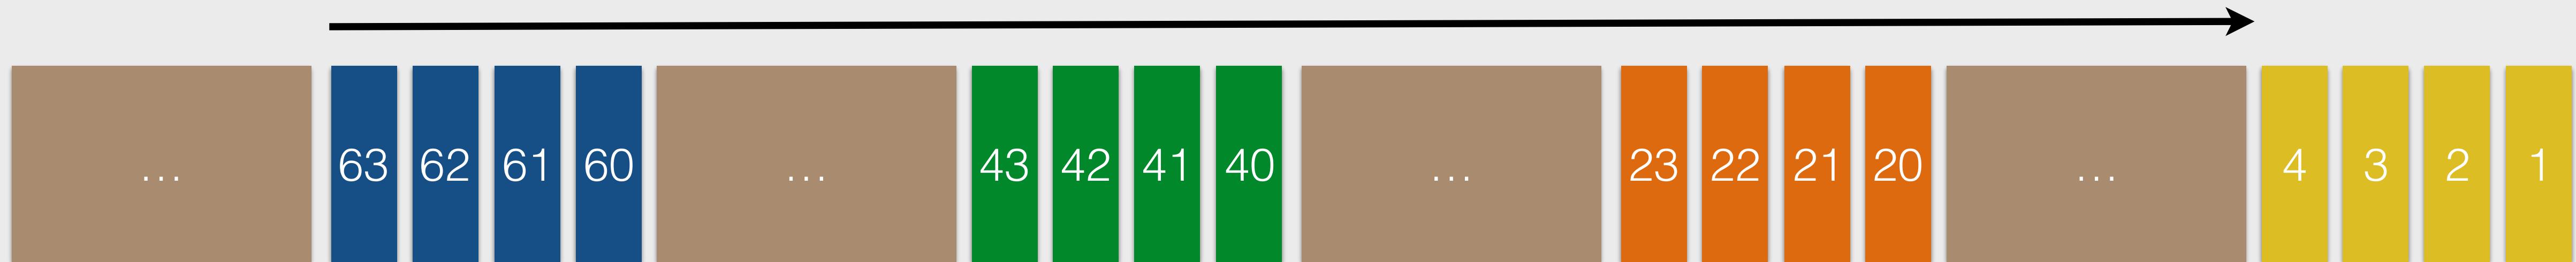
# APACHE PULSAR - SEGMENT STORAGE



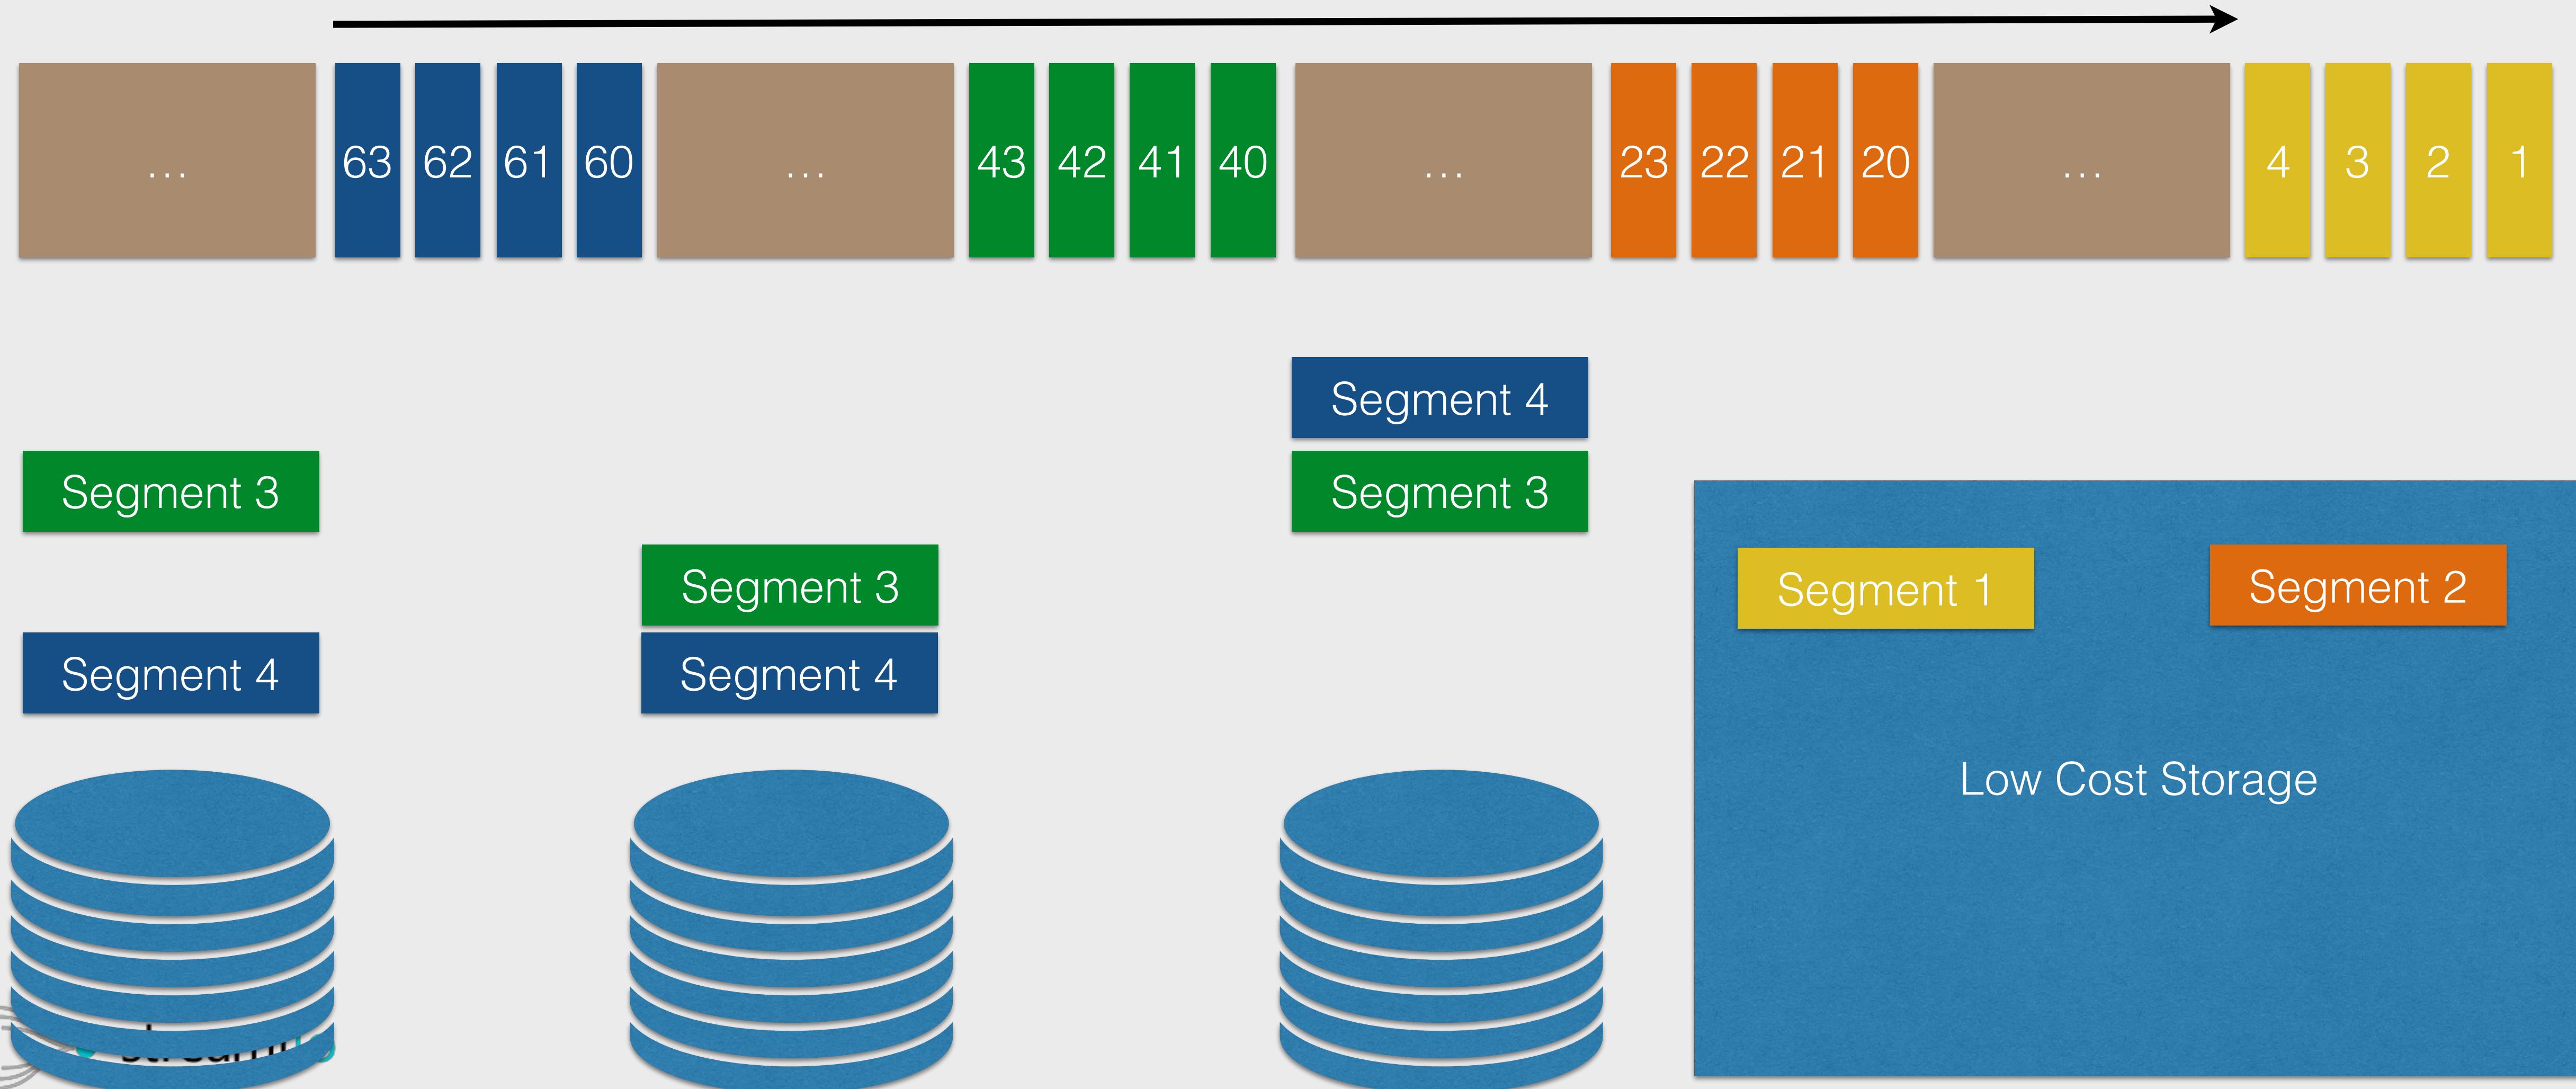
# APACHE PULSAR - RESILIENCY



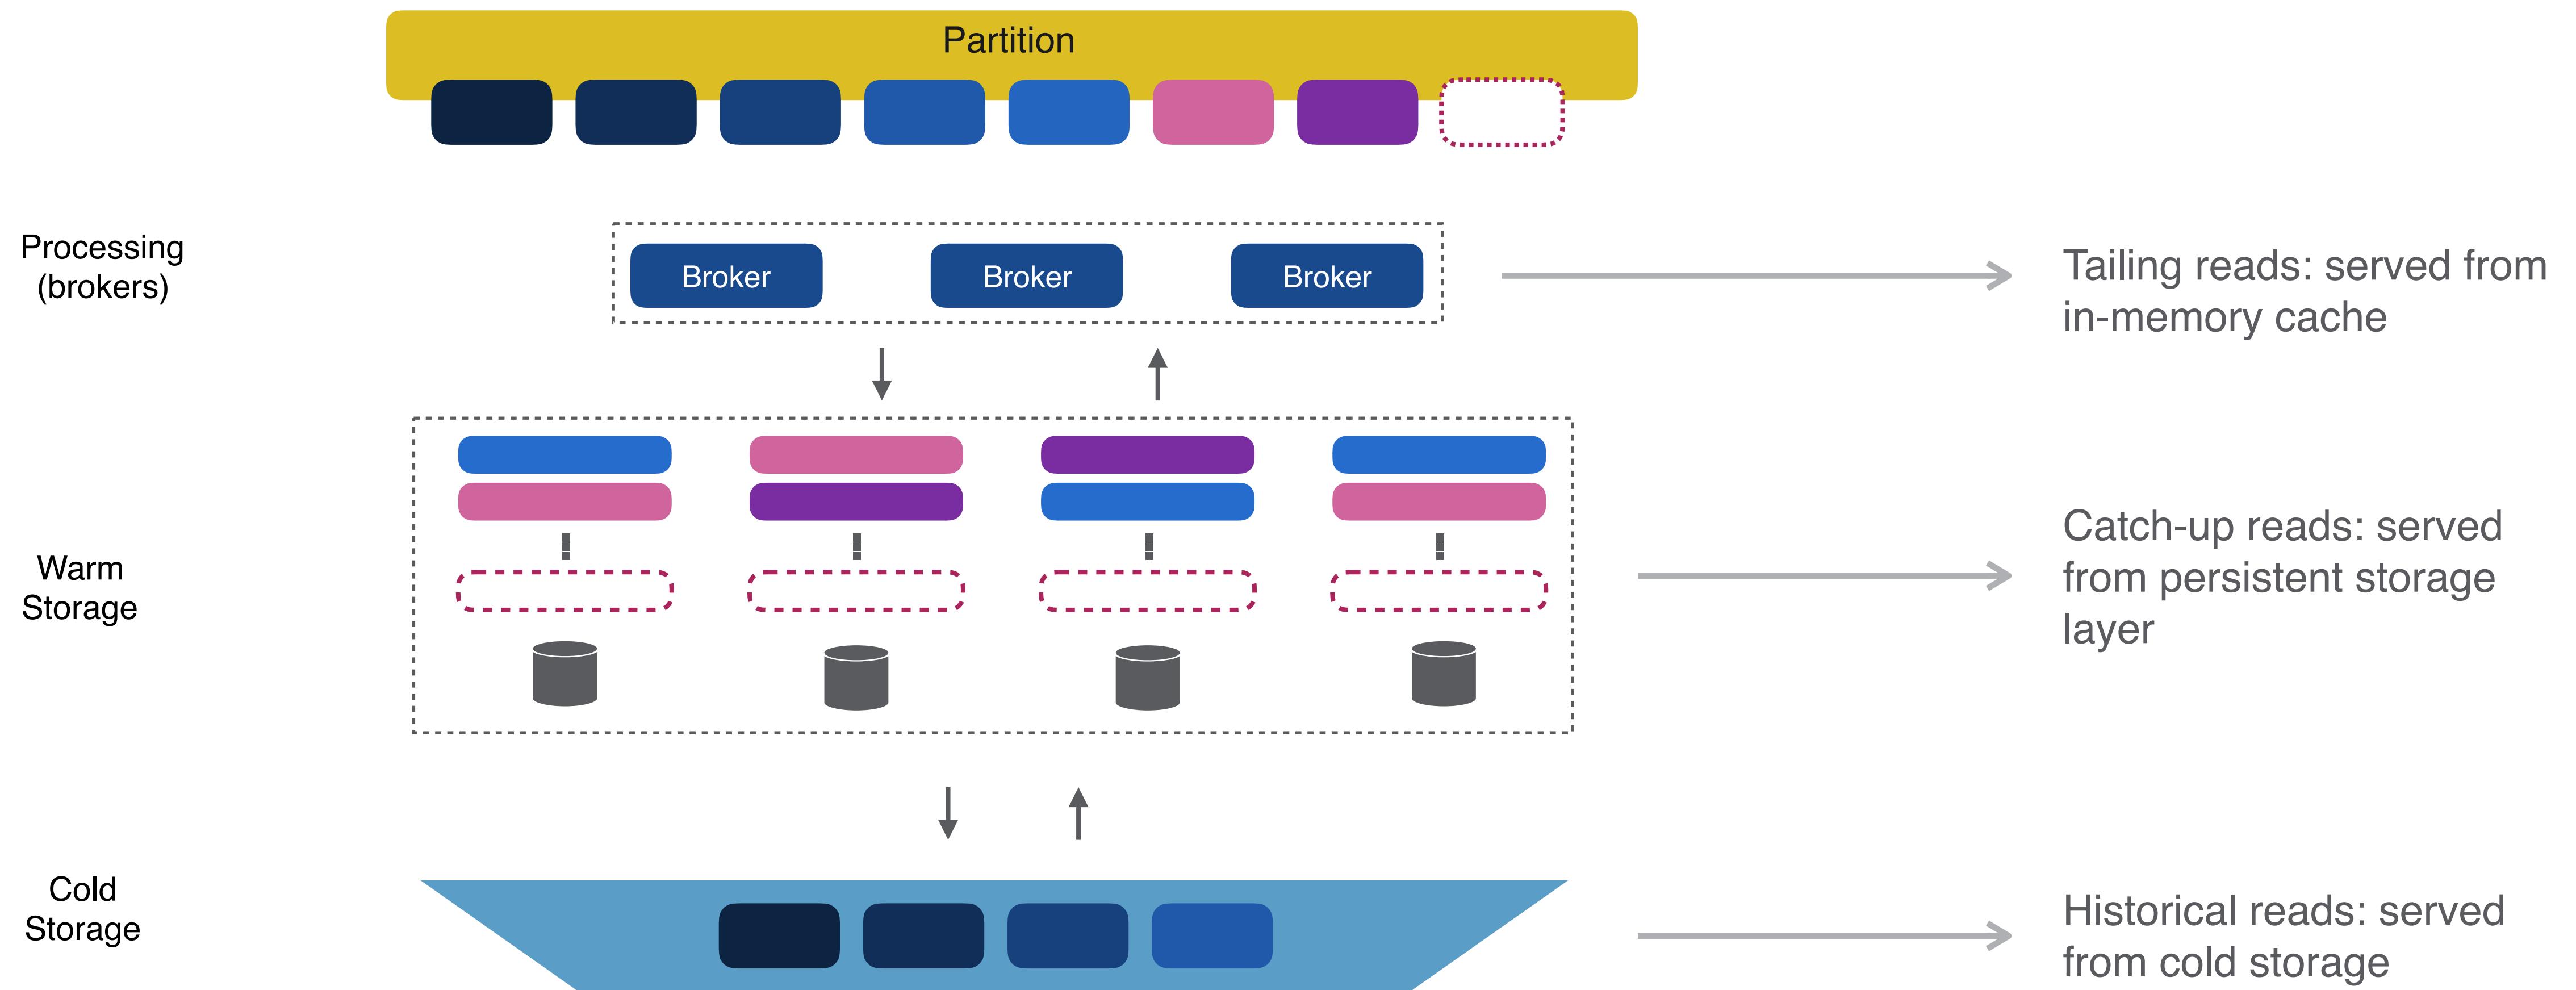
# APACHE PULSAR - SEAMLESS CLUSTER EXPANSION



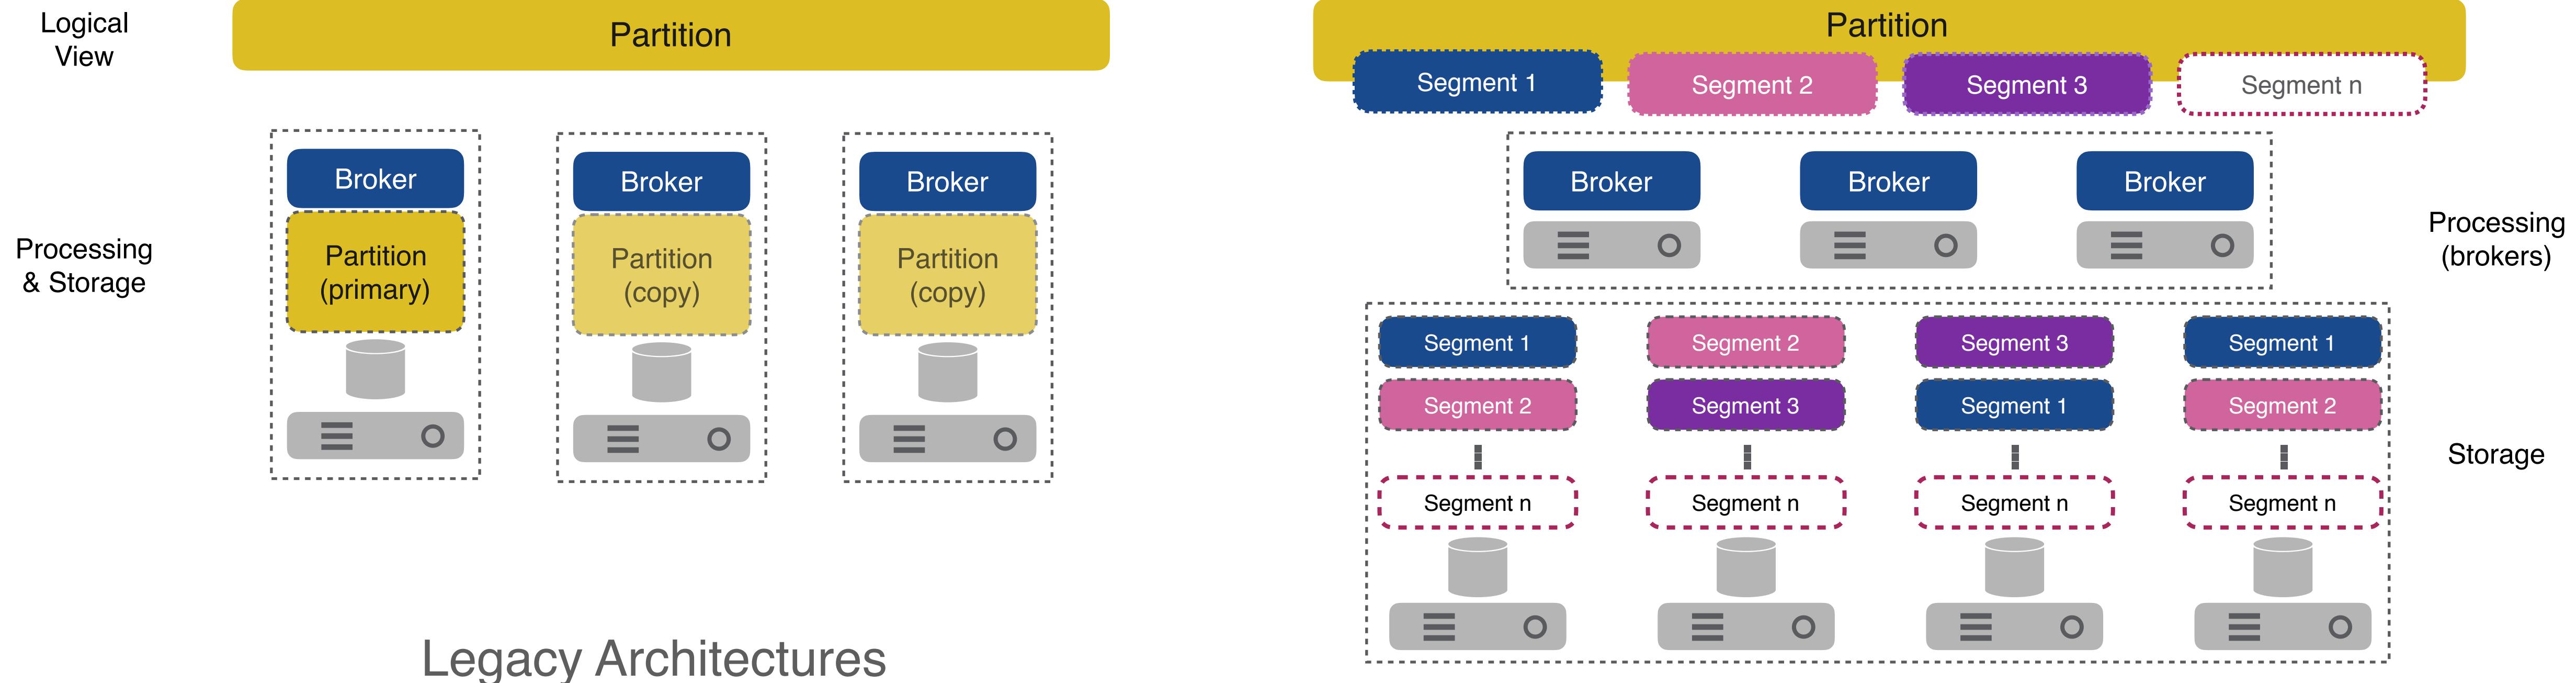
# APACHE PULSAR - TIERED STORAGE



# Multi-tiered storage and serving



# PARTITIONS VS SEGMENTS - WHY SHOULD YOU CARE?



## Legacy Architectures

- Storage co-resident with processing
- Partition-centric
- Cumbersome to scale--data redistribution, performance impact

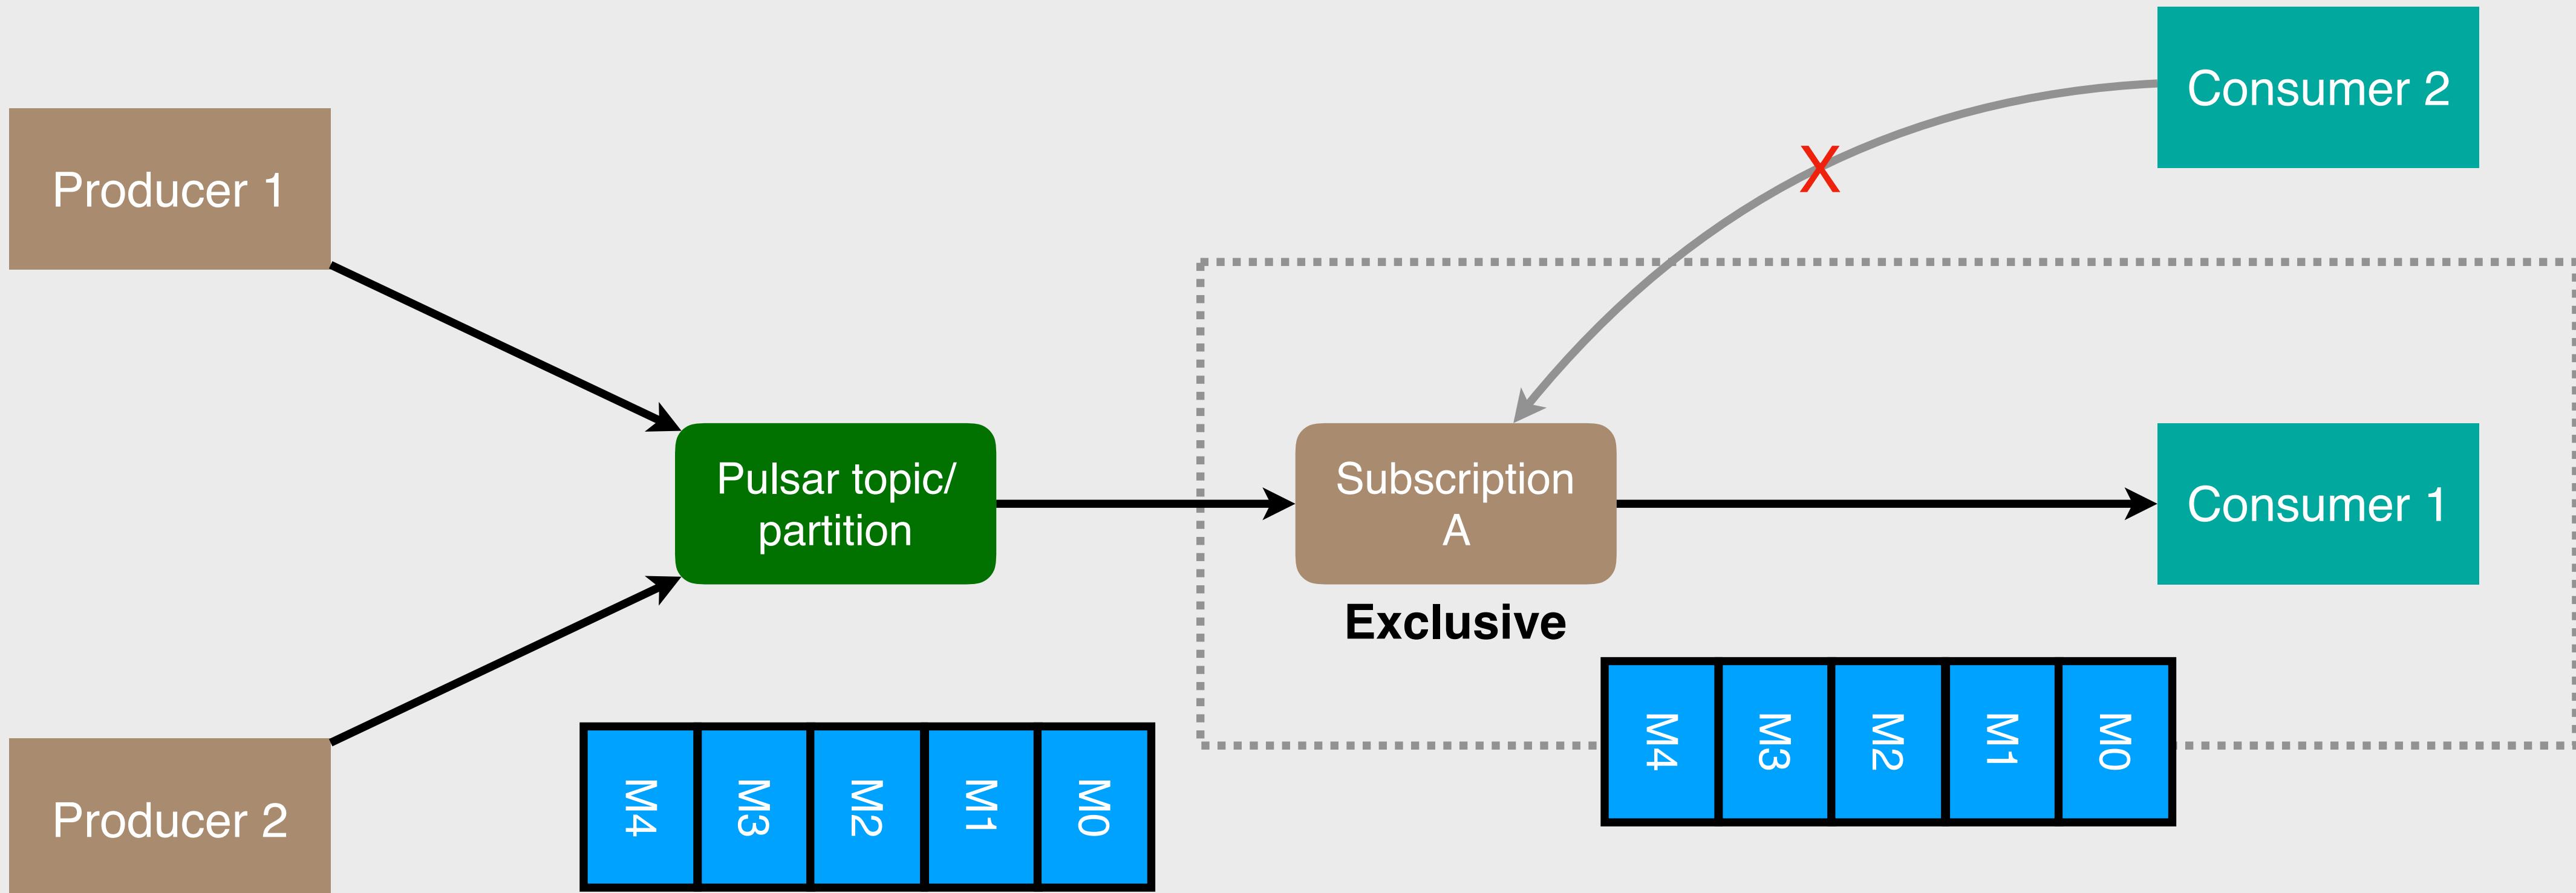
## Apache Pulsar

- Storage decoupled from processing
- Partitions stored as segments
- Flexible, easy scalability

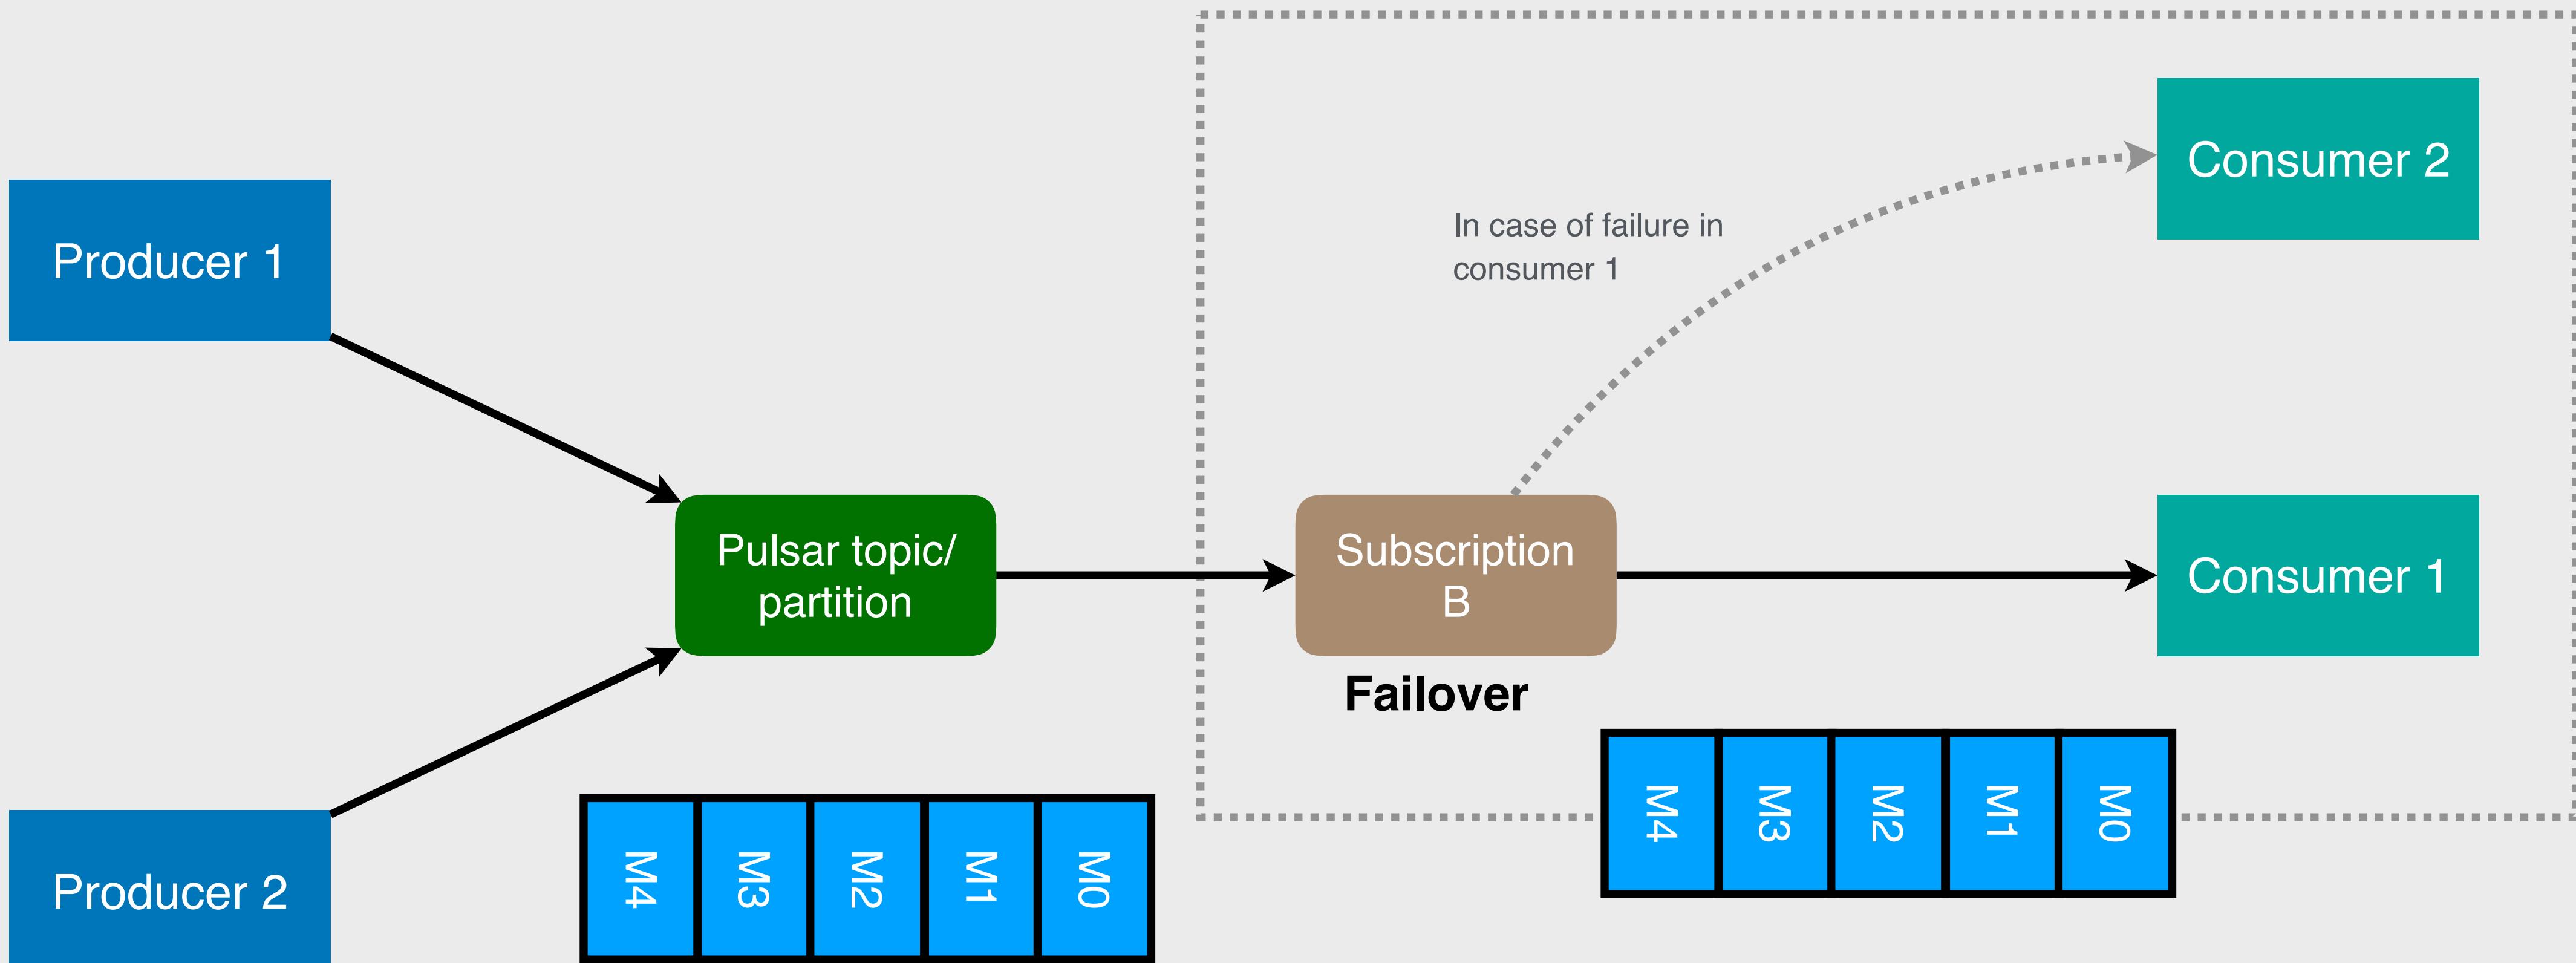
# PARTITIONS VS SEGMENTS - WHY SHOULD YOU CARE?

- ◆ In Kafka, partitions are assigned to brokers “permanently”
- ◆ A single partition is stored entirely in a single node
- ◆ Retention is limited by a single node storage capacity
- ◆ Failure recovery and capacity expansion require expensive “rebalancing”
- ◆ Rebalancing has a big impact over the system, affecting regular traffic

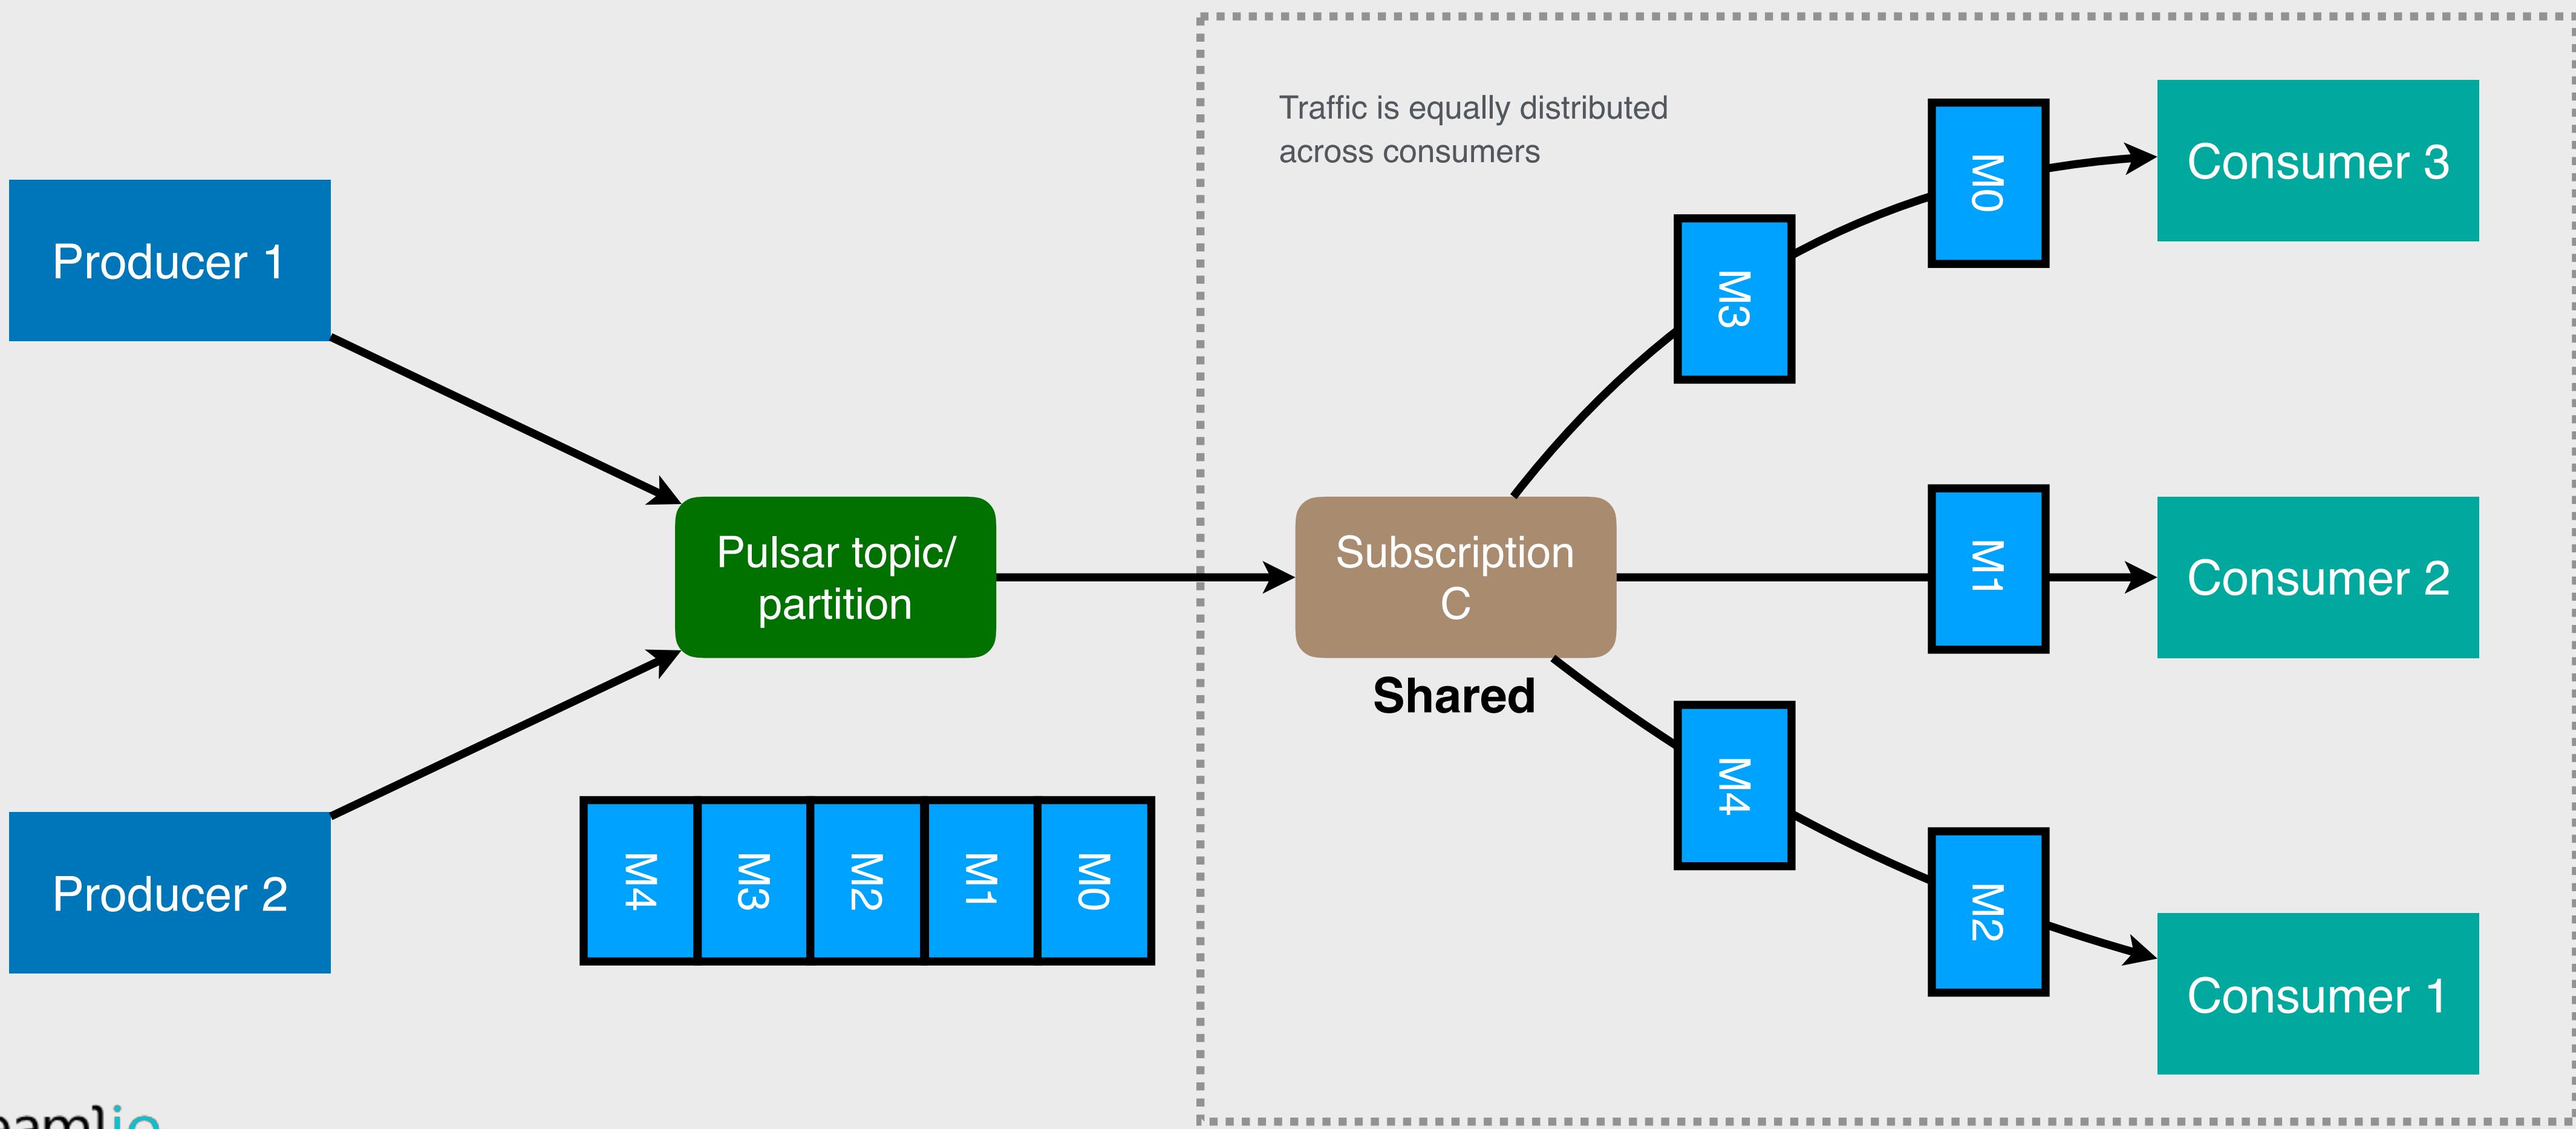
# UNIFIED MESSAGING MODEL - STREAMING



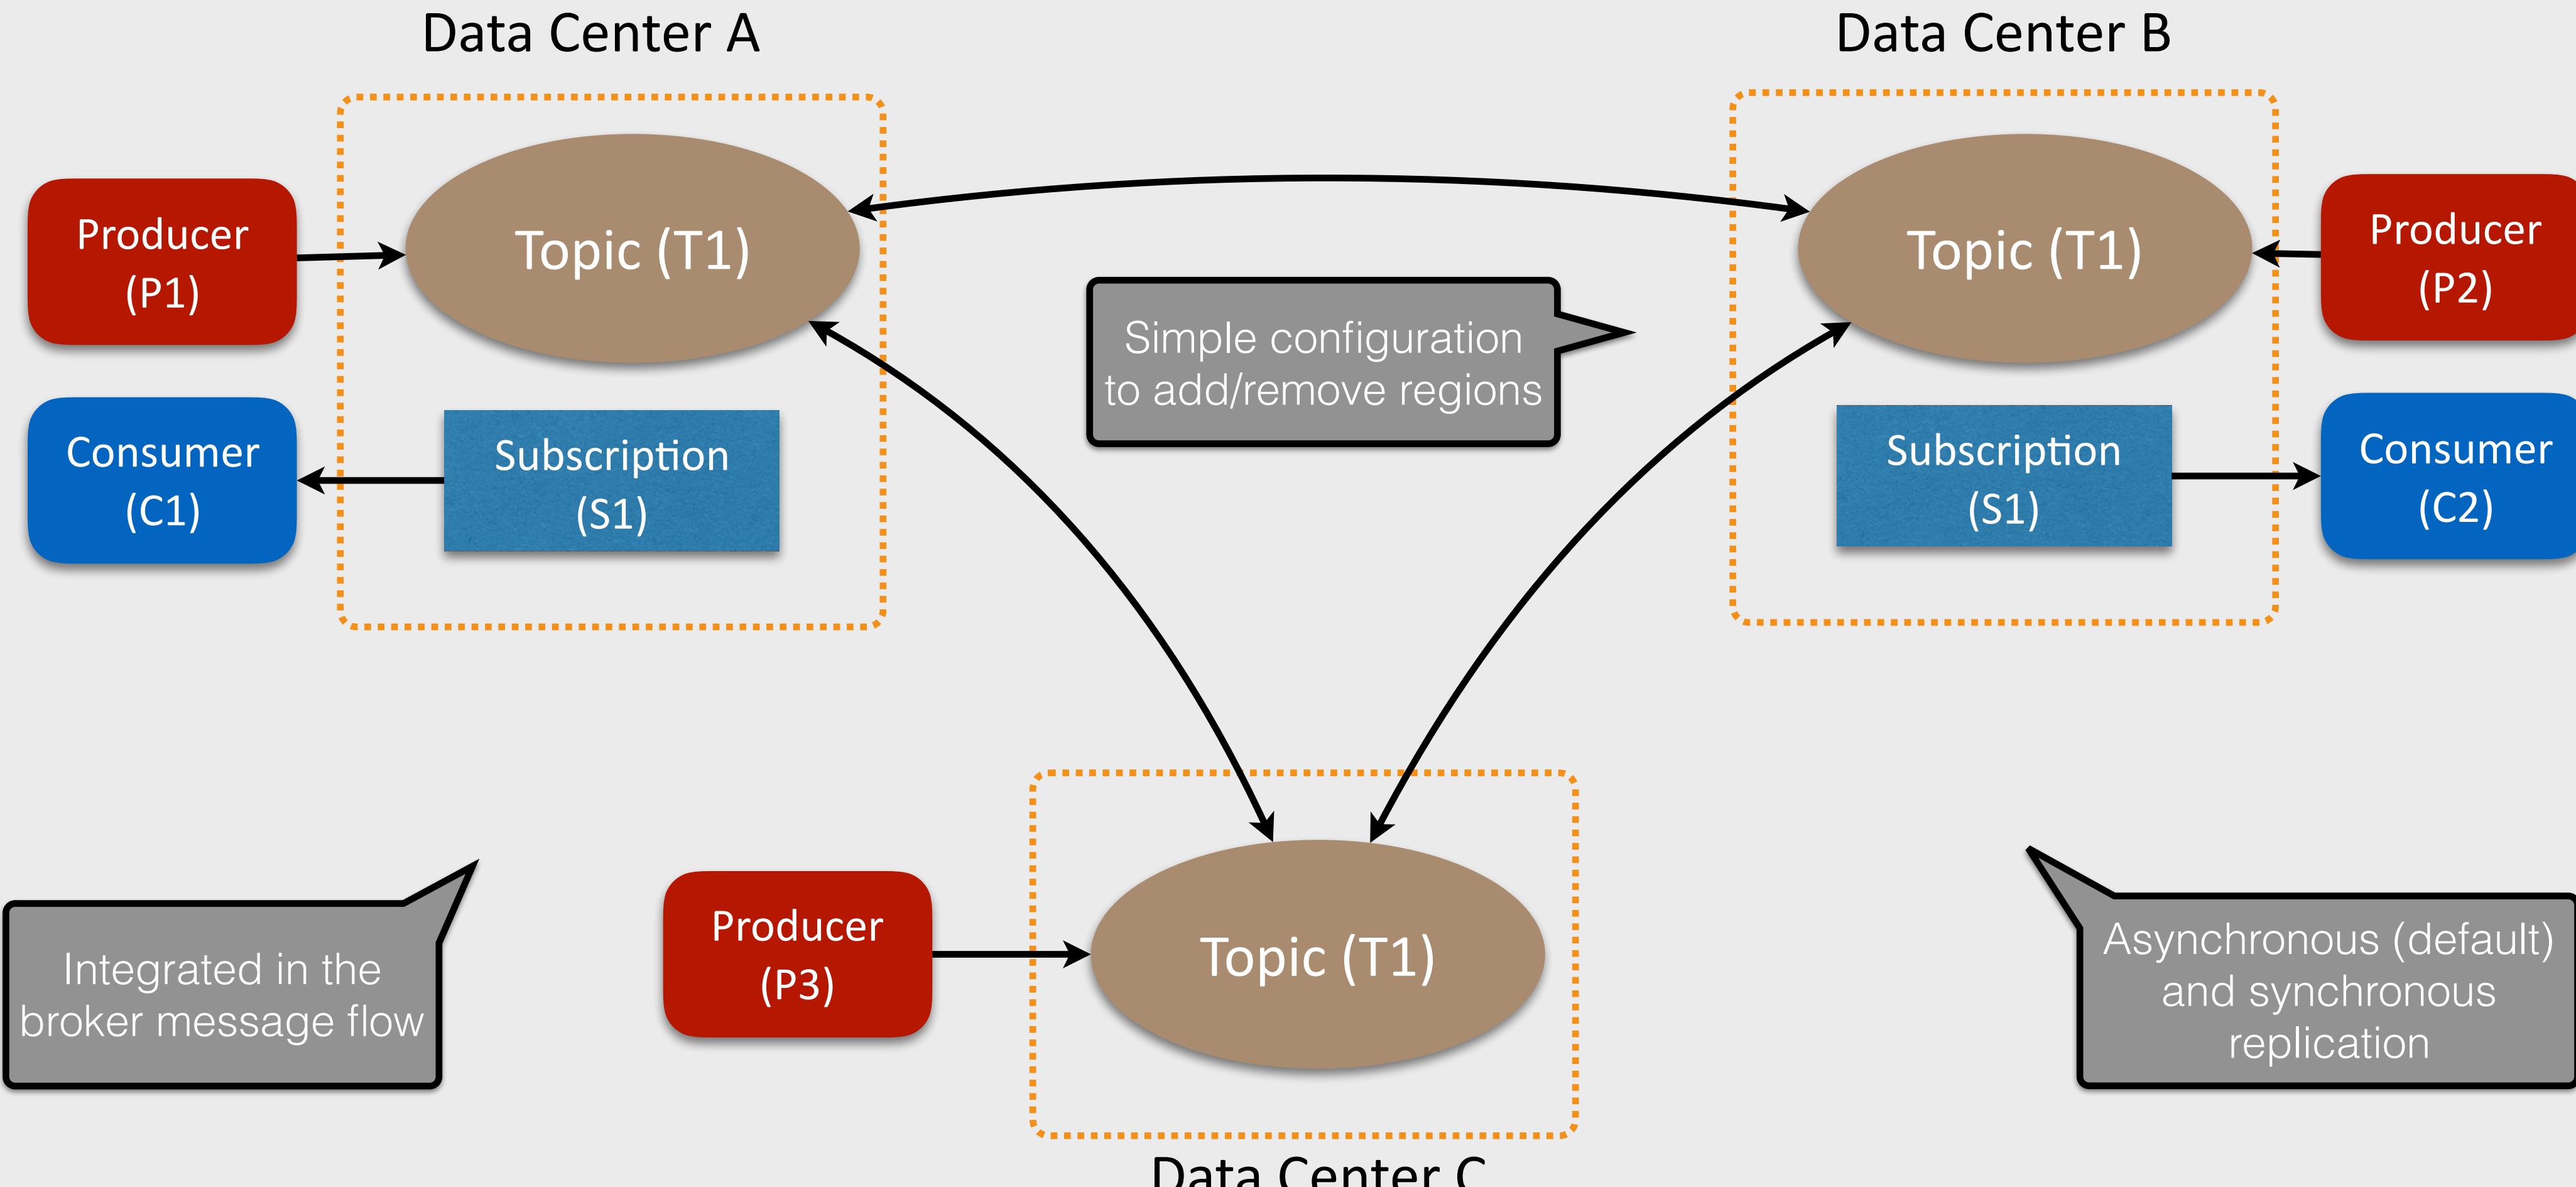
# UNIFIED MESSAGING MODEL - STREAMING



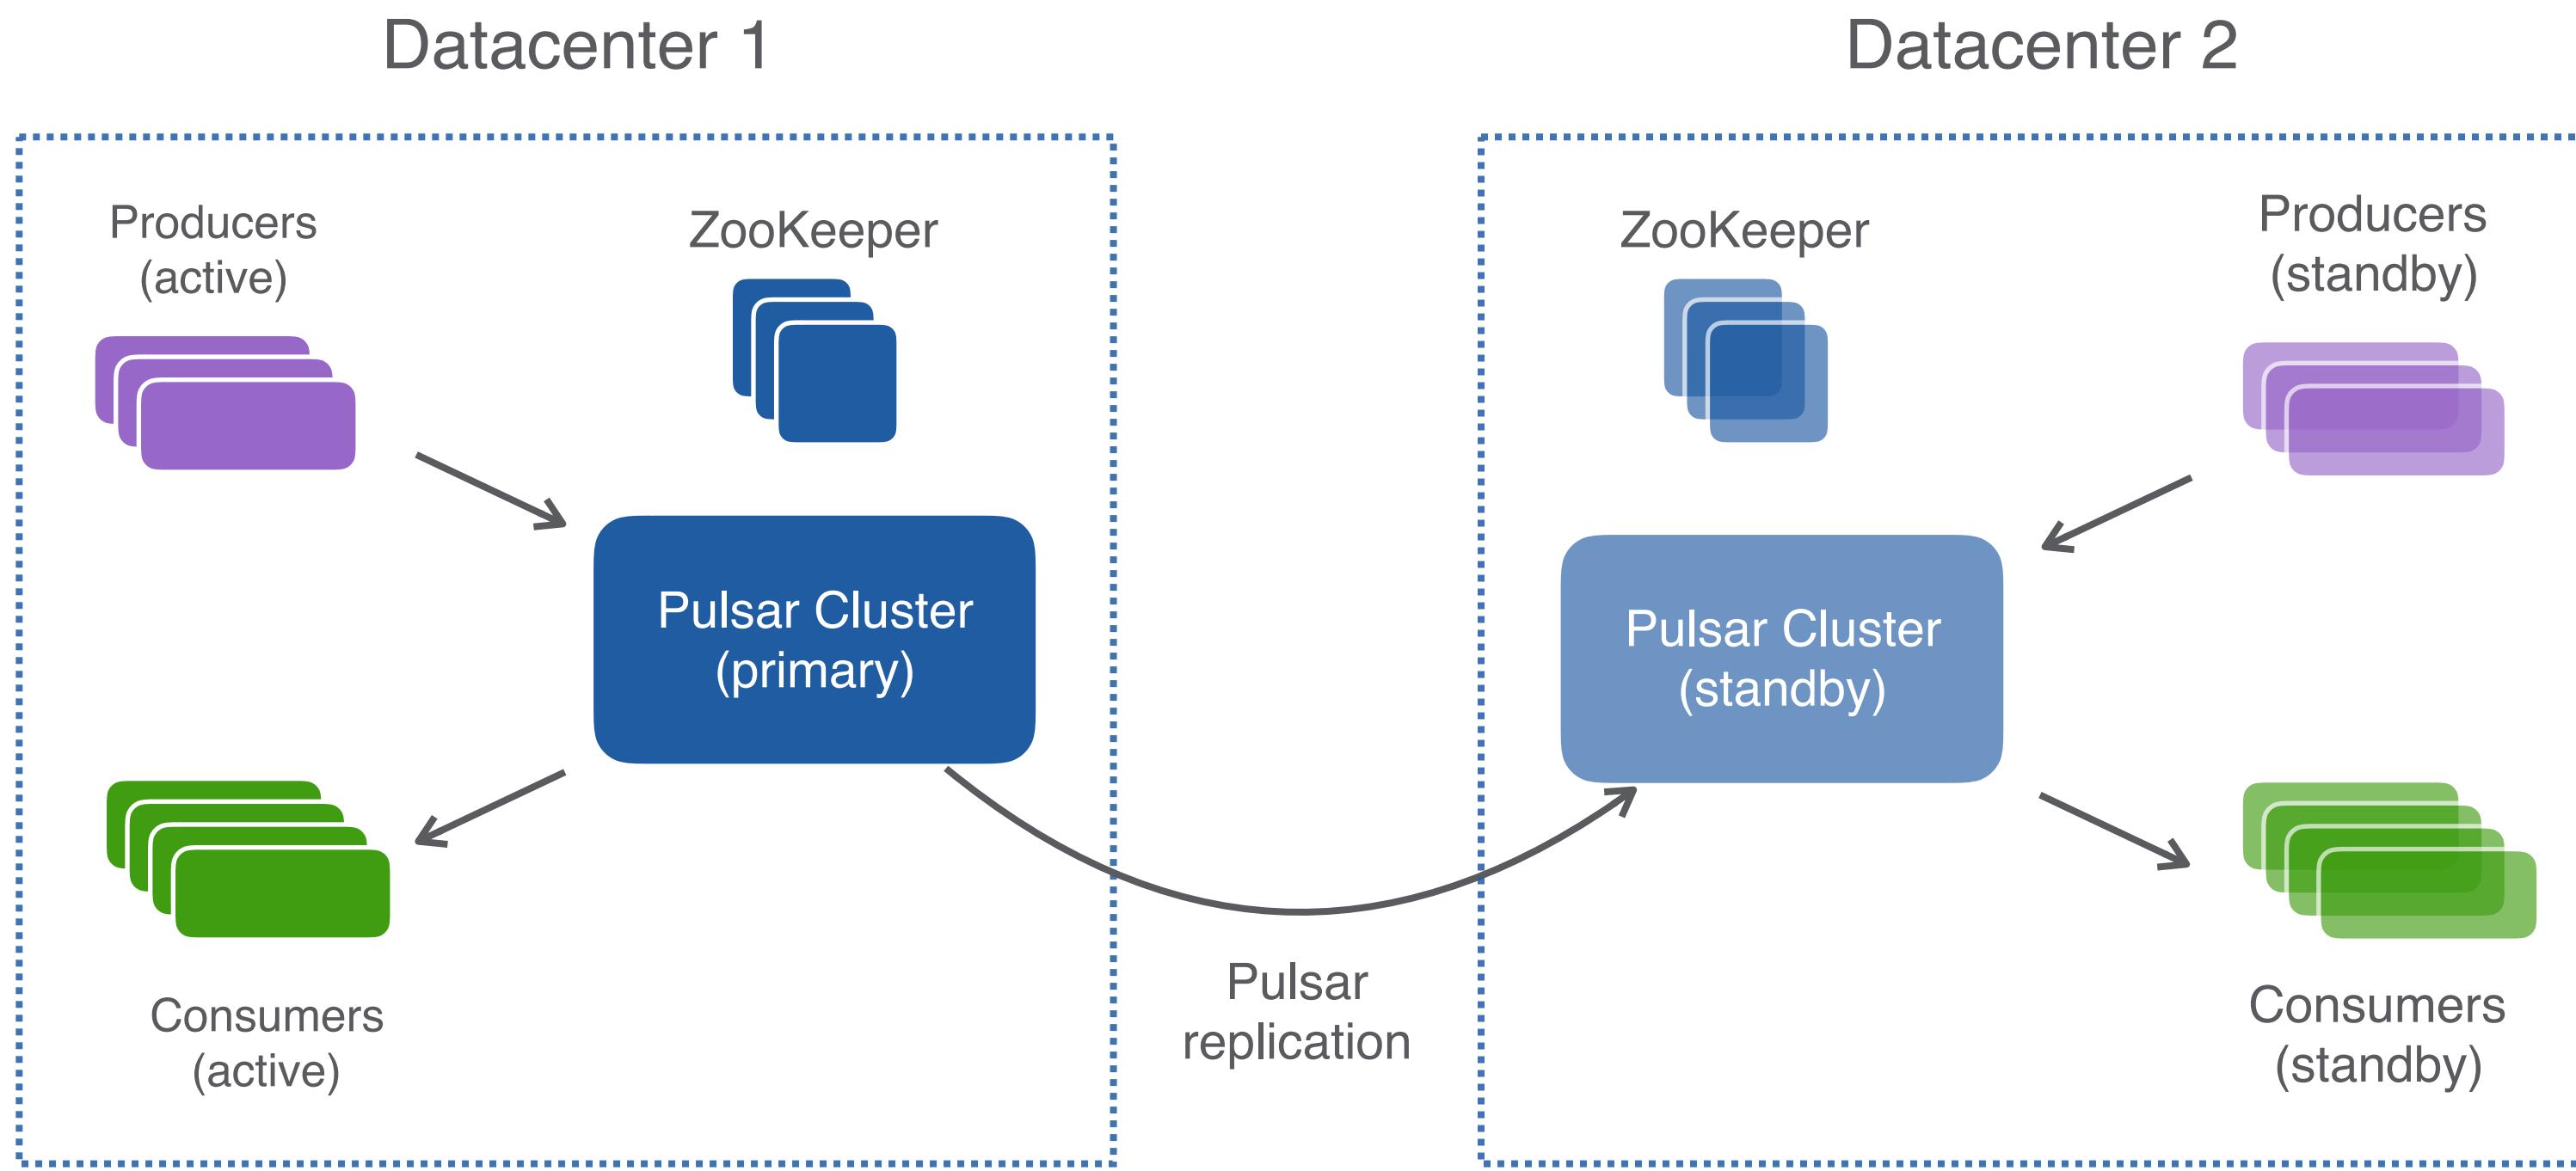
# UNIFIED MESSAGING MODEL - QUEUING



# DISASTER RECOVERY

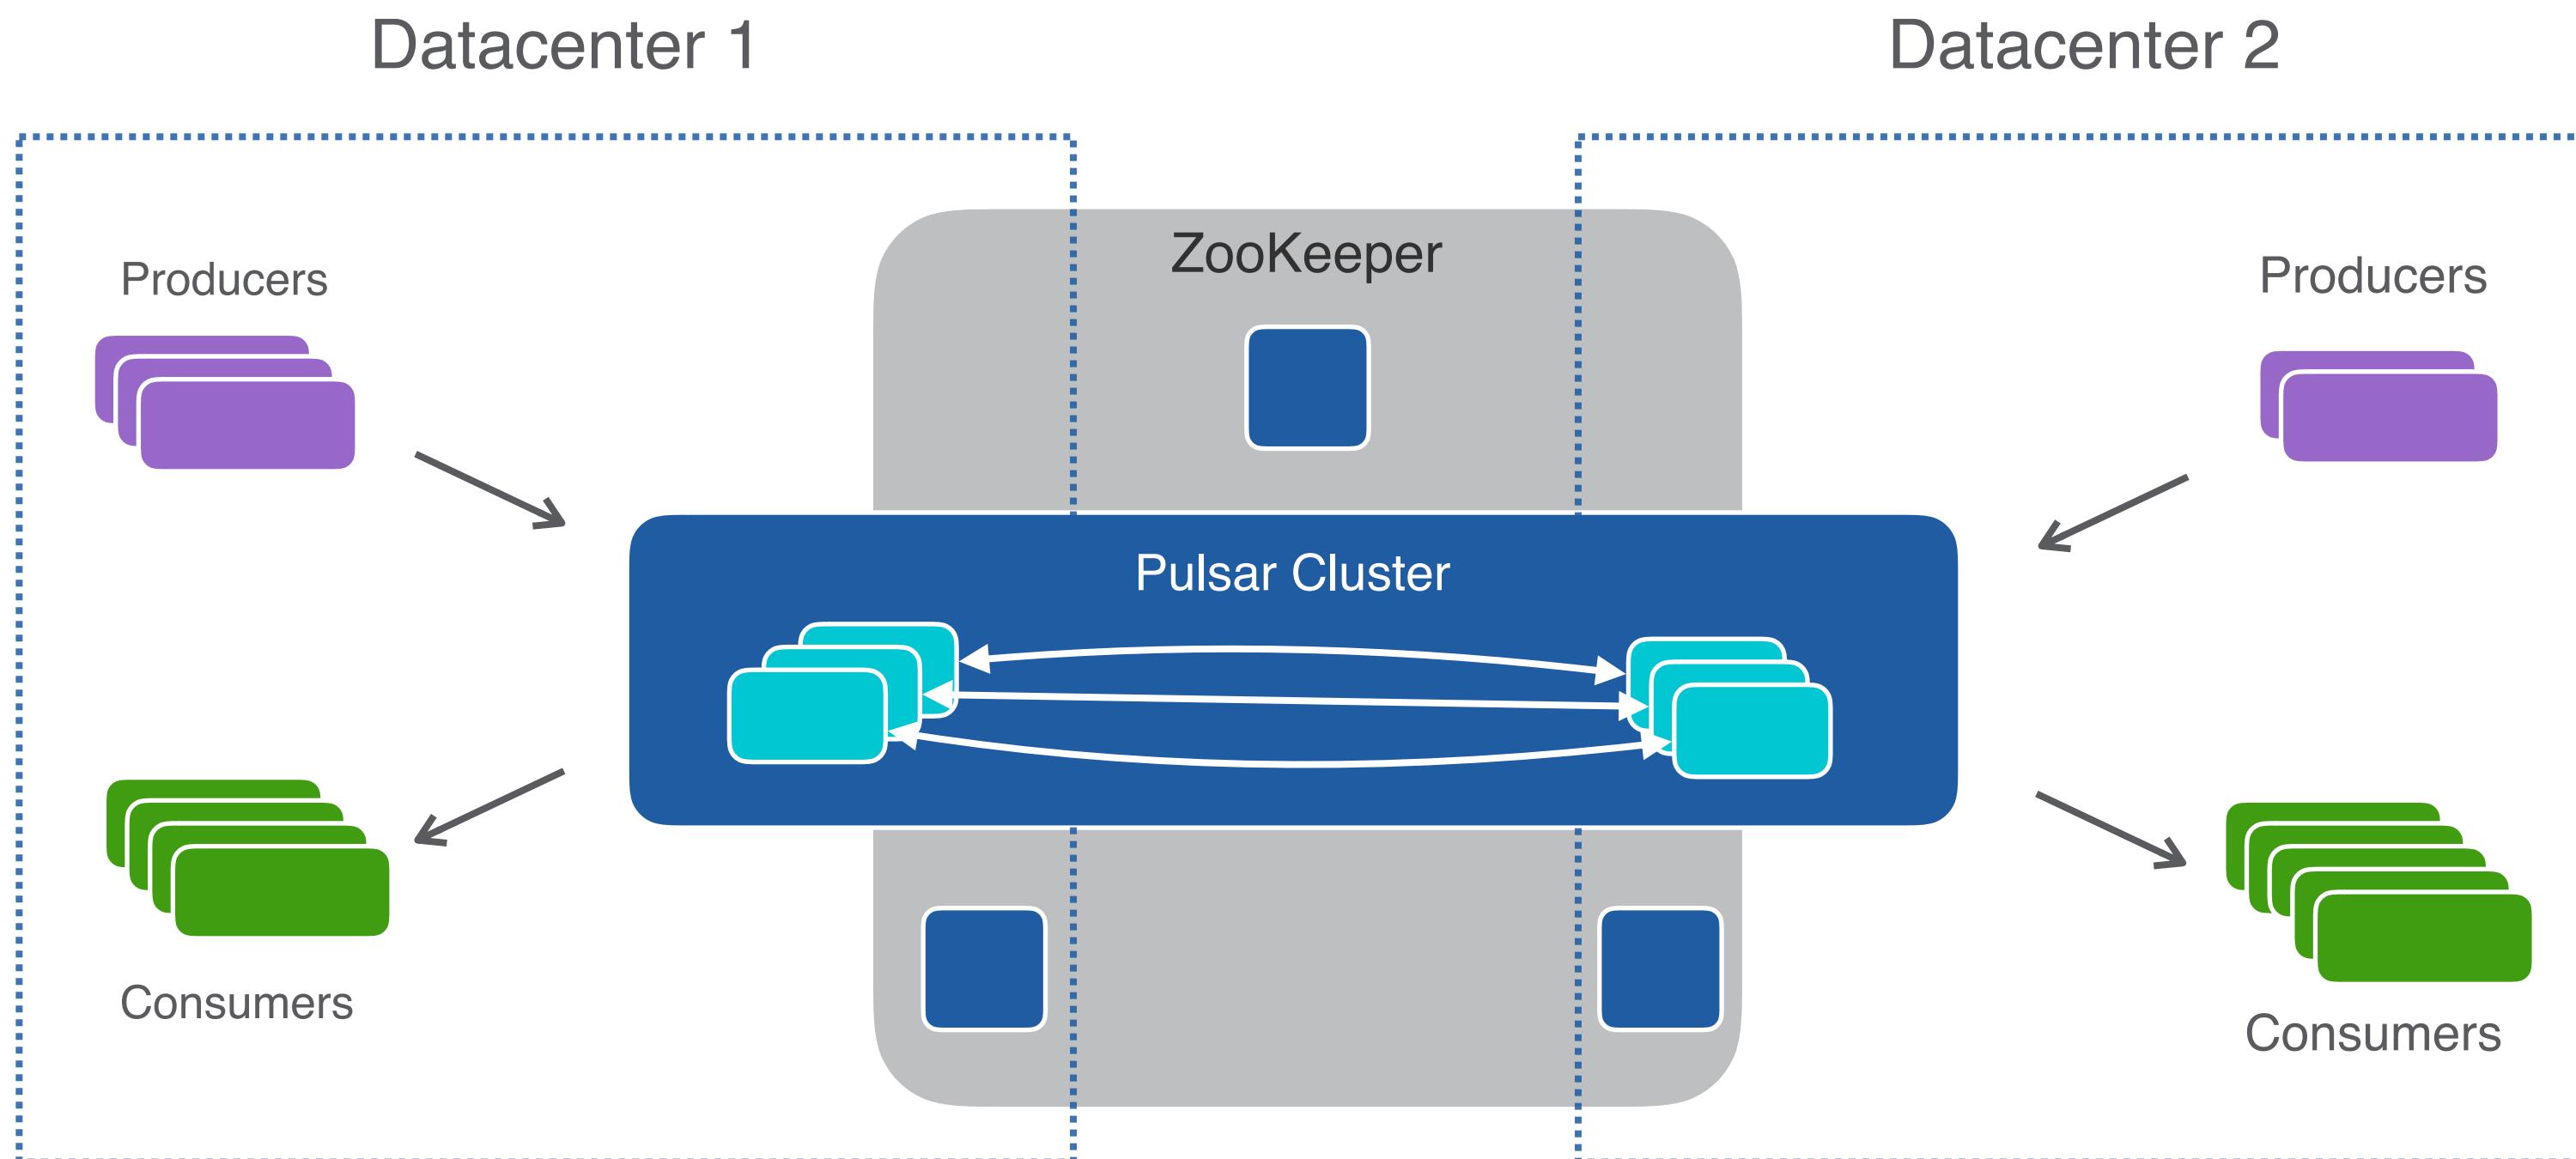


# Asynchronous replication example



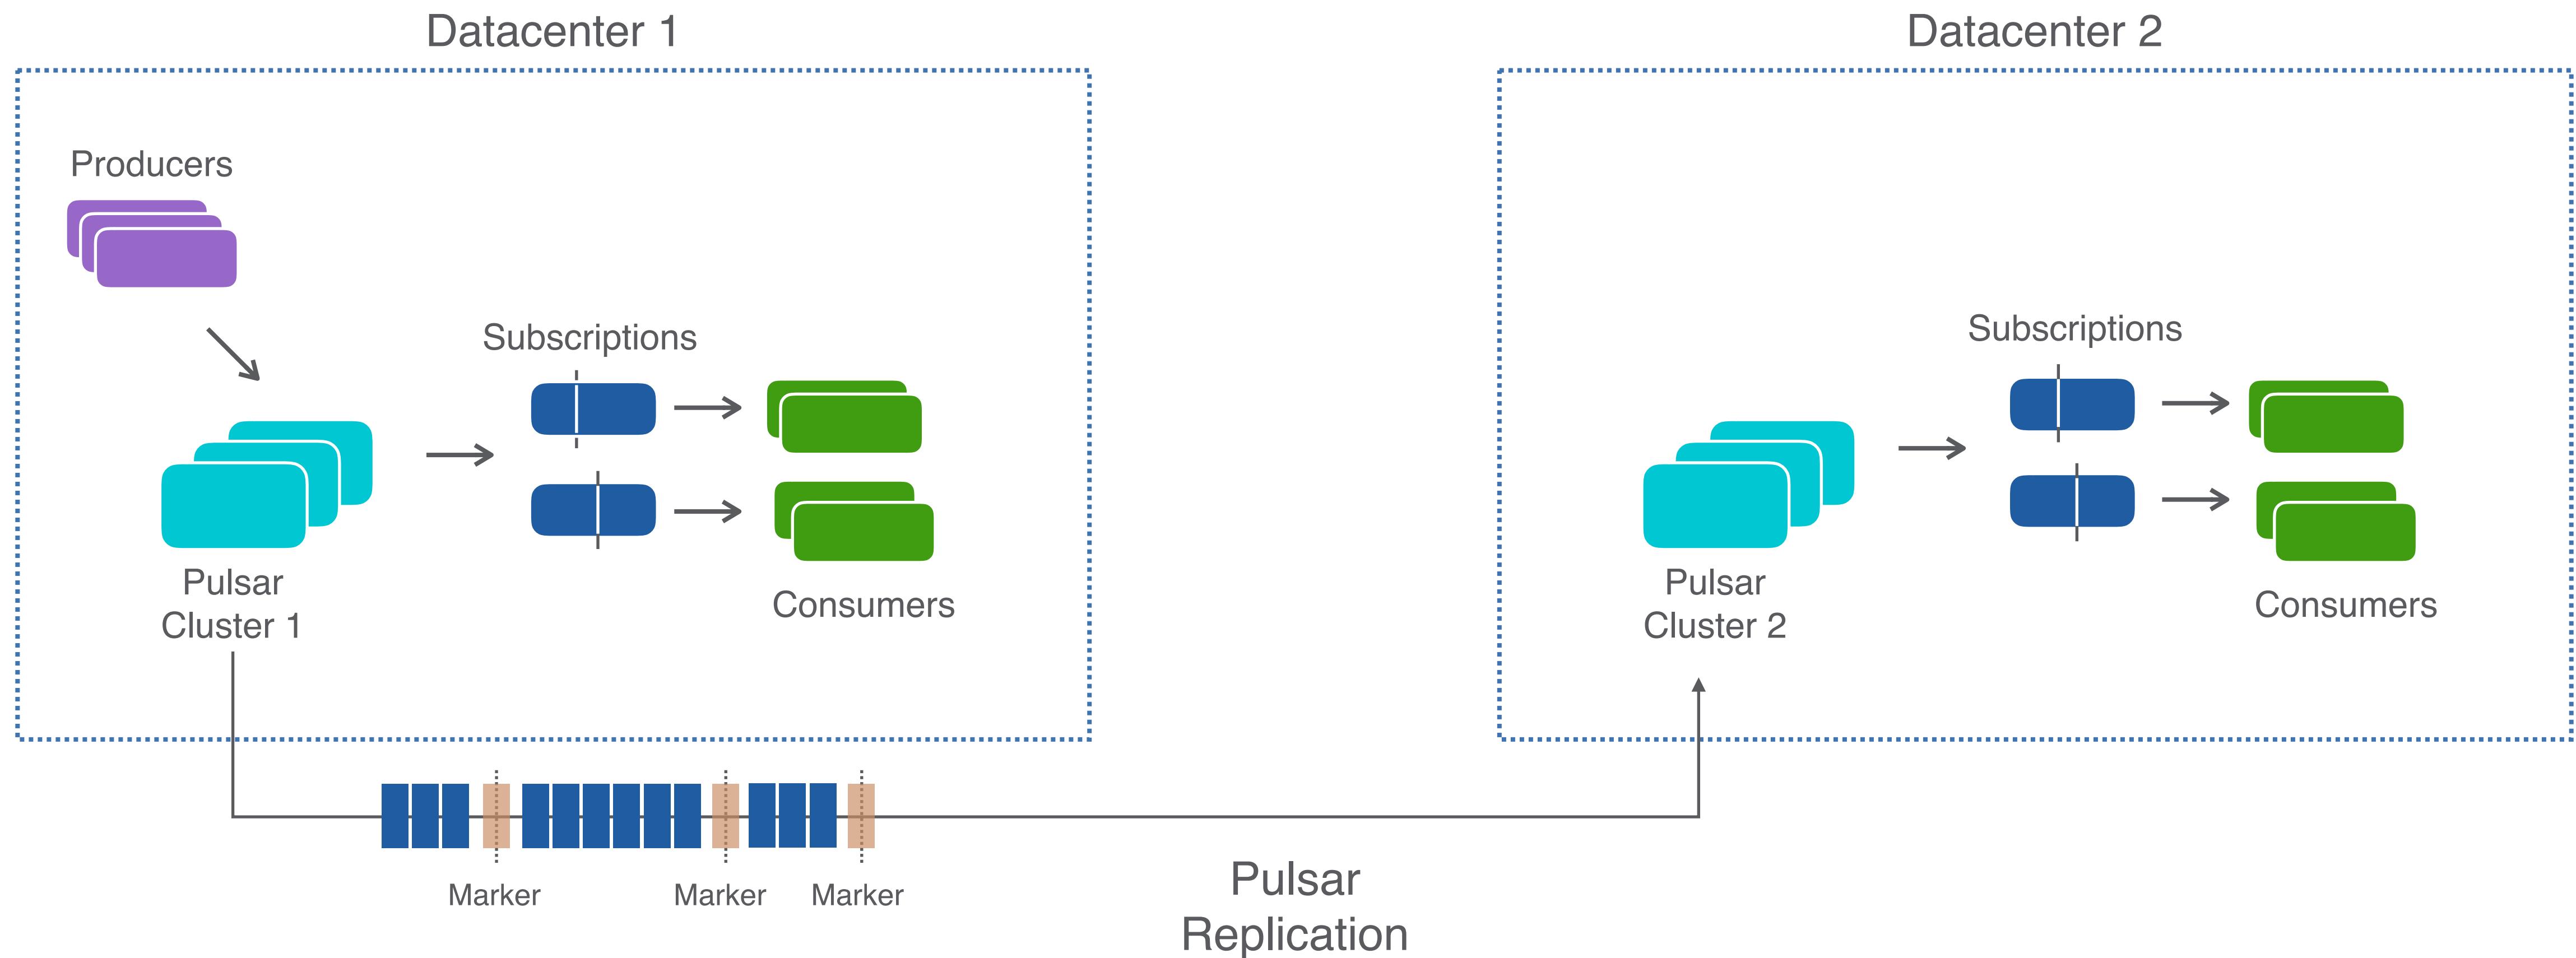
- Two independent clusters, primary and standby
- Configured tenants and namespaces replicate to standby
- Data published to primary is asynchronously replicated to standby
- Producers and consumers restarted in second datacenter upon primary failure

# Synchronous replication example

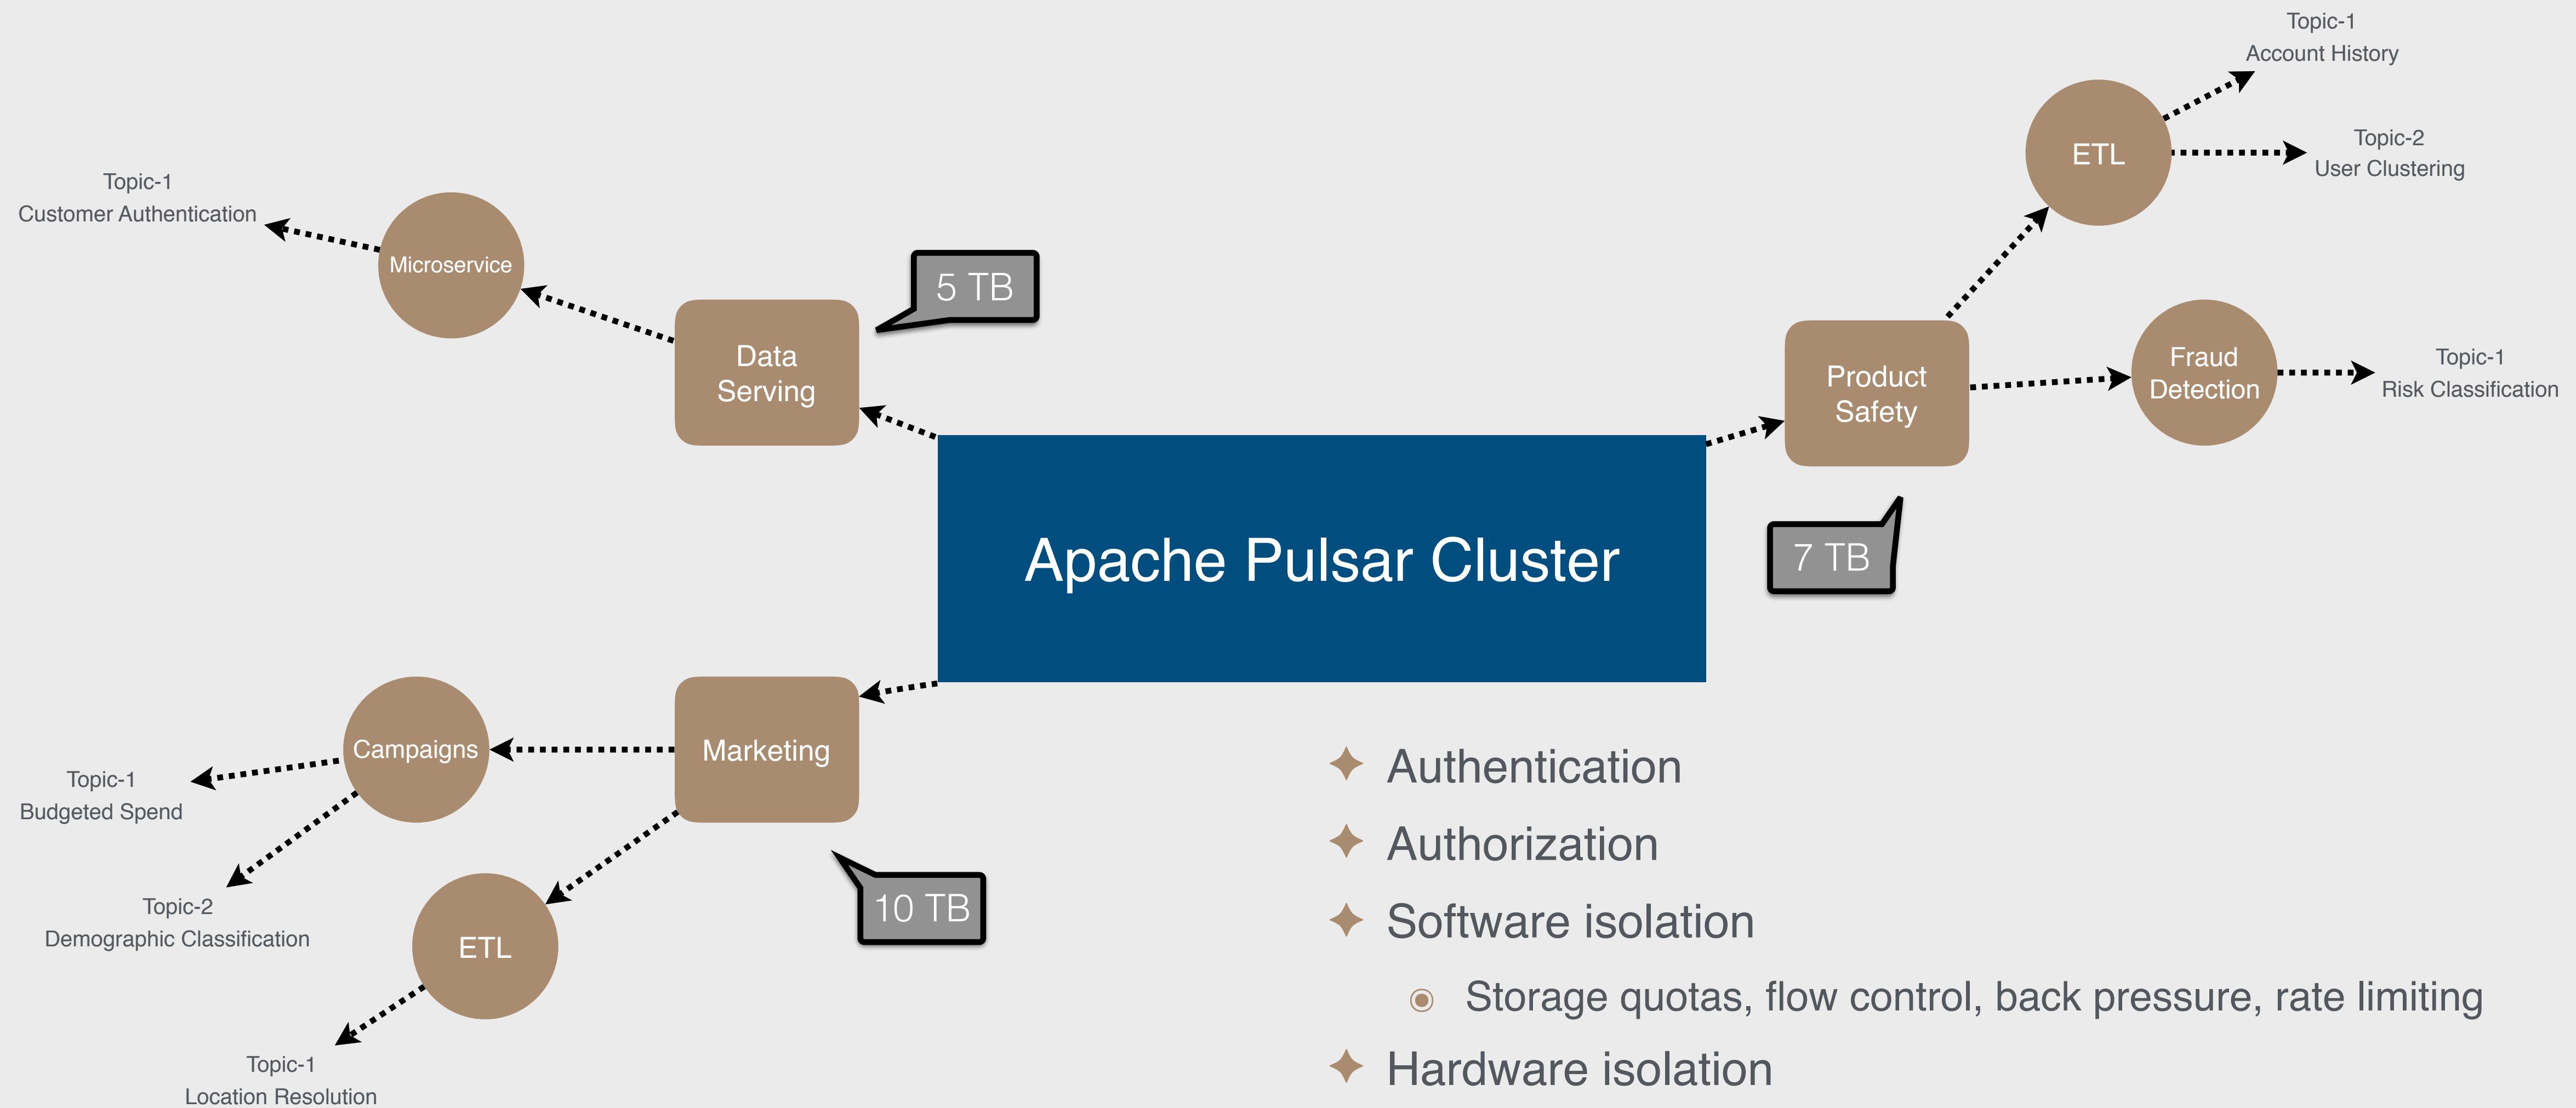


- Each topic owned by one broker at a time, i.e. in one datacenter
- ZooKeeper cluster spread across multiple locations
- Broker commits writes to bookies in both datacenters
- In event of datacenter failure, broker in surviving datacenter assumes ownership of topic

# Replicated subscriptions

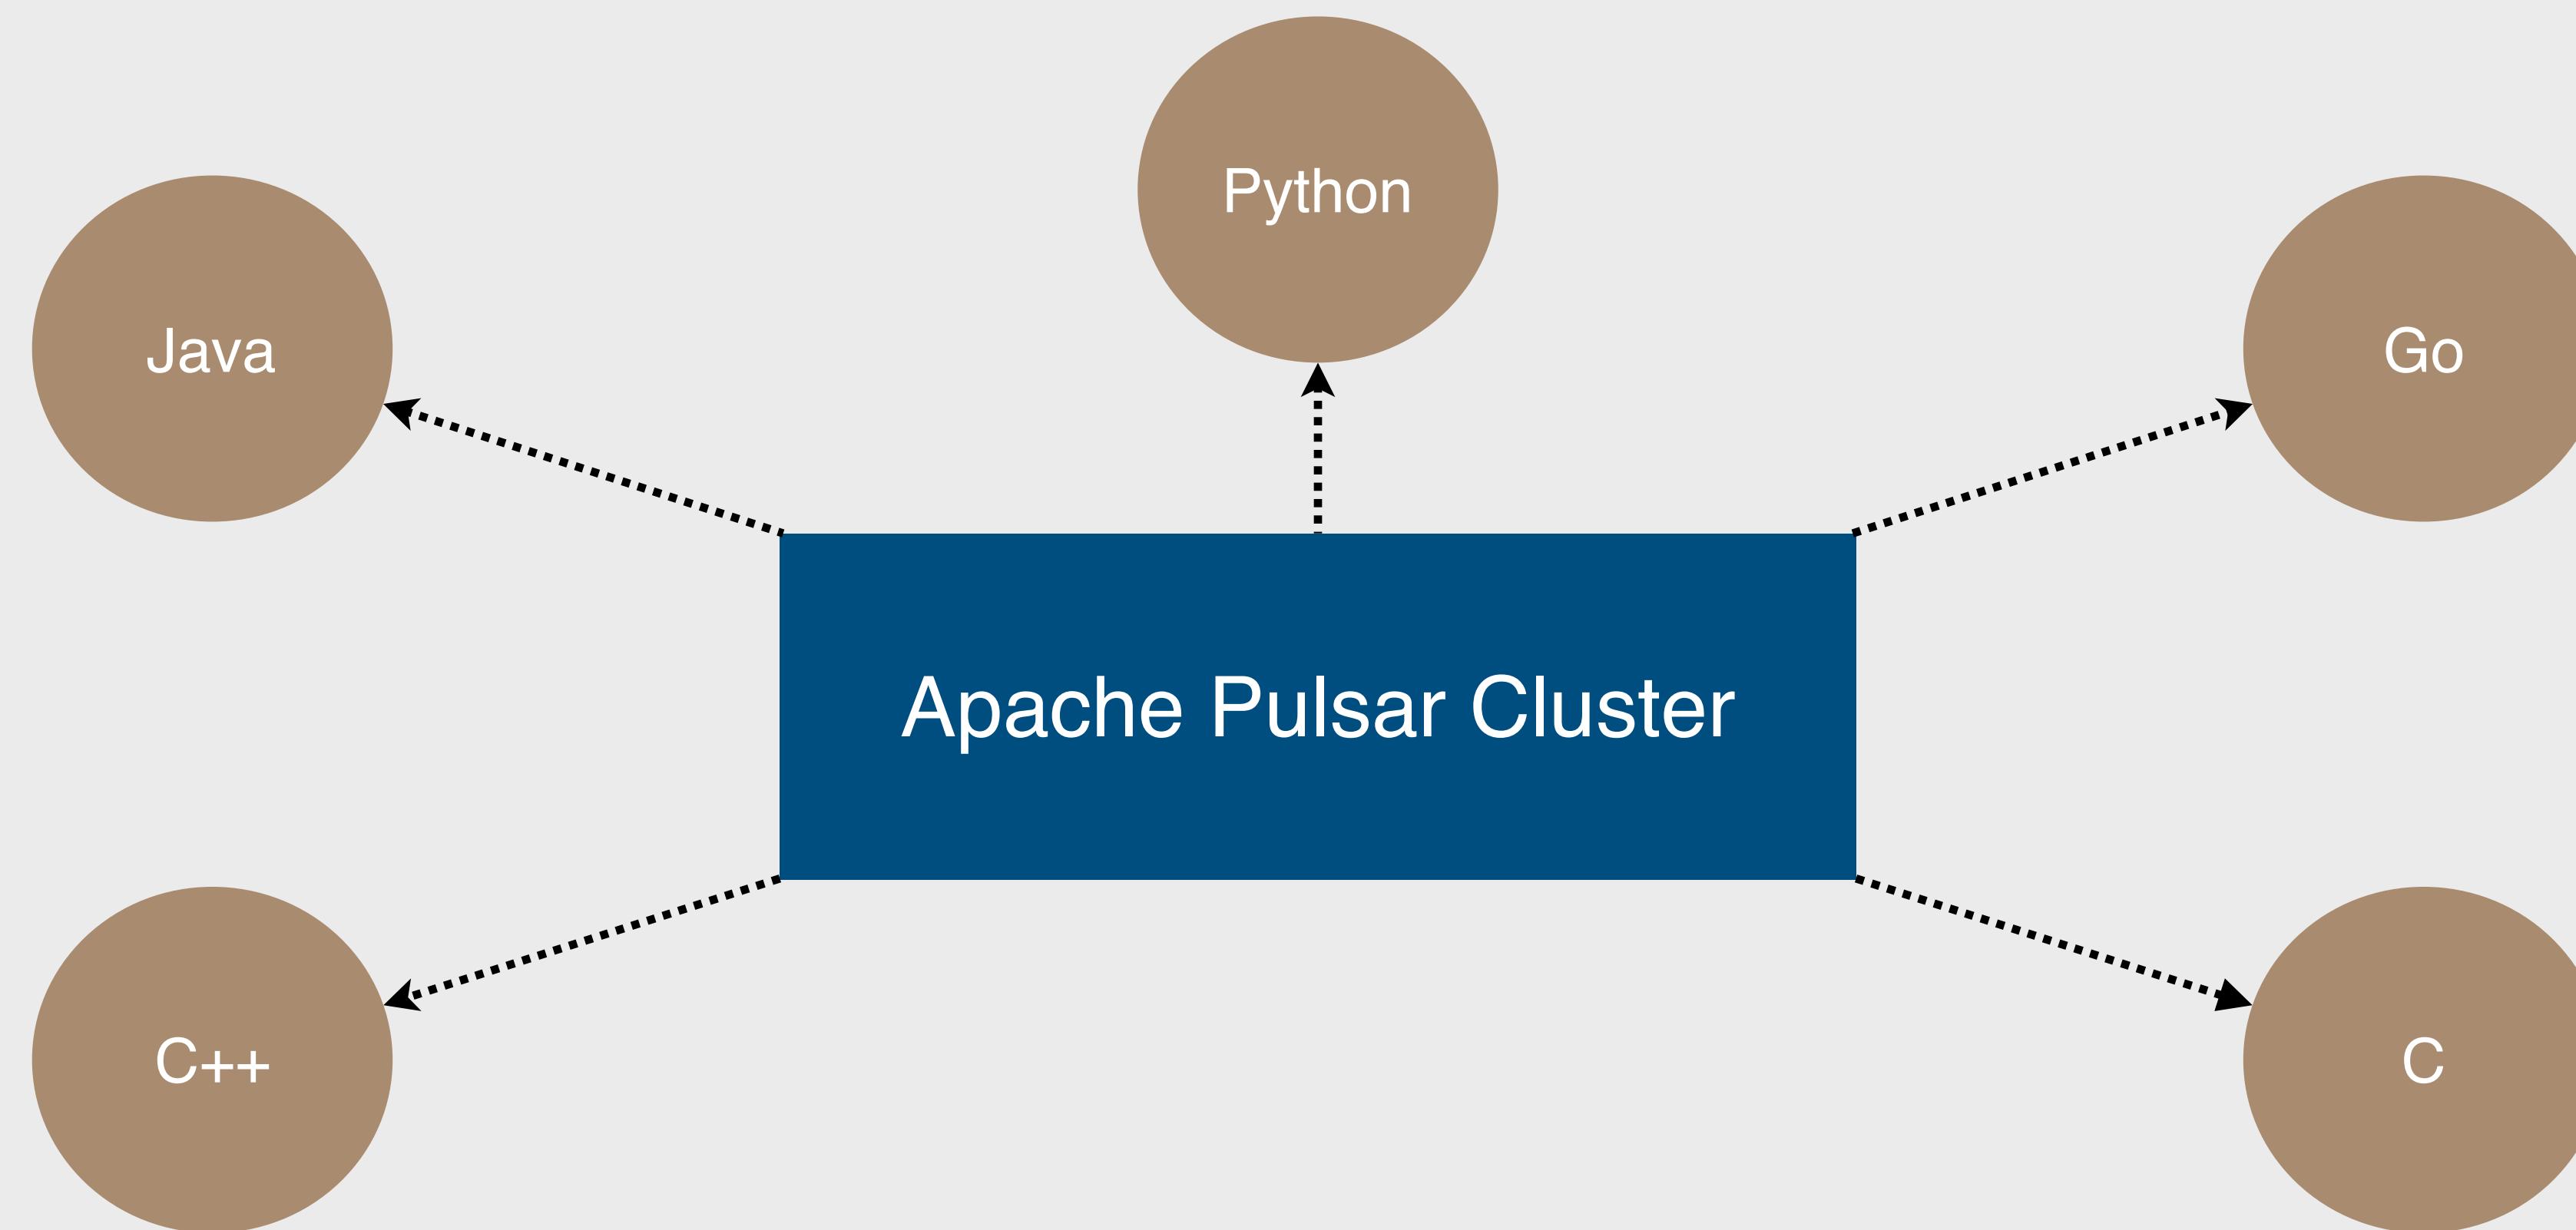


# MULTITENANCY - CLOUD NATIVE



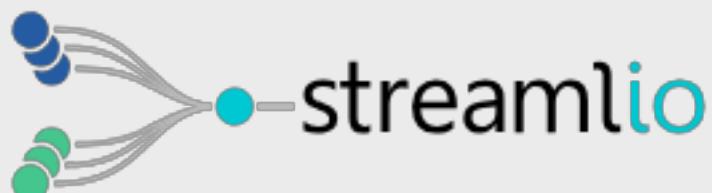
- ◆ Authentication
- ◆ Authorization
- ◆ Software isolation
  - Storage quotas, flow control, back pressure, rate limiting
- ◆ Hardware isolation
  - Constrain some tenants on a subset of brokers/bookies

# PULSAR CLIENTS



# PULSAR PRODUCER

```
PulsarClient client = PulsarClient.create(  
        "http://broker.usw.example.com:8080");  
  
Producer producer = client.createProducer(  
        "persistent://my-property/us-west/my-namespace/my-topic");  
  
// handles retries in case of failure  
producer.send("my-message".getBytes());  
  
// Async version:  
producer.sendAsync("my-message".getBytes()).thenRun(() -> {  
    // Message was persisted  
});
```



# PULSAR CONSUMER

```
PulsarClient client = PulsarClient.create(  
        "http://broker.usw.example.com:8080");  
  
Consumer consumer = client.subscribe(  
        "persistent://my-property/us-west/my-namespace/my-topic",  
        "my-subscription-name");  
  
while (true) {  
    // Wait for a message  
    Message msg = consumer.receive();  
  
    System.out.println("Received message: " + msg.getData());  
  
    // Acknowledge the message so that it can be deleted by broker  
    consumer.acknowledge(msg);  
}
```

# SCHEMA REGISTRY

- ◆ Provides type safety to applications built on top of Pulsar
- ◆ Two approaches
  - ◆ Client side - type safety enforcement up to the application
  - ◆ Server side - system enforces type safety and ensures that producers and consumers remain synced
- ◆ Schema registry enables clients to upload data schemas on a topic basis.
- ◆ Schemas dictate which data types are recognized as valid for that topic

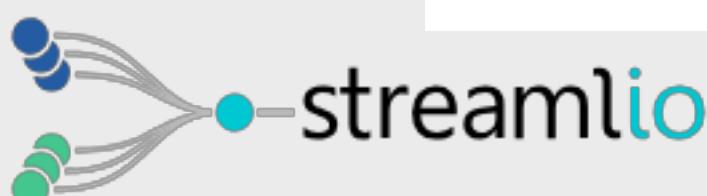
# PULSAR SCHEMAS - HOW DO THEY WORK?

- ◆ Enforced at the topic level
- ◆ Pulsar schemas consists of
  - ◆ Name - Name refers to the topic to which the schema is applied
  - ◆ Payload - Binary representation of the schema
  - ◆ Schema type - JSON, Protobuf and Avro
  - ◆ User defined properties - Map of strings to strings (application specific - e.g git hash of the schema)

# SCHEMA VERSIONING

```
PulsarClient client = PulsarClient.builder()  
    .serviceUrl("http://broker.usw.example.com:6650")  
    .build()  
  
Producer<SensorReading> producer = client.newProducer(JSONSchema.of(SensorReading.class))  
    .topic("sensor-data")  
    .sendTimeout(3, TimeUnit.SECONDS)  
    .create()
```

| Scenario   | What happens  |
|--|---|
| No schema exists for the topic   | Producer is created using the given schema                                |
| Schema already exists; producer connects using the same schema that's already stored | Schema is transmitted to the broker, determines that it is already stored |
| Schema already exists; producer connects using a new schema that is compatible       | Schema is transmitted, compatibility determined and stored as new schema  |



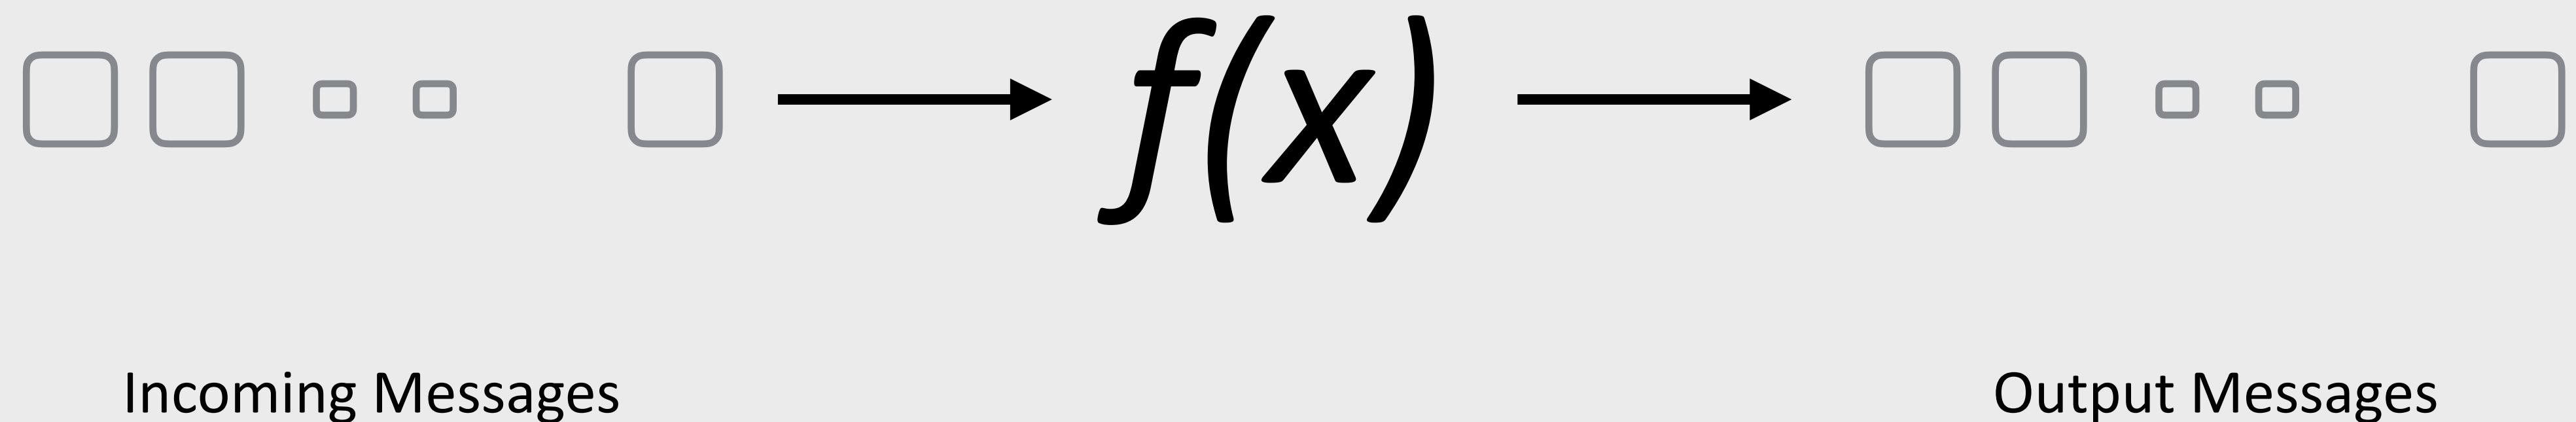
# Processing framework

# HOW TO PROCESS DATA MODELED AS STREAMS

- ◆ Consume data as it is produced (pub/sub)
- ◆ Light weight compute - transform and react to data as it arrives
- ◆ Heavy weight compute - continuous data processing
- ◆ Interactive query of stored streams

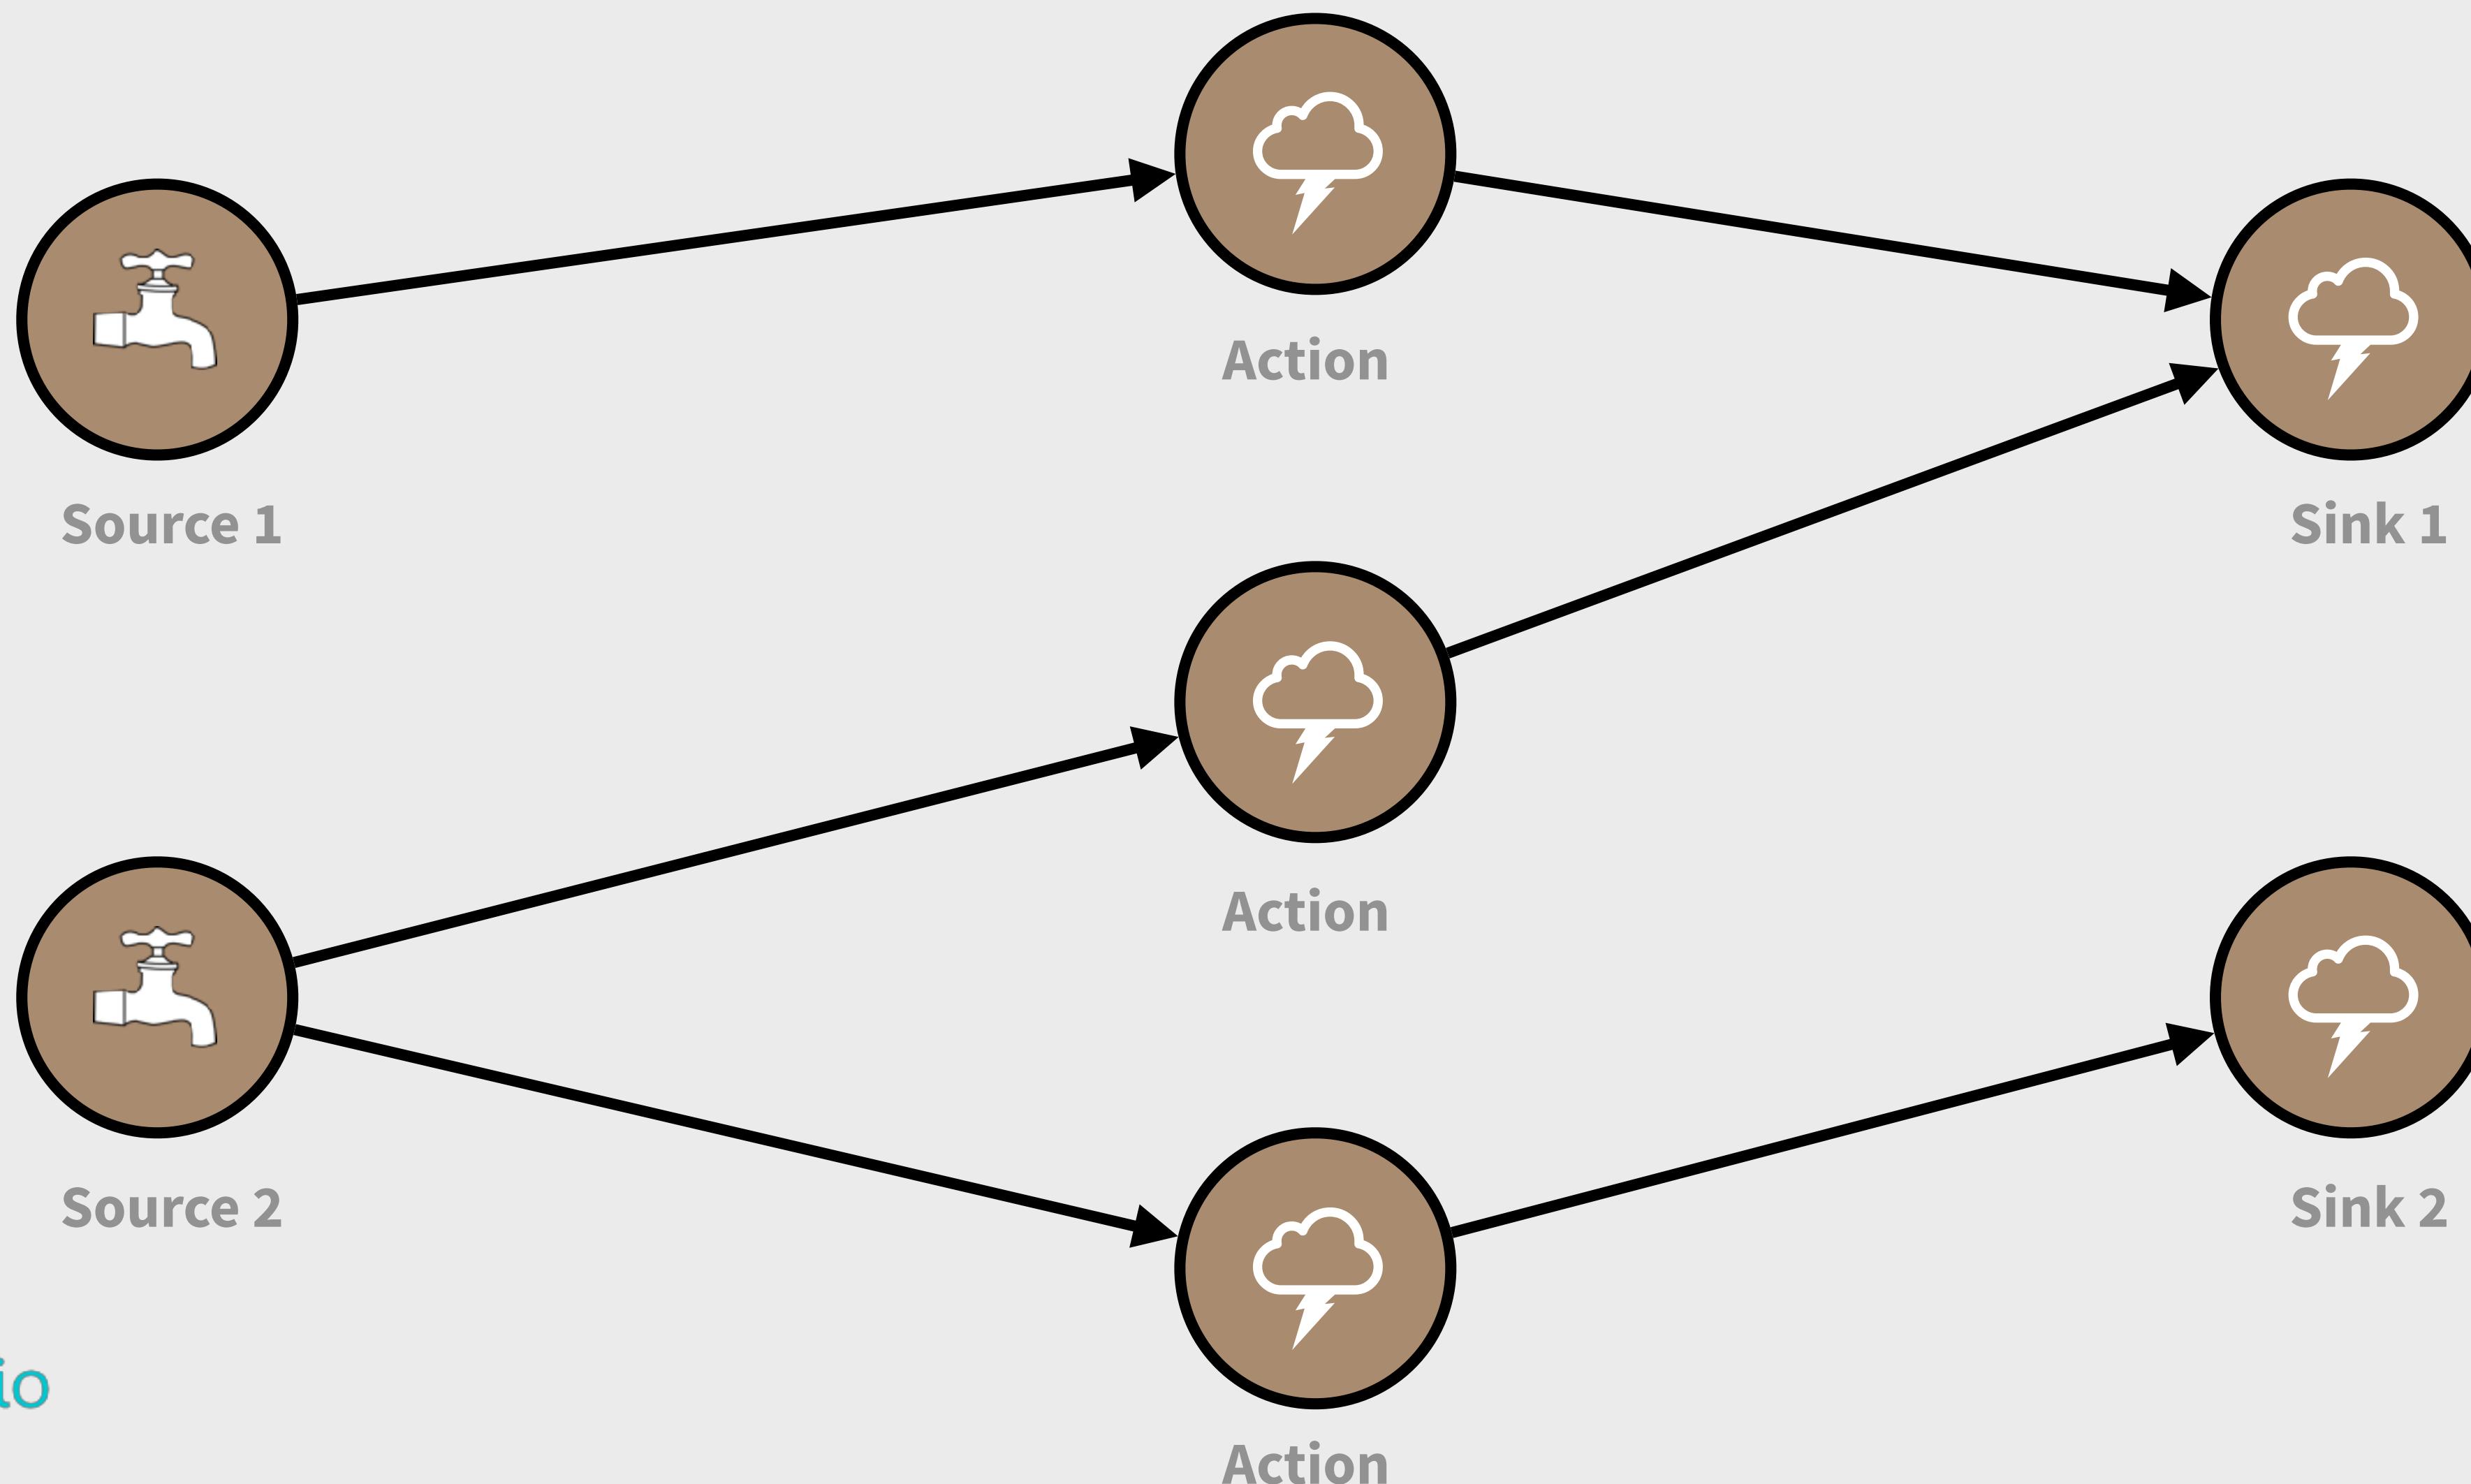
# LIGHT WEIGHT COMPUTE

ABSTRACT VIEW OF COMPUTE REPRESENTATION



# TRADITIONAL COMPUTE REPRESENTATION

DAG



# REALIZING COMPUTATION - EXPLICIT CODE

STITCHED BY PROGRAMMERS

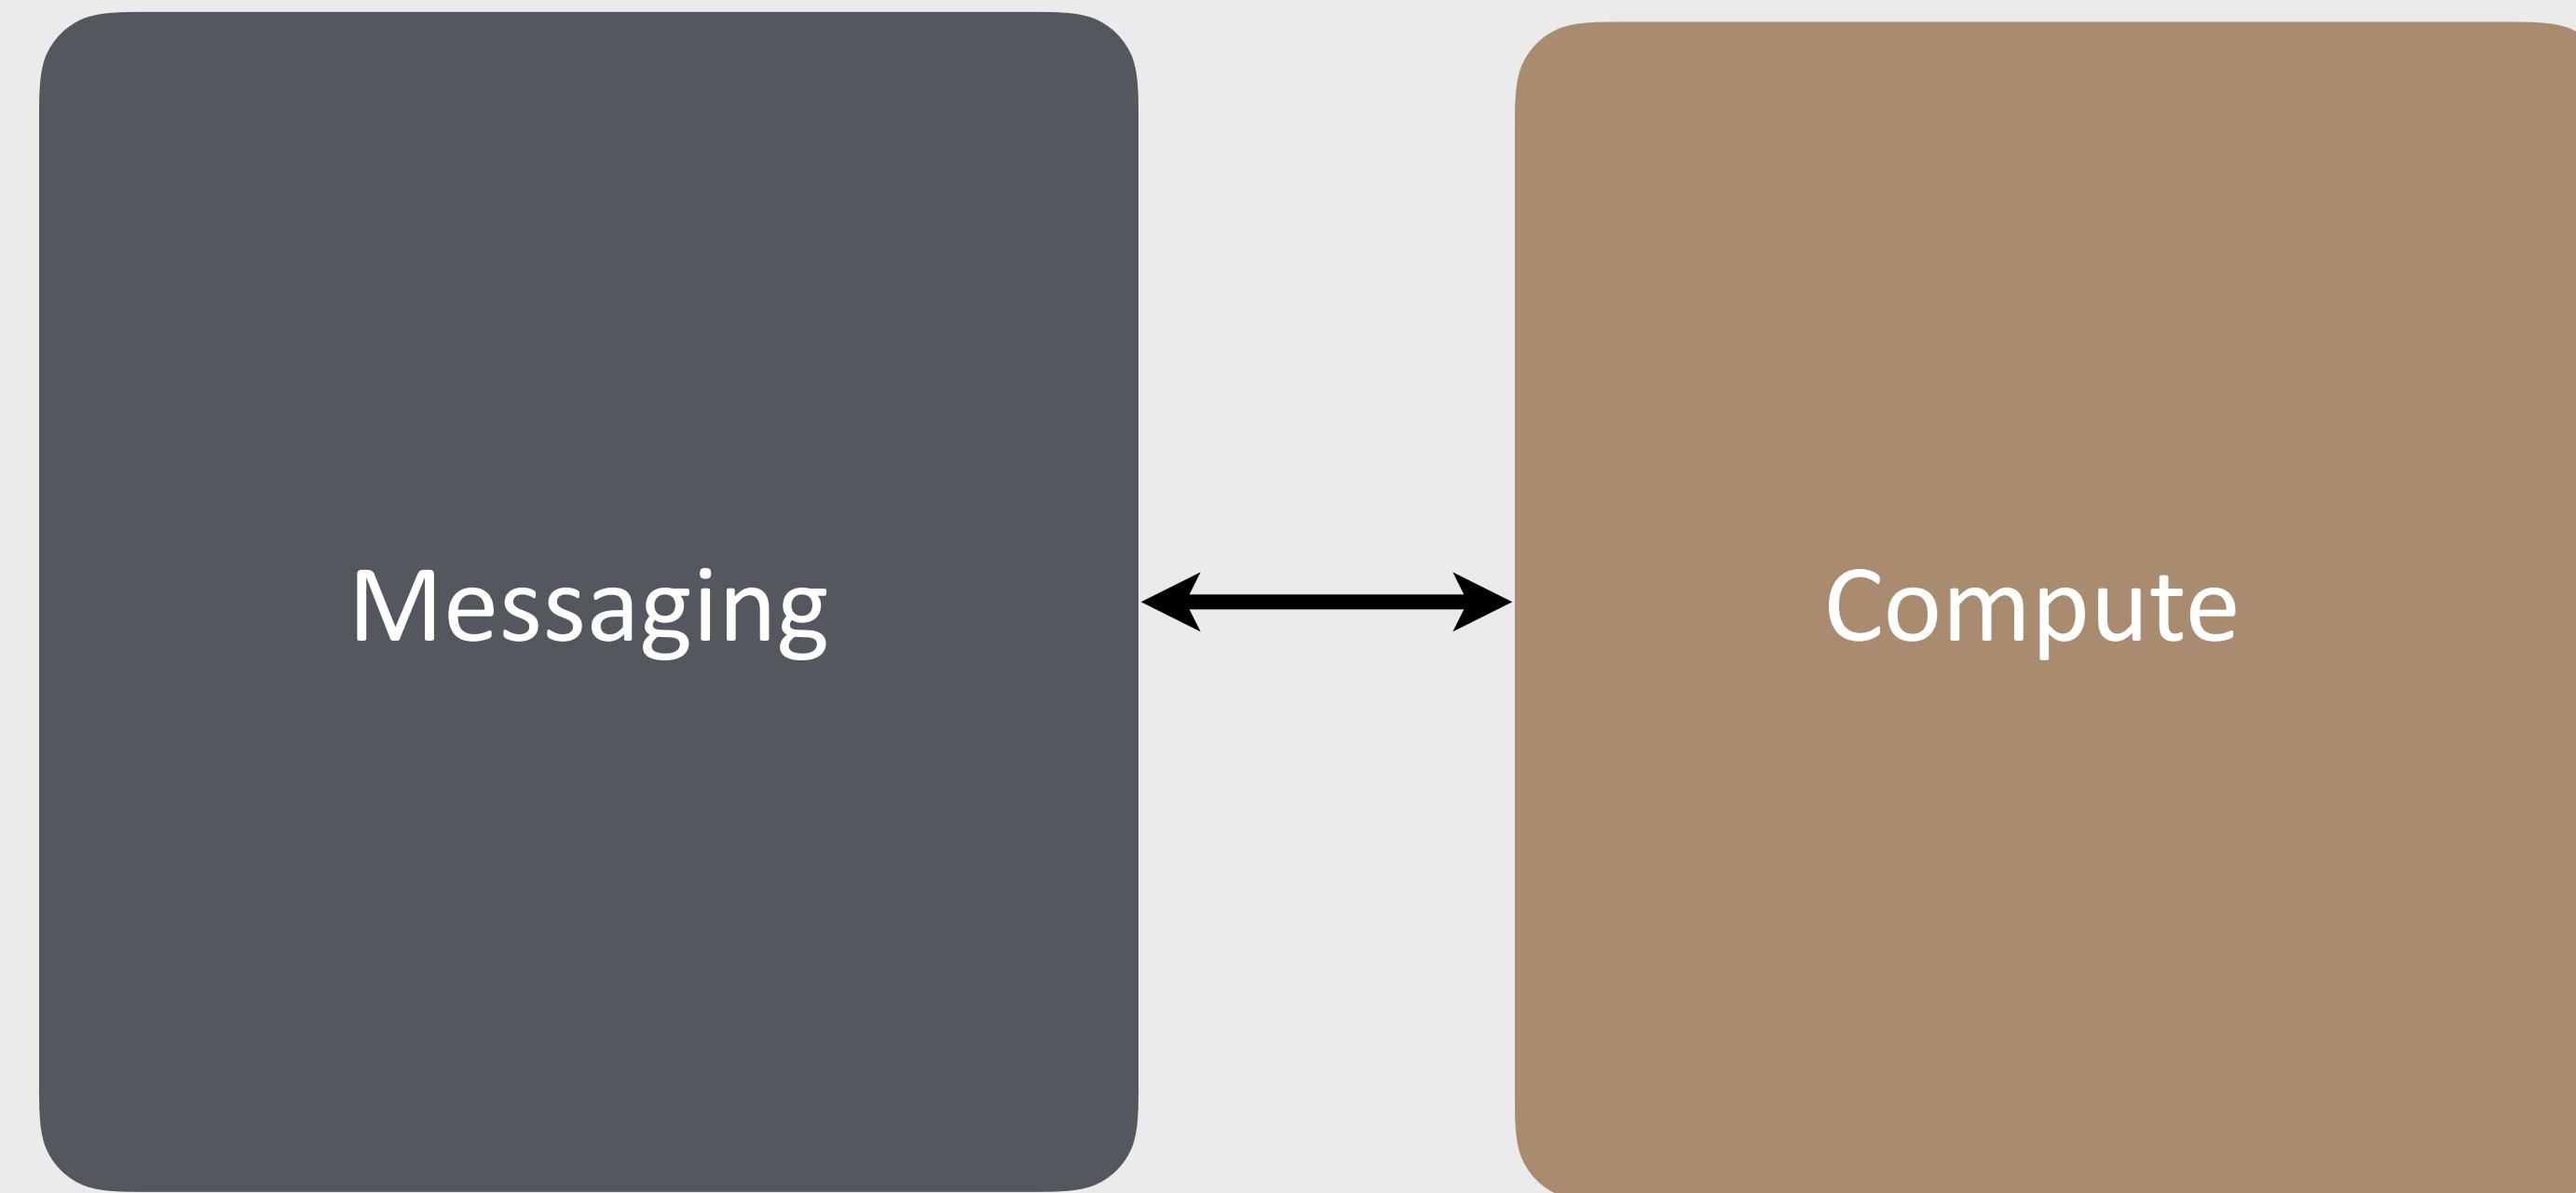
```
public static class SplitSentence extends BaseBasicBolt {  
    @Override  
    public void declareOutputFields(OutputFieldsDeclarer declarer) {  
        declarer.declare(new Fields("word"));  
    }  
  
    @Override  
    public Map<String, Object> getComponentConfiguration() {  
        return null;  
    }  
  
    public void execute(Tuple tuple, BasicOutputCollector  
basicOutputCollector) {  
        String sentence = tuple.getStringByField("sentence");  
        String words[] = sentence.split(" ");  
        for (String w : words) {  
            basicOutputCollector.emit(new Values(w));  
        }  
    }  
}
```



# REALIZING COMPUTATION - FUNCTIONAL

```
Builder.newBuilder()
    .newSource(() -> StreamletUtils.randomFromList(SENTENCES))
    .flatMap(sentence -> Arrays.asList(sentence.toLowerCase().split("|\s+")))
    .reduceByKeyAndWindow(word -> word, word -> 1,
        WindowConfig.TumblingCountWindow(50),
        (x, y) -> x + y);
```

# TRADITIONAL REAL TIME - SEPARATE SYSTEMS



# TRADITIONAL REAL TIME SYSTEMS

## DEVELOPER EXPERIENCE

- ◆ Powerful API but complicated
- ◆ Does everyone really need to learn functional programming?
- ◆ Configurable and scalable but management overhead
- ◆ Edge systems have resource and management constraints

# TRADITIONAL REAL TIME SYSTEMS

## OPERATIONAL EXPERIENCE

- ◆ Multiple systems to operate
  - ◆ IoT deployments routinely have thousands of edge systems
- ◆ Semantic differences
  - ◆ Mismatch and duplication between systems
  - ◆ Creates developer and operator friction

# LESSONS LEARNT - USE CASES

- ◆ Data transformations
- ◆ Data classification
- ◆ Data enrichment
- ◆ Data routing
- ◆ Data extraction and loading
- ◆ Real time aggregation
- ◆ Microservices



Significant set of processing tasks are exceedingly simple

# EMERGENCE OF CLOUD - SERVERLESS

- ◆ Simple function API
- ◆ Functions are submitted to the system
- ◆ Runs per events
- ◆ Composition APIs to do complex things
- ◆ Wildly popular

# SERVERLESS VS STREAMING

- ◆ Both are event driven architectures
- ◆ Both can be used for analytics and data serving
- ◆ Both have composition APIs
  - Configuration based for serverless
  - DSL based for streaming
- ◆ Serverless typically does not guarantee ordering
- ◆ Serverless is pay per action

# STREAM NATIVE COMPUTE USING FUNCTIONS

APPLYING INSIGHT GAINED FROM SERVERLESS

- ◆ Simplest possible API -function or a procedure
- ◆ Support for multi language
- ◆ Use of native API for each language
- ◆ Scale developers
- ◆ Use of message bus native concepts - input and output as topics
- ◆ Flexible runtime - simple standalone applications vs managed system applications



# PULSAR FUNCTIONS

SDK LESS API

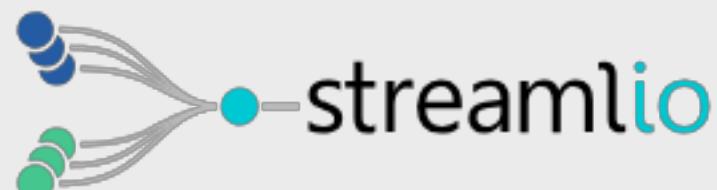
```
import java.util.function.Function;
public class ExclamationFunction implements Function<String, String> {
    @Override
    public String apply(String input) {
        return input + "!";
    }
}
```



# PULSAR FUNCTIONS

SDK API

```
import org.apache.pulsar.functions.api.PulsarFunction;
import org.apache.pulsar.functions.api.Context;
public class ExclamationFunction implements PulsarFunction<String, String> {
    @Override
    public String process(String input, Context context) {
        return input + "!";
    }
}
```



# PULSAR FUNCTIONS

- ◆ Function executed for every message of input topic
- ◆ Support for multiple topics as inputs
- ◆ Function output goes into output topic - can be void topic as well
- ◆ SerDe takes care of serialization/deserialization of messages
  - Custom SerDe can be provided by the users
  - Integration with schema registry

# PROCESSING GUARANTEES

## ◆ ATMOST\_ONCE

- Message acked to Pulsar as soon as we receive it

## ◆ ATLEAST\_ONCE

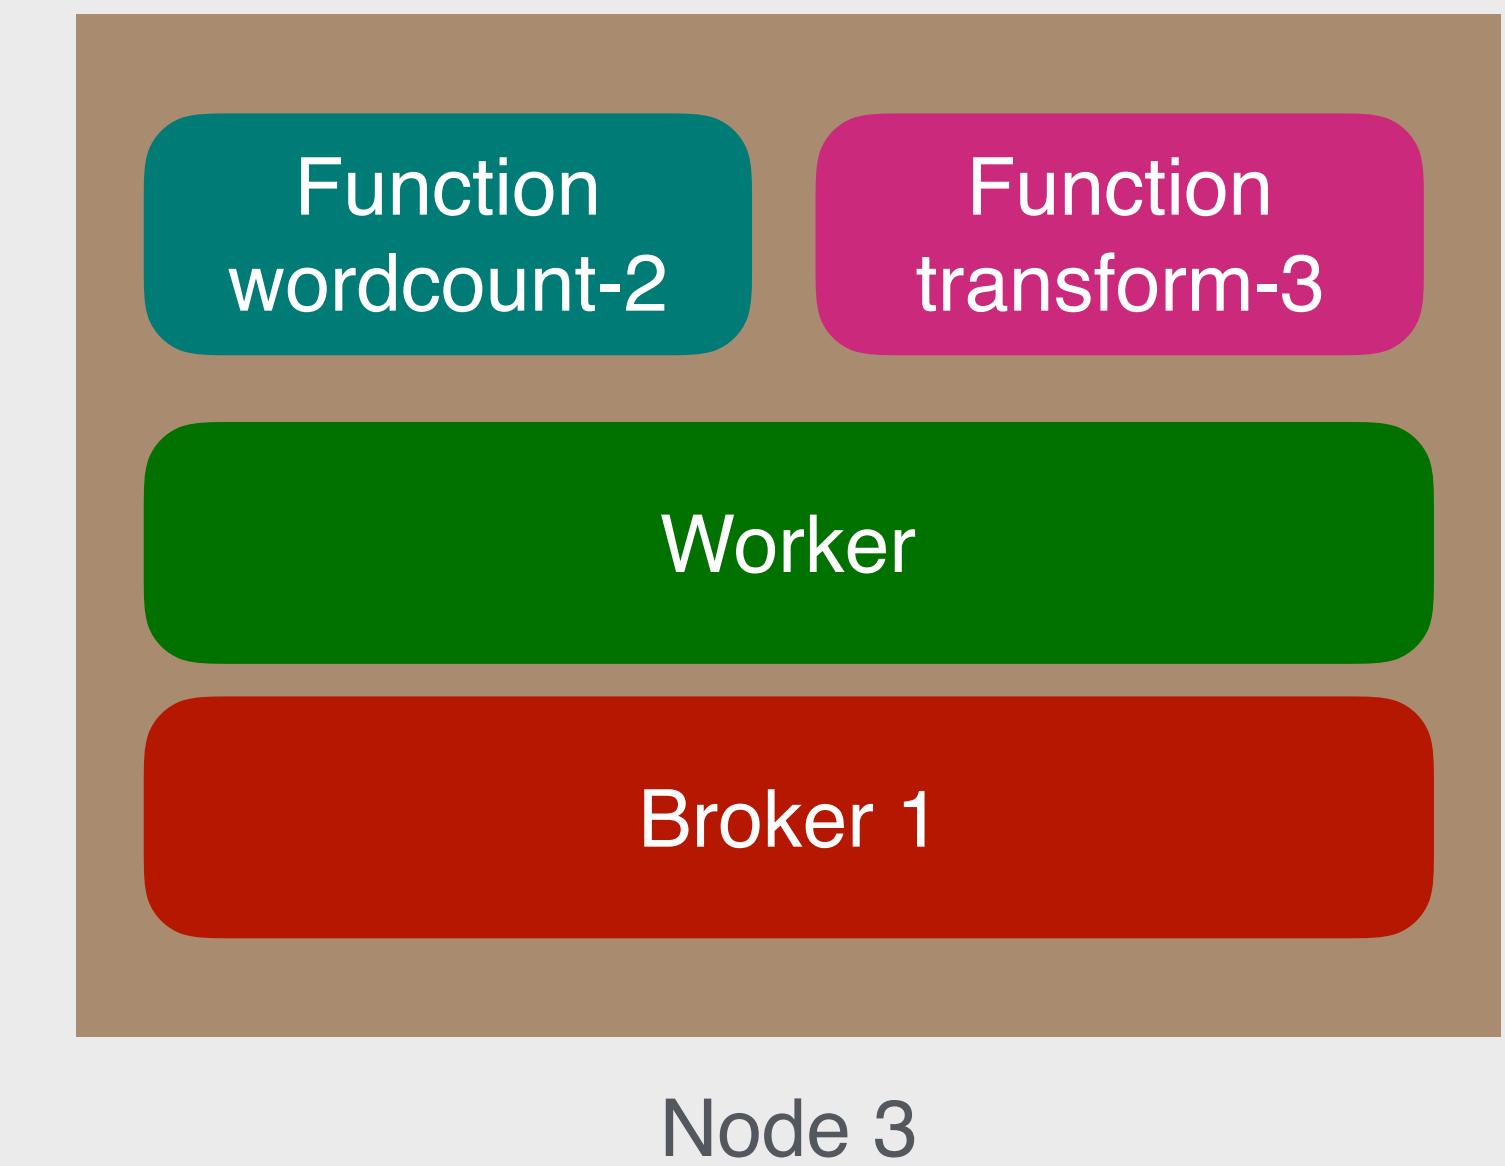
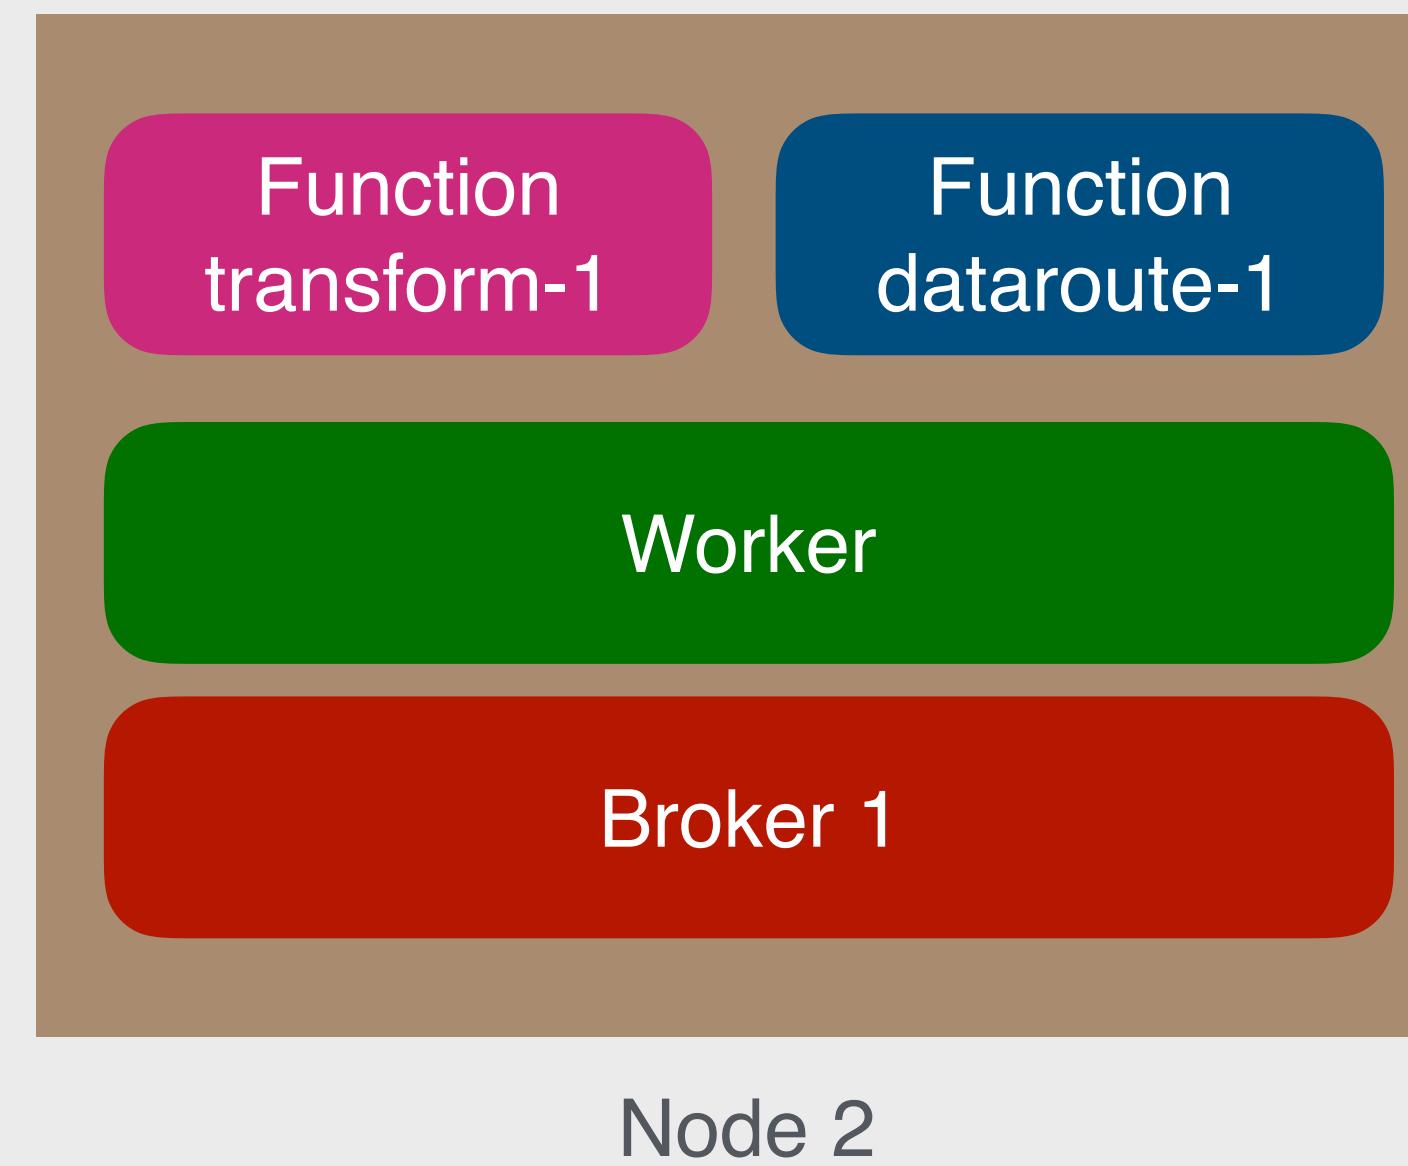
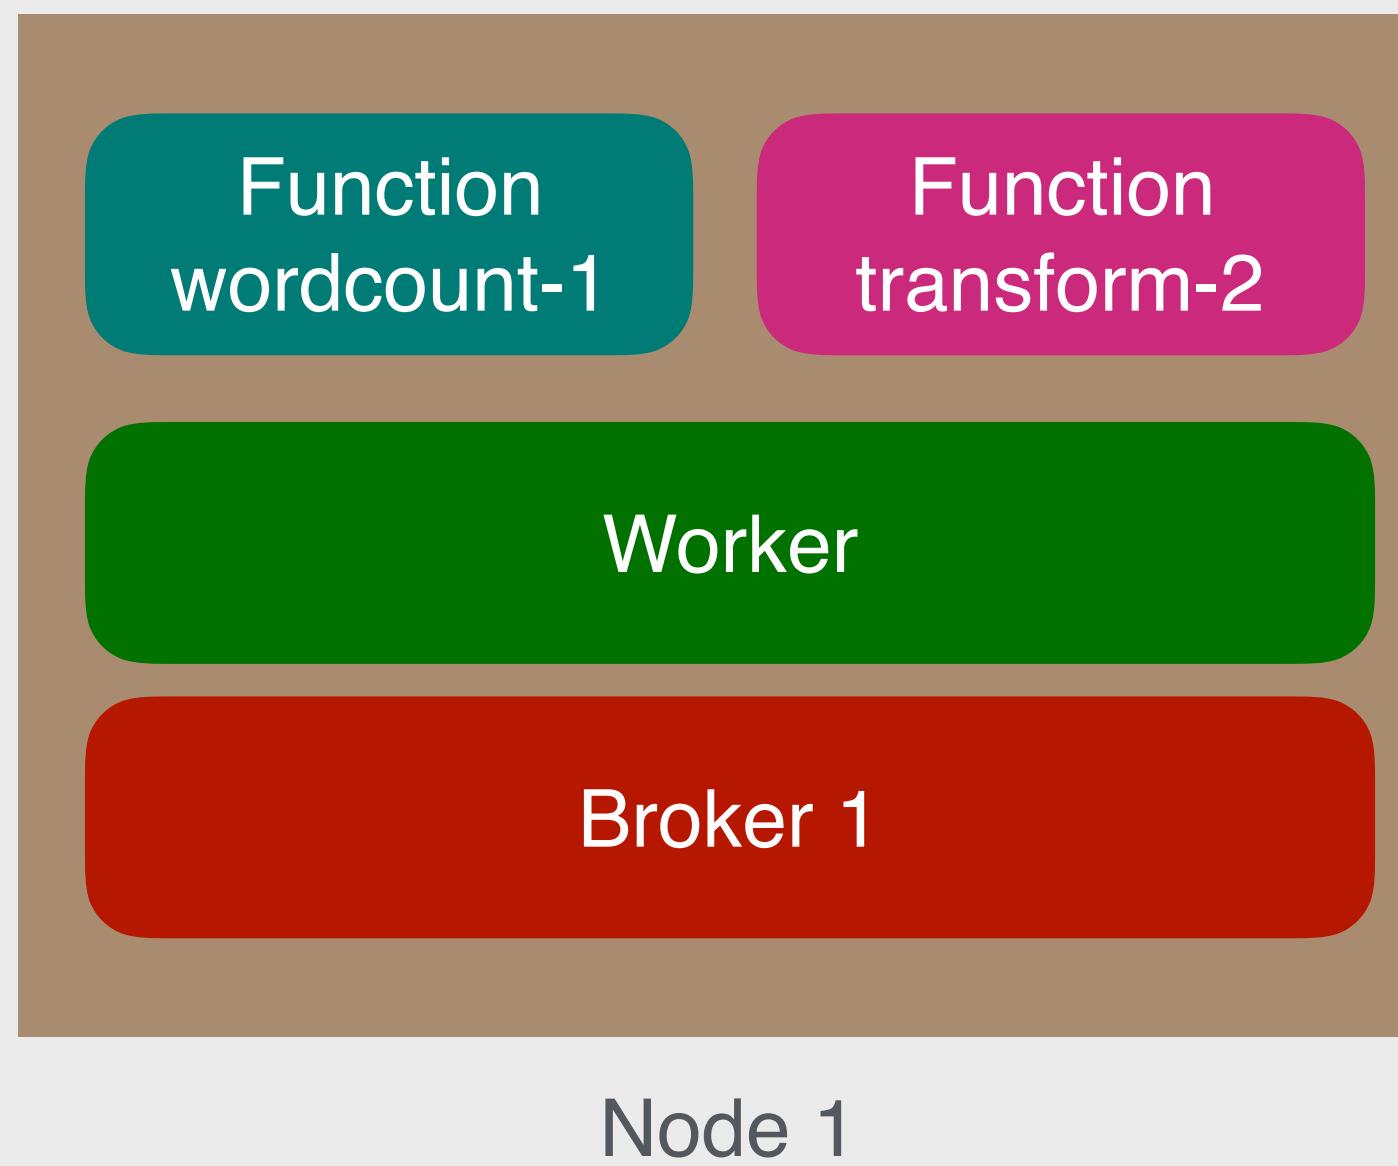
- Message acked to Pulsar after the function completes
- Default behavior - don't want people to loose data

## ◆ EFFECTIVELY\_ONCE

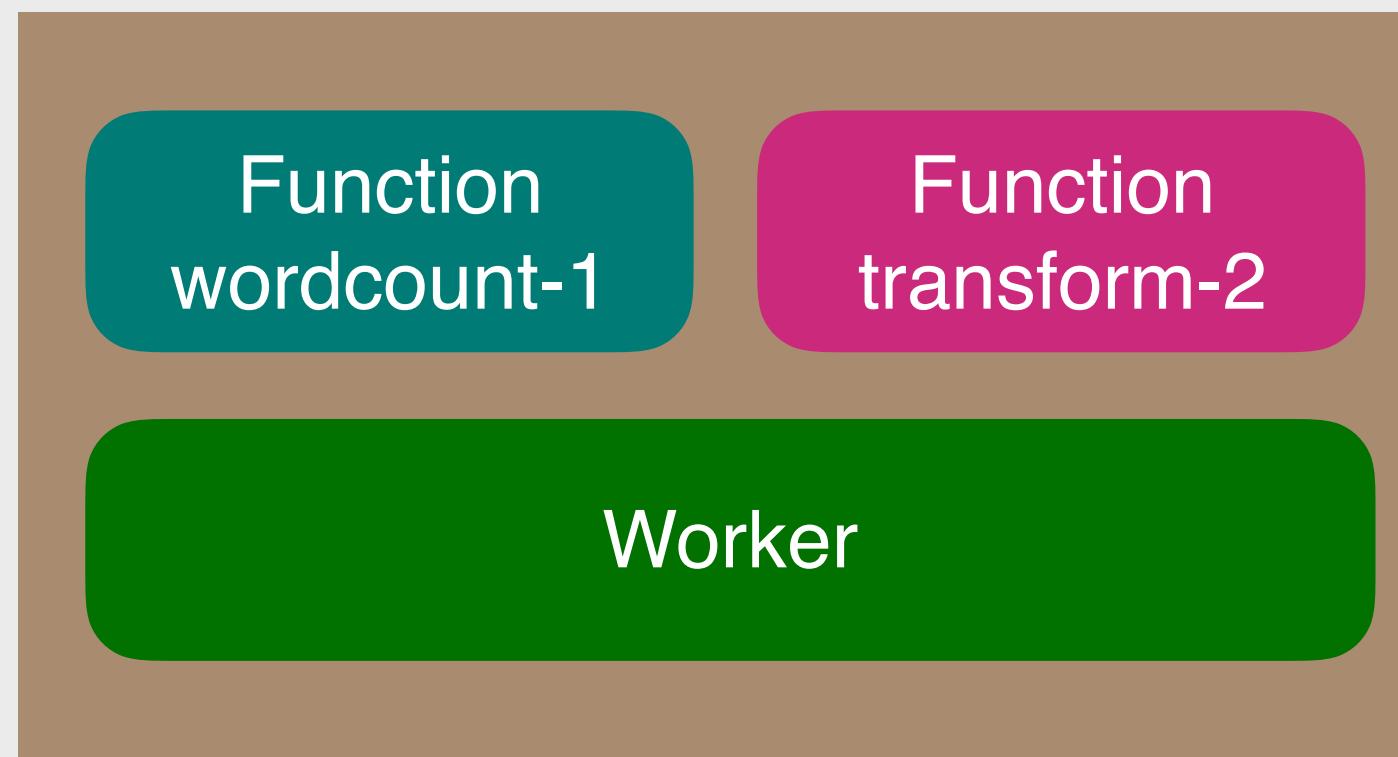
- Uses Pulsar's inbuilt effectively once semantics

## ◆ Controlled at runtime by user

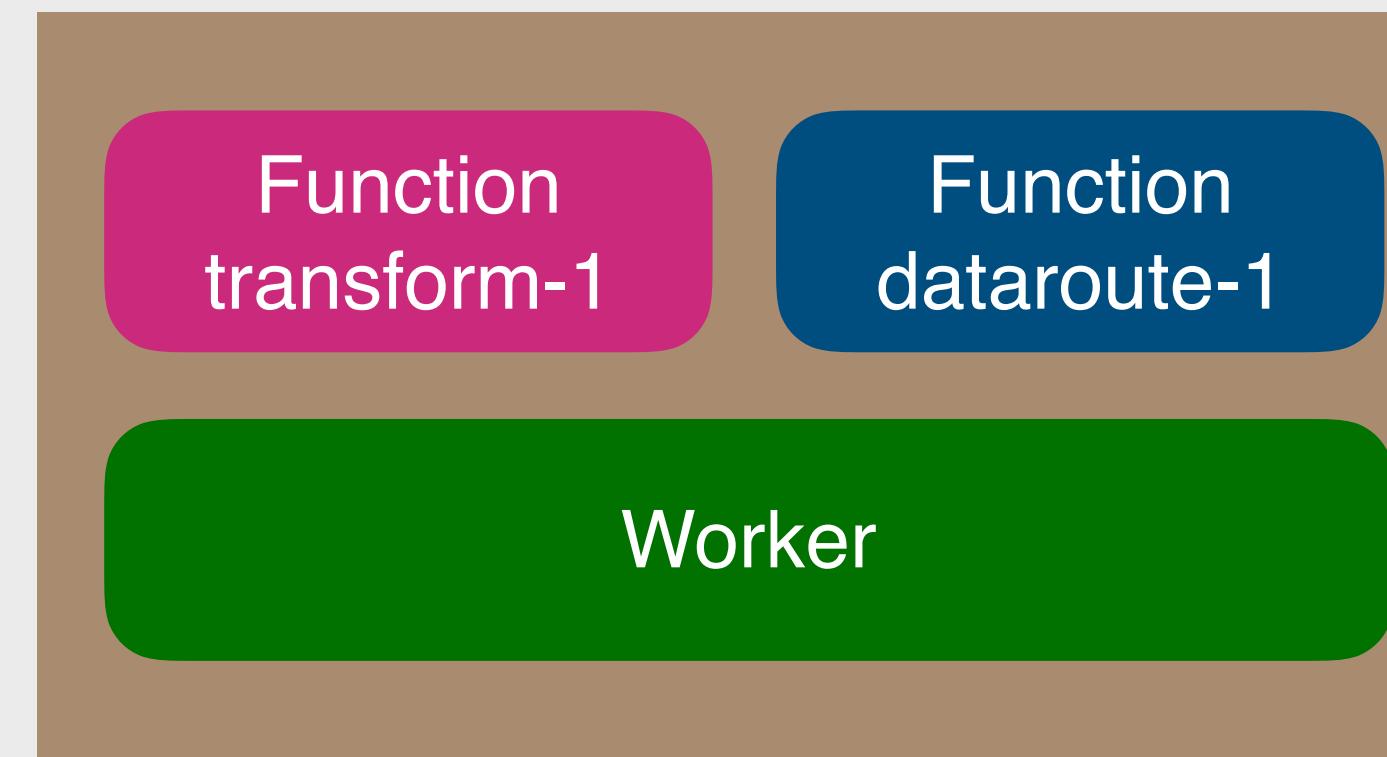
# DEPLOYING FUNCTIONS - BROKER



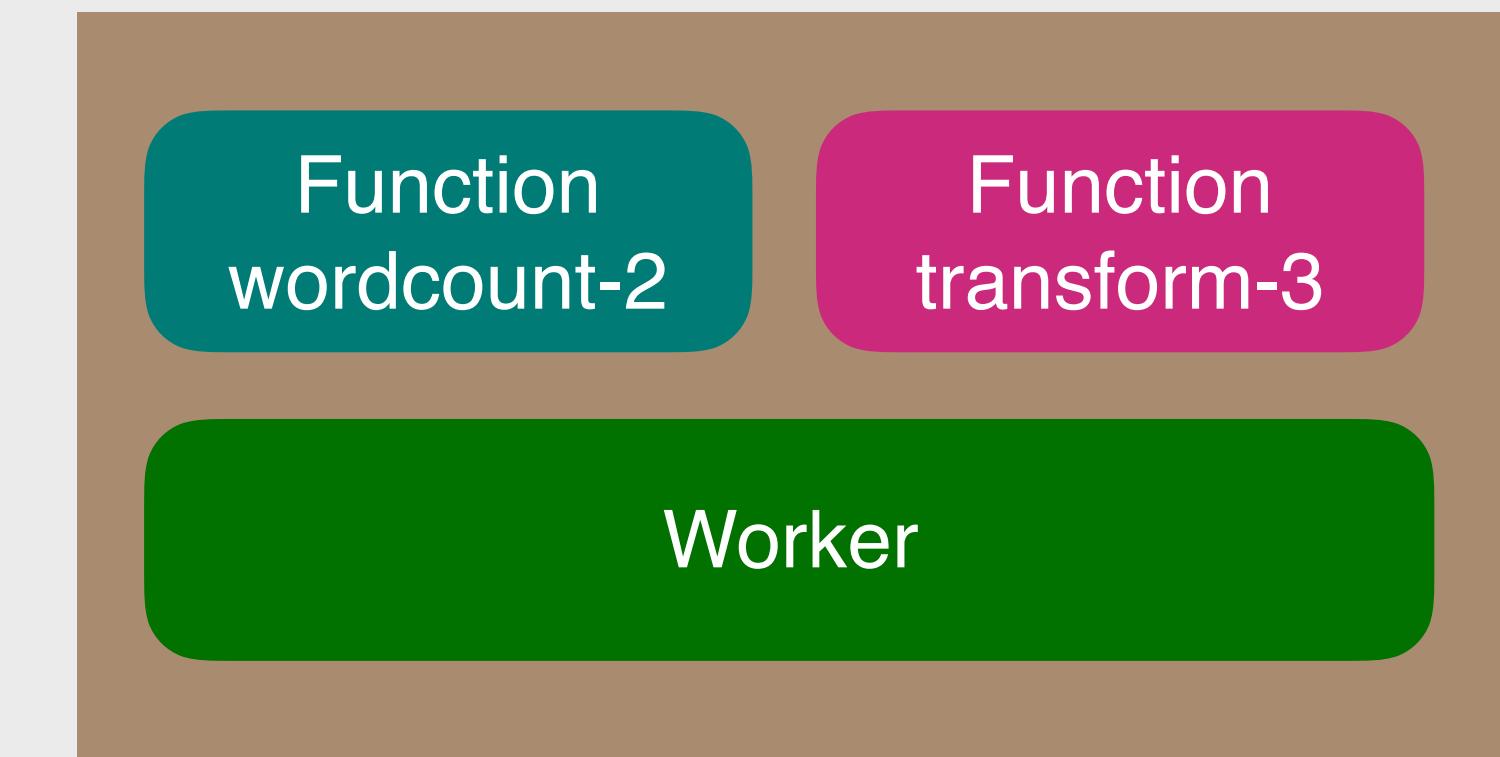
# DEPLOYING FUNCTIONS - WORKER NODES



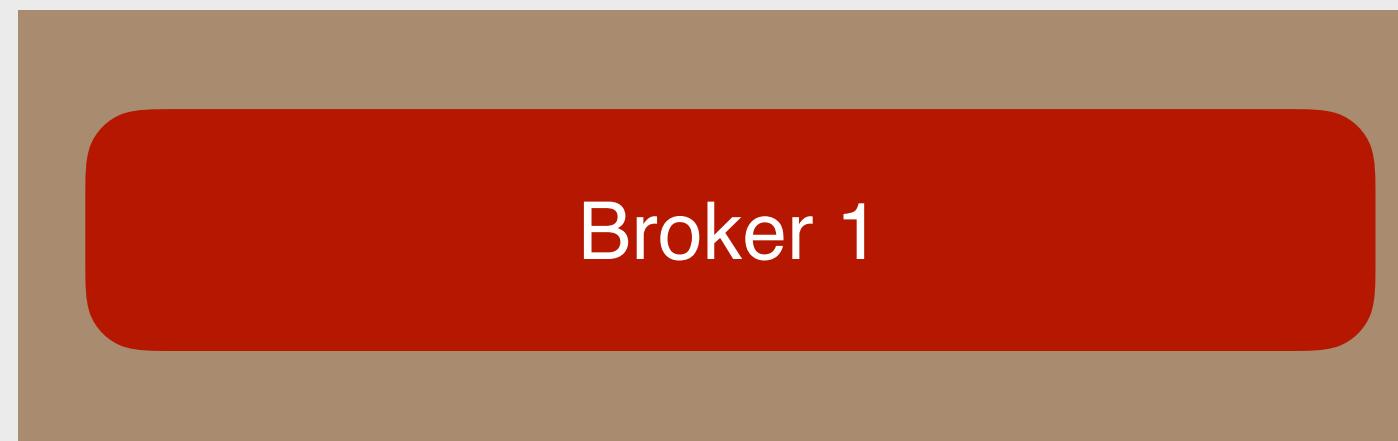
Node 1



Node 2



Node 3



Node 4

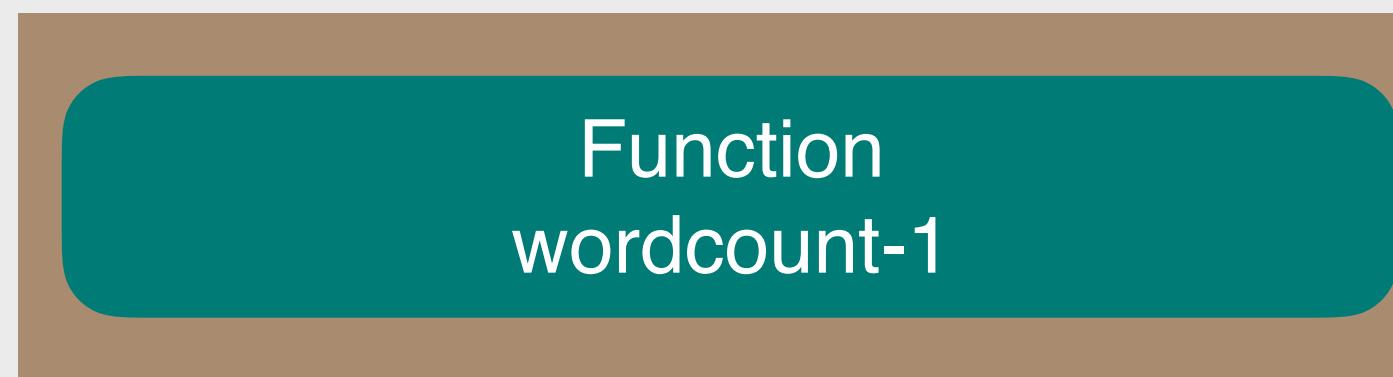


Node 5

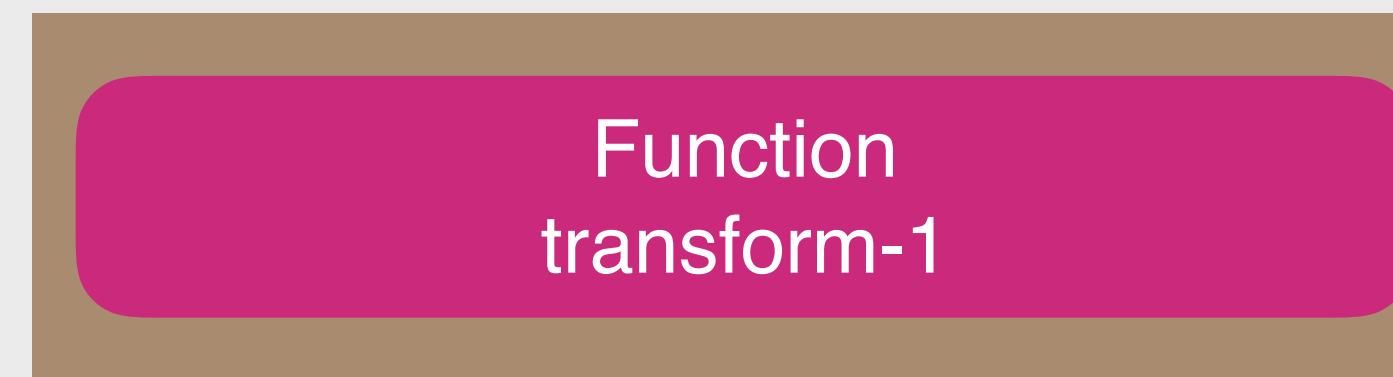


Node 6

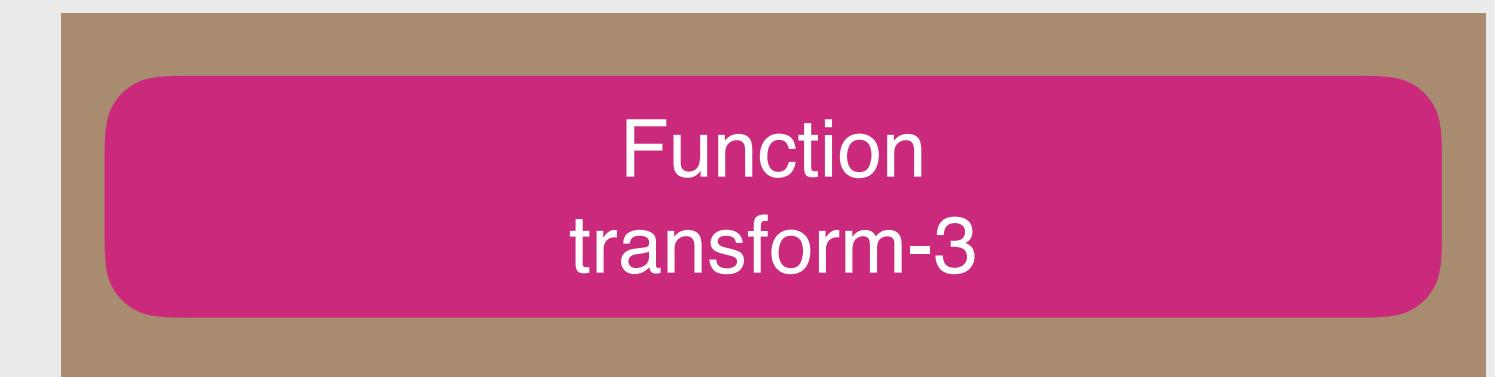
# DEPLOYING FUNCTIONS - KUBERNETES



Pod 1



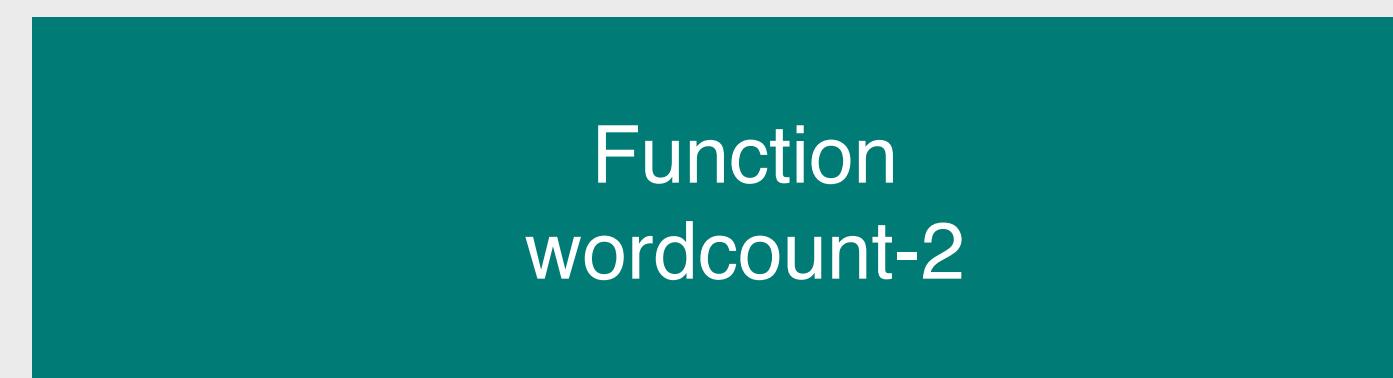
Pod 2



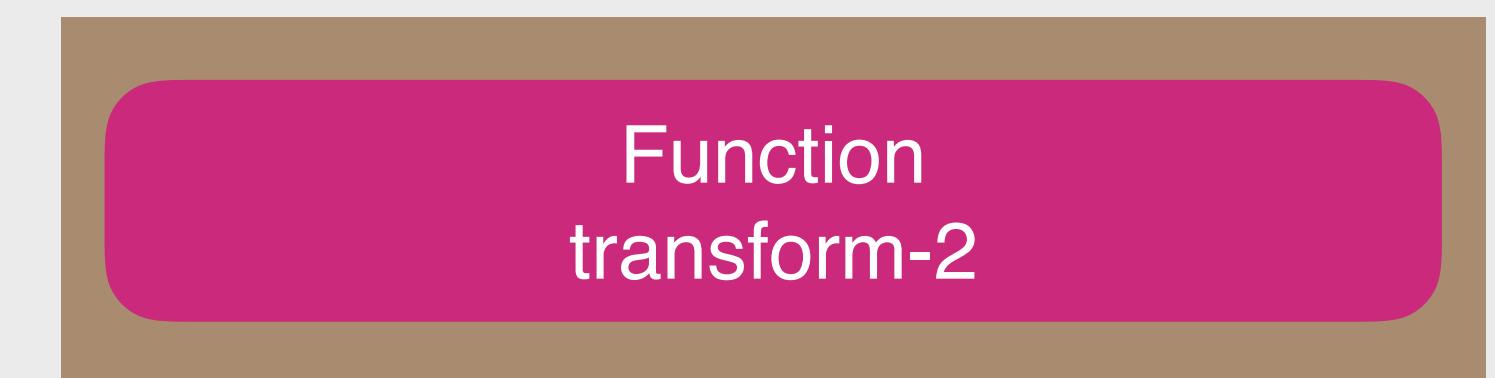
Pod 3



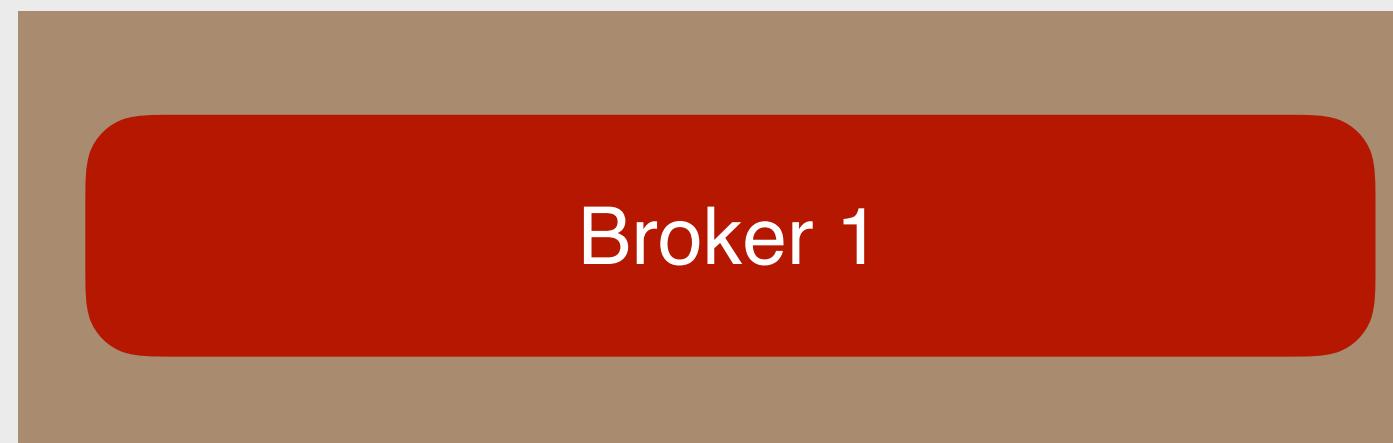
Pod 4



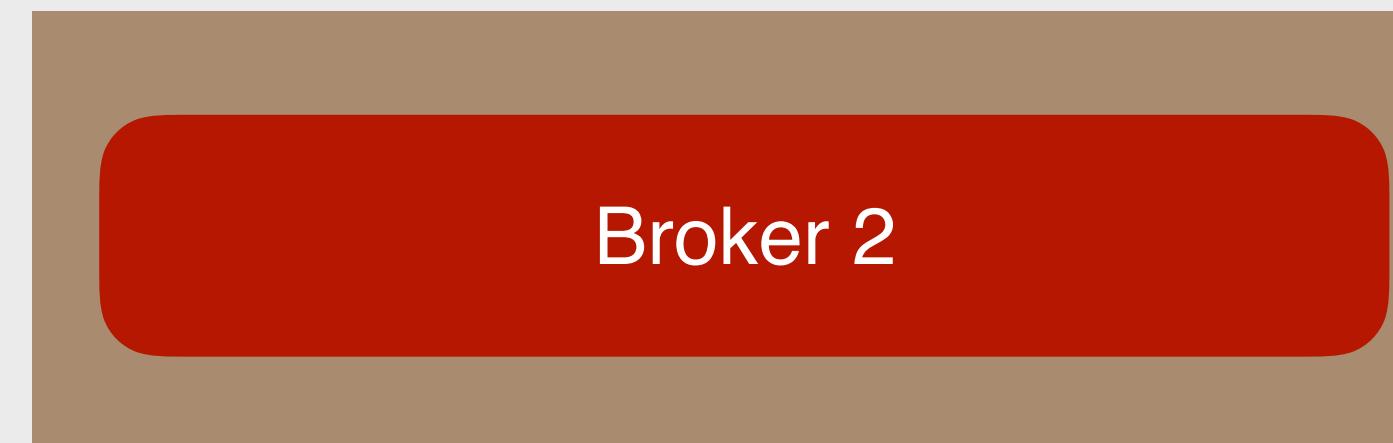
Pod 5



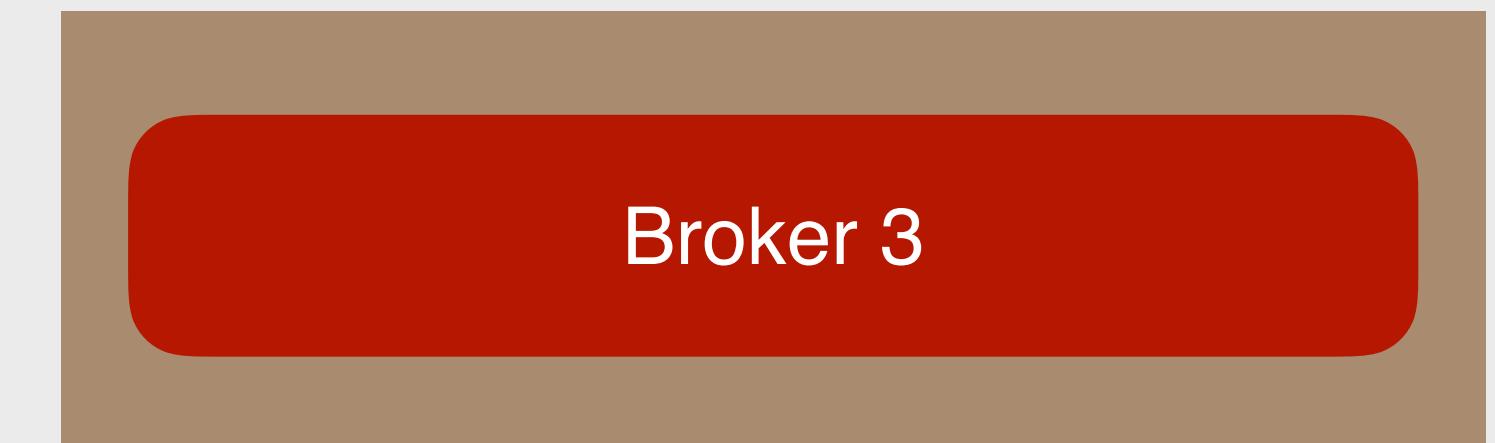
Pod 6



Pod 7



Pod 8



Pod 9

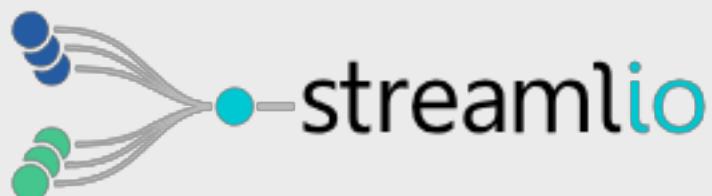
# BUILT-IN STATE MANAGEMENT IN FUNCTIONS

- ◆ Functions can store state in inbuilt storage
  - Framework provides a simple library to store and retrieve state
- ◆ Support server side operations like counters
- ◆ Simplified application development
  - No need to standup an extra system

# DISTRIBUTED STATE IN FUNCTIONS

```
import org.apache.pulsar.functions.api.Context;
import org.apache.pulsar.functions.api.PulsarFunction;

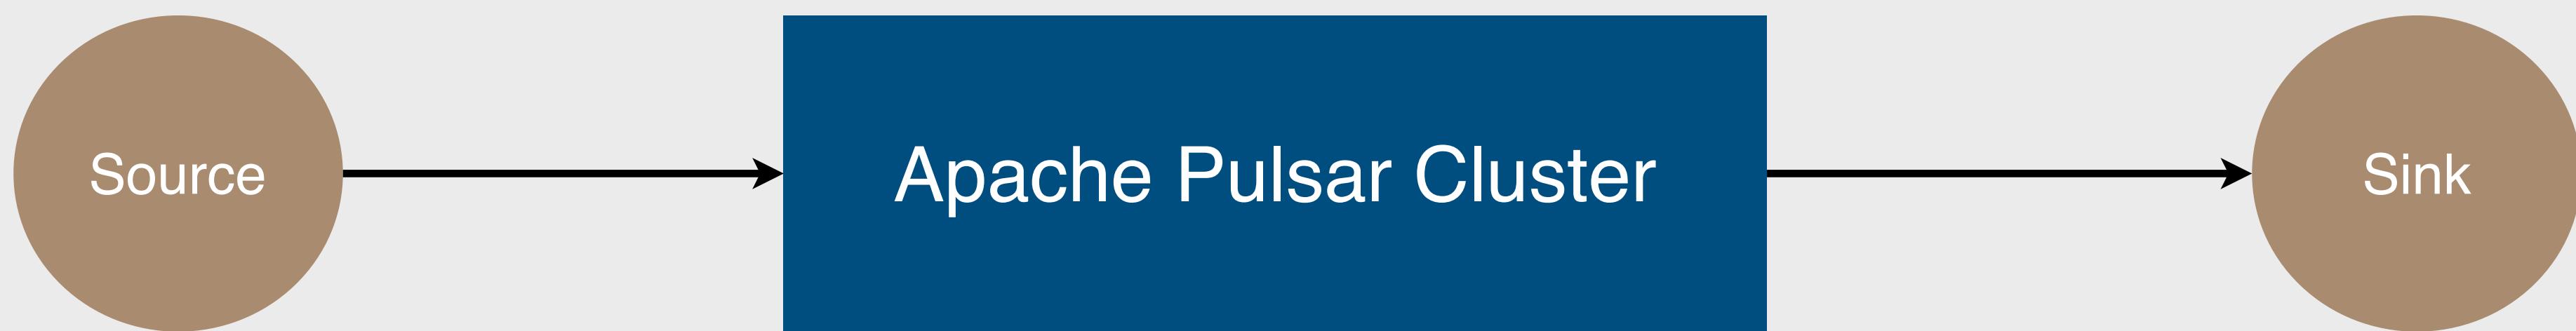
public class CounterFunction implements PulsarFunction<String, Void> {
    @Override
    public Void process(String input, Context context) throws Exception {
        for (String word : input.split("\\.")) {
            context.incrCounter(word, 1);
        }
        return null;
    }
}
```



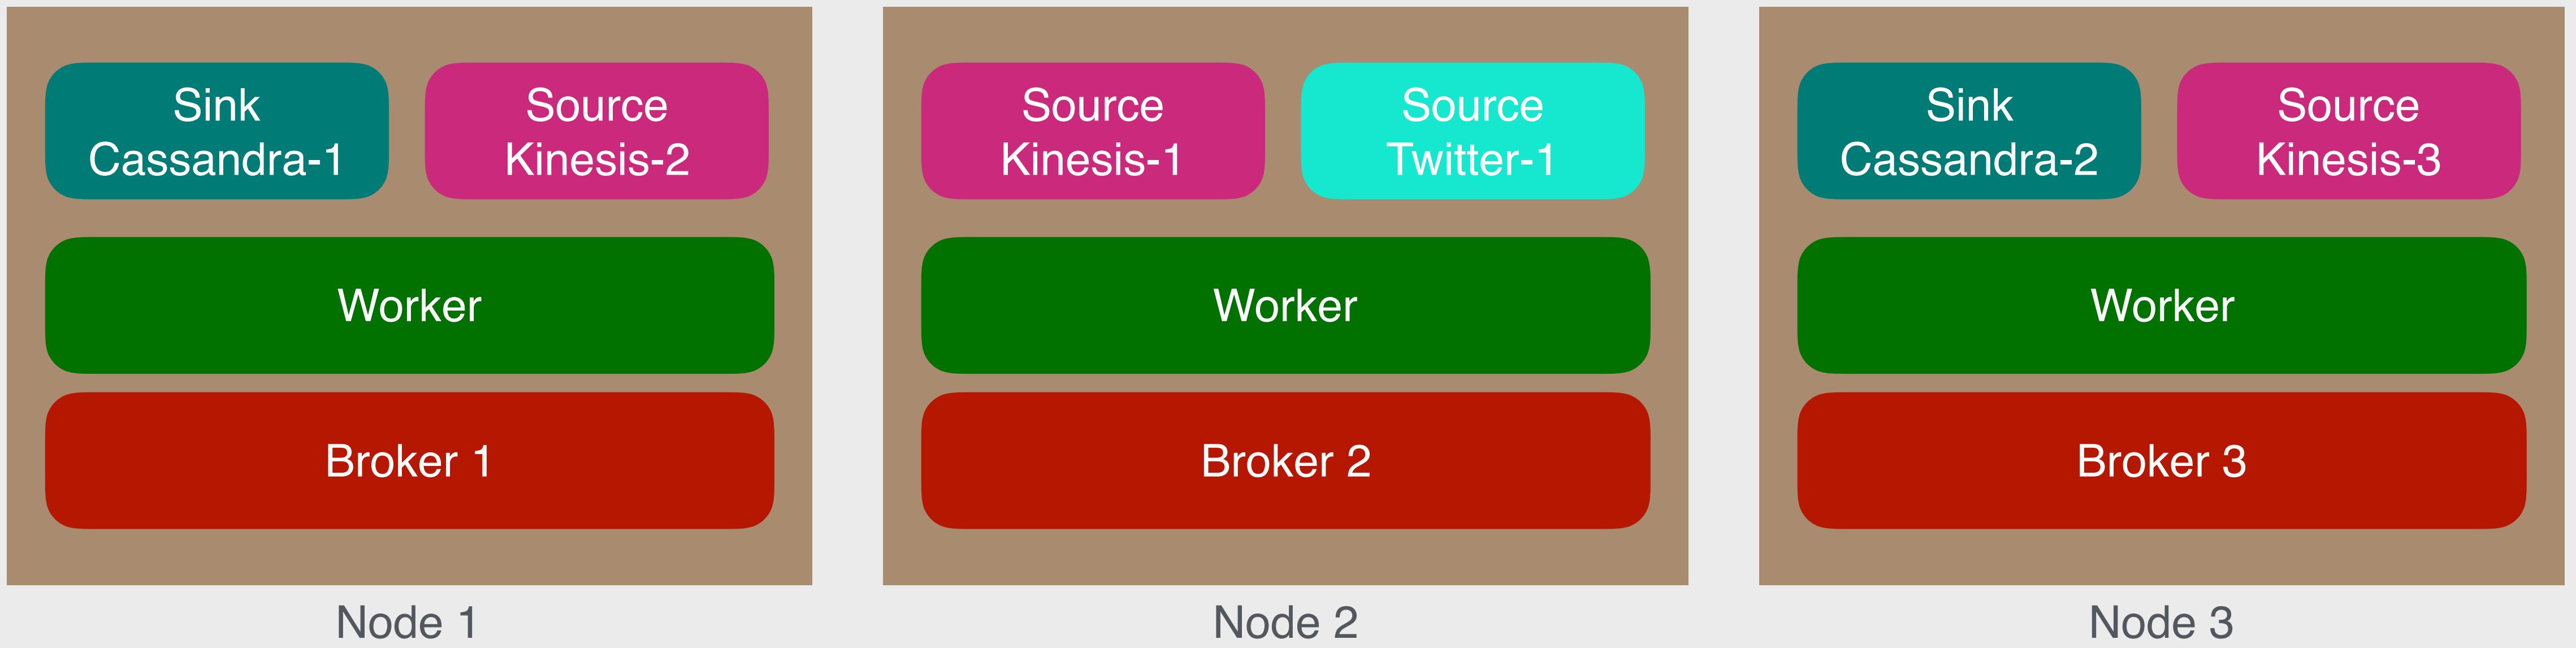
# PULSAR - DATA IN AND OUT

- ◆ Users can write custom code using Pulsar producer and consumer API
- ◆ Challenges
  - Where should the application to publish data or consume data from Pulsar?
  - How should the application to publish data or consume data from Pulsar?
- ◆ Current systems have no organized and fault tolerant way to run applications that ingress and egress data from and to external systems

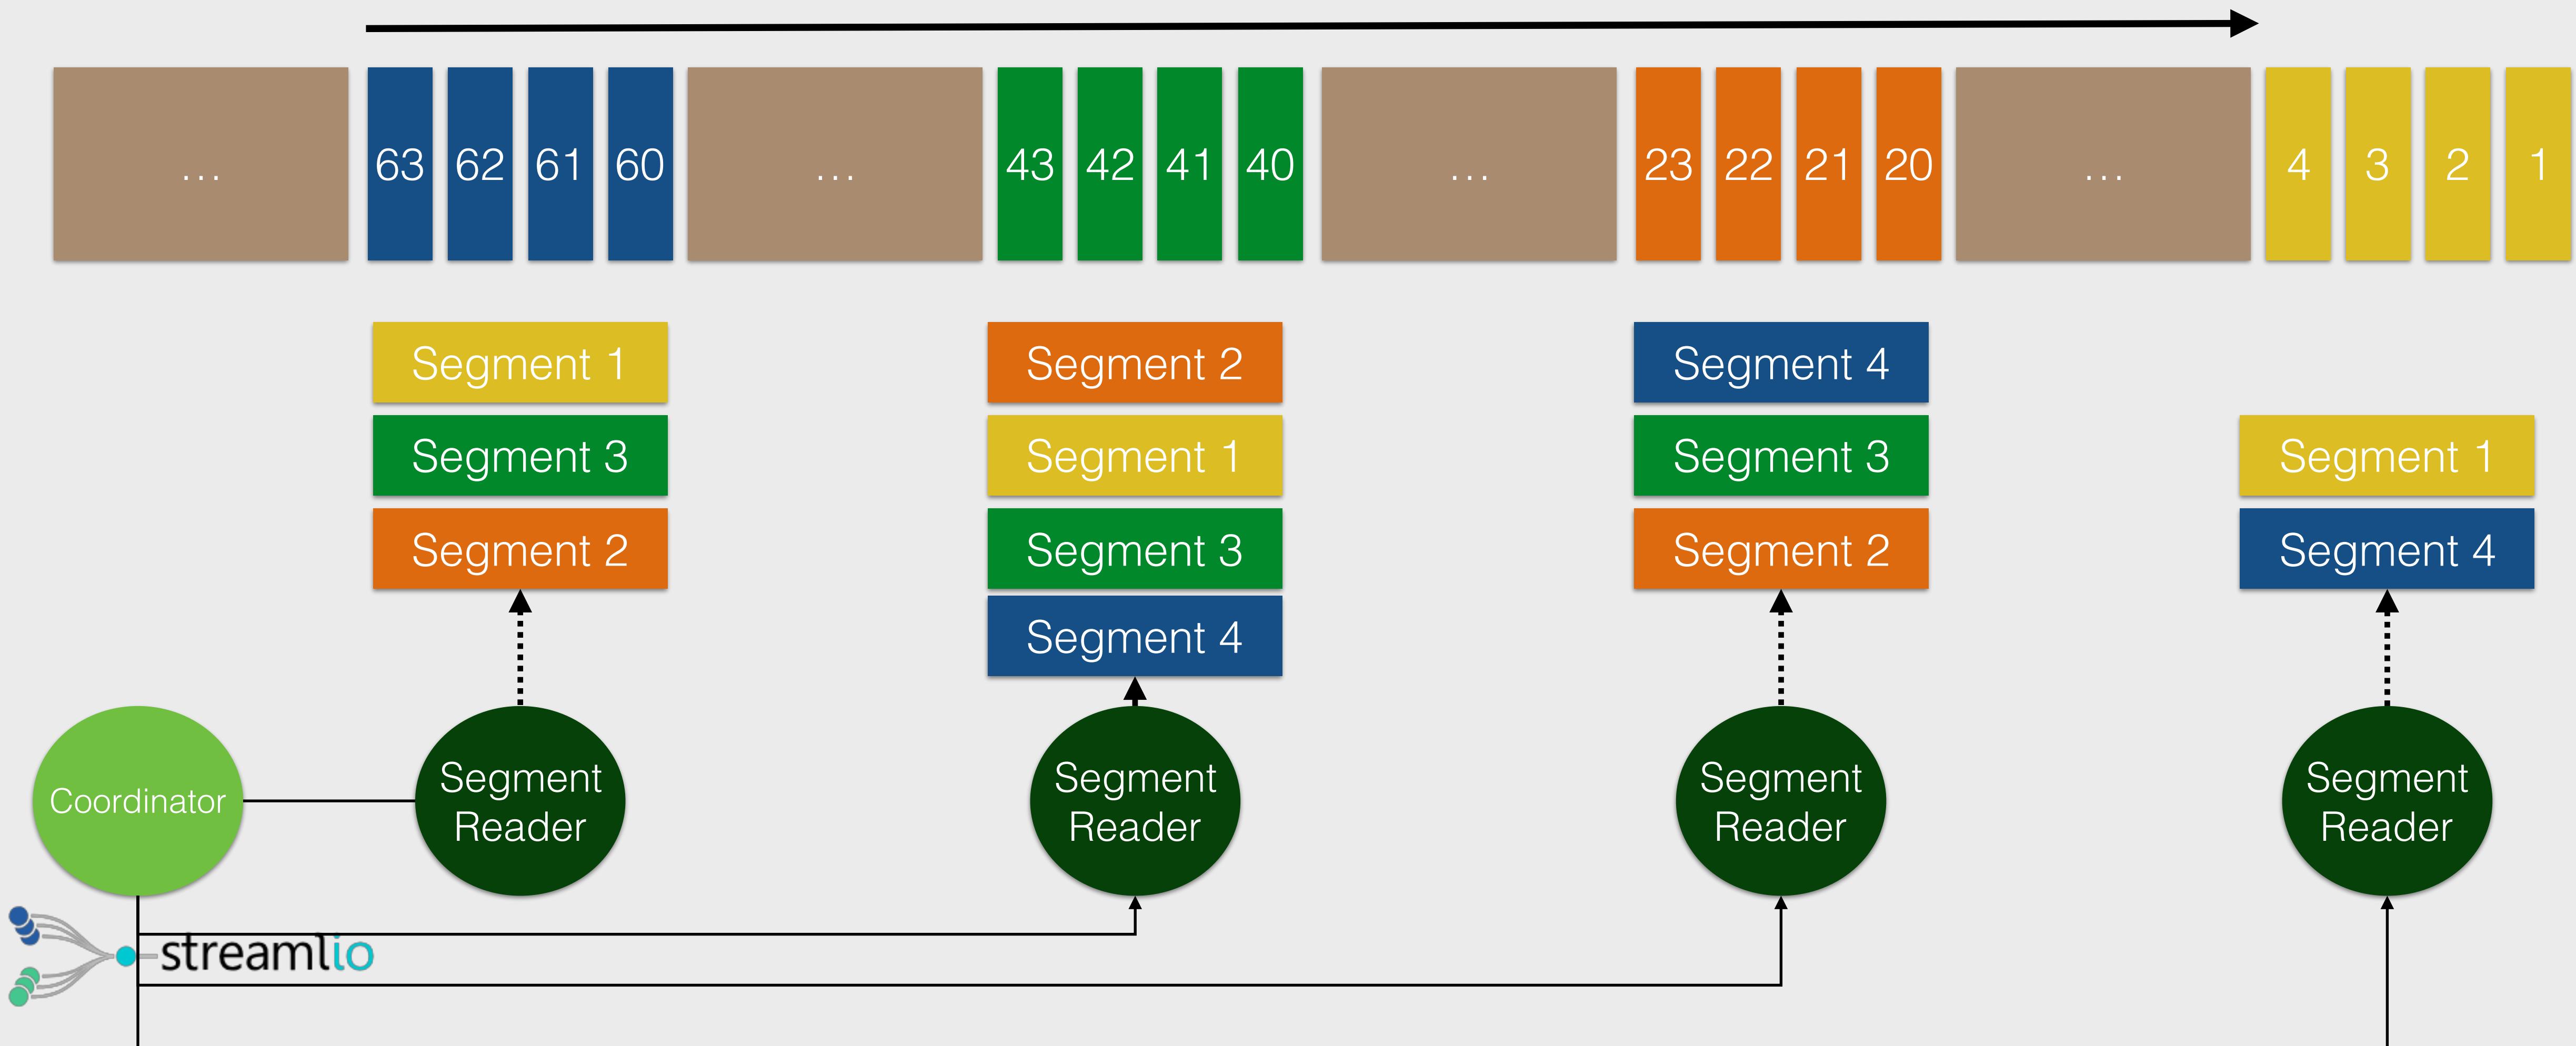
# PULSAR IO TO THE RESCUE



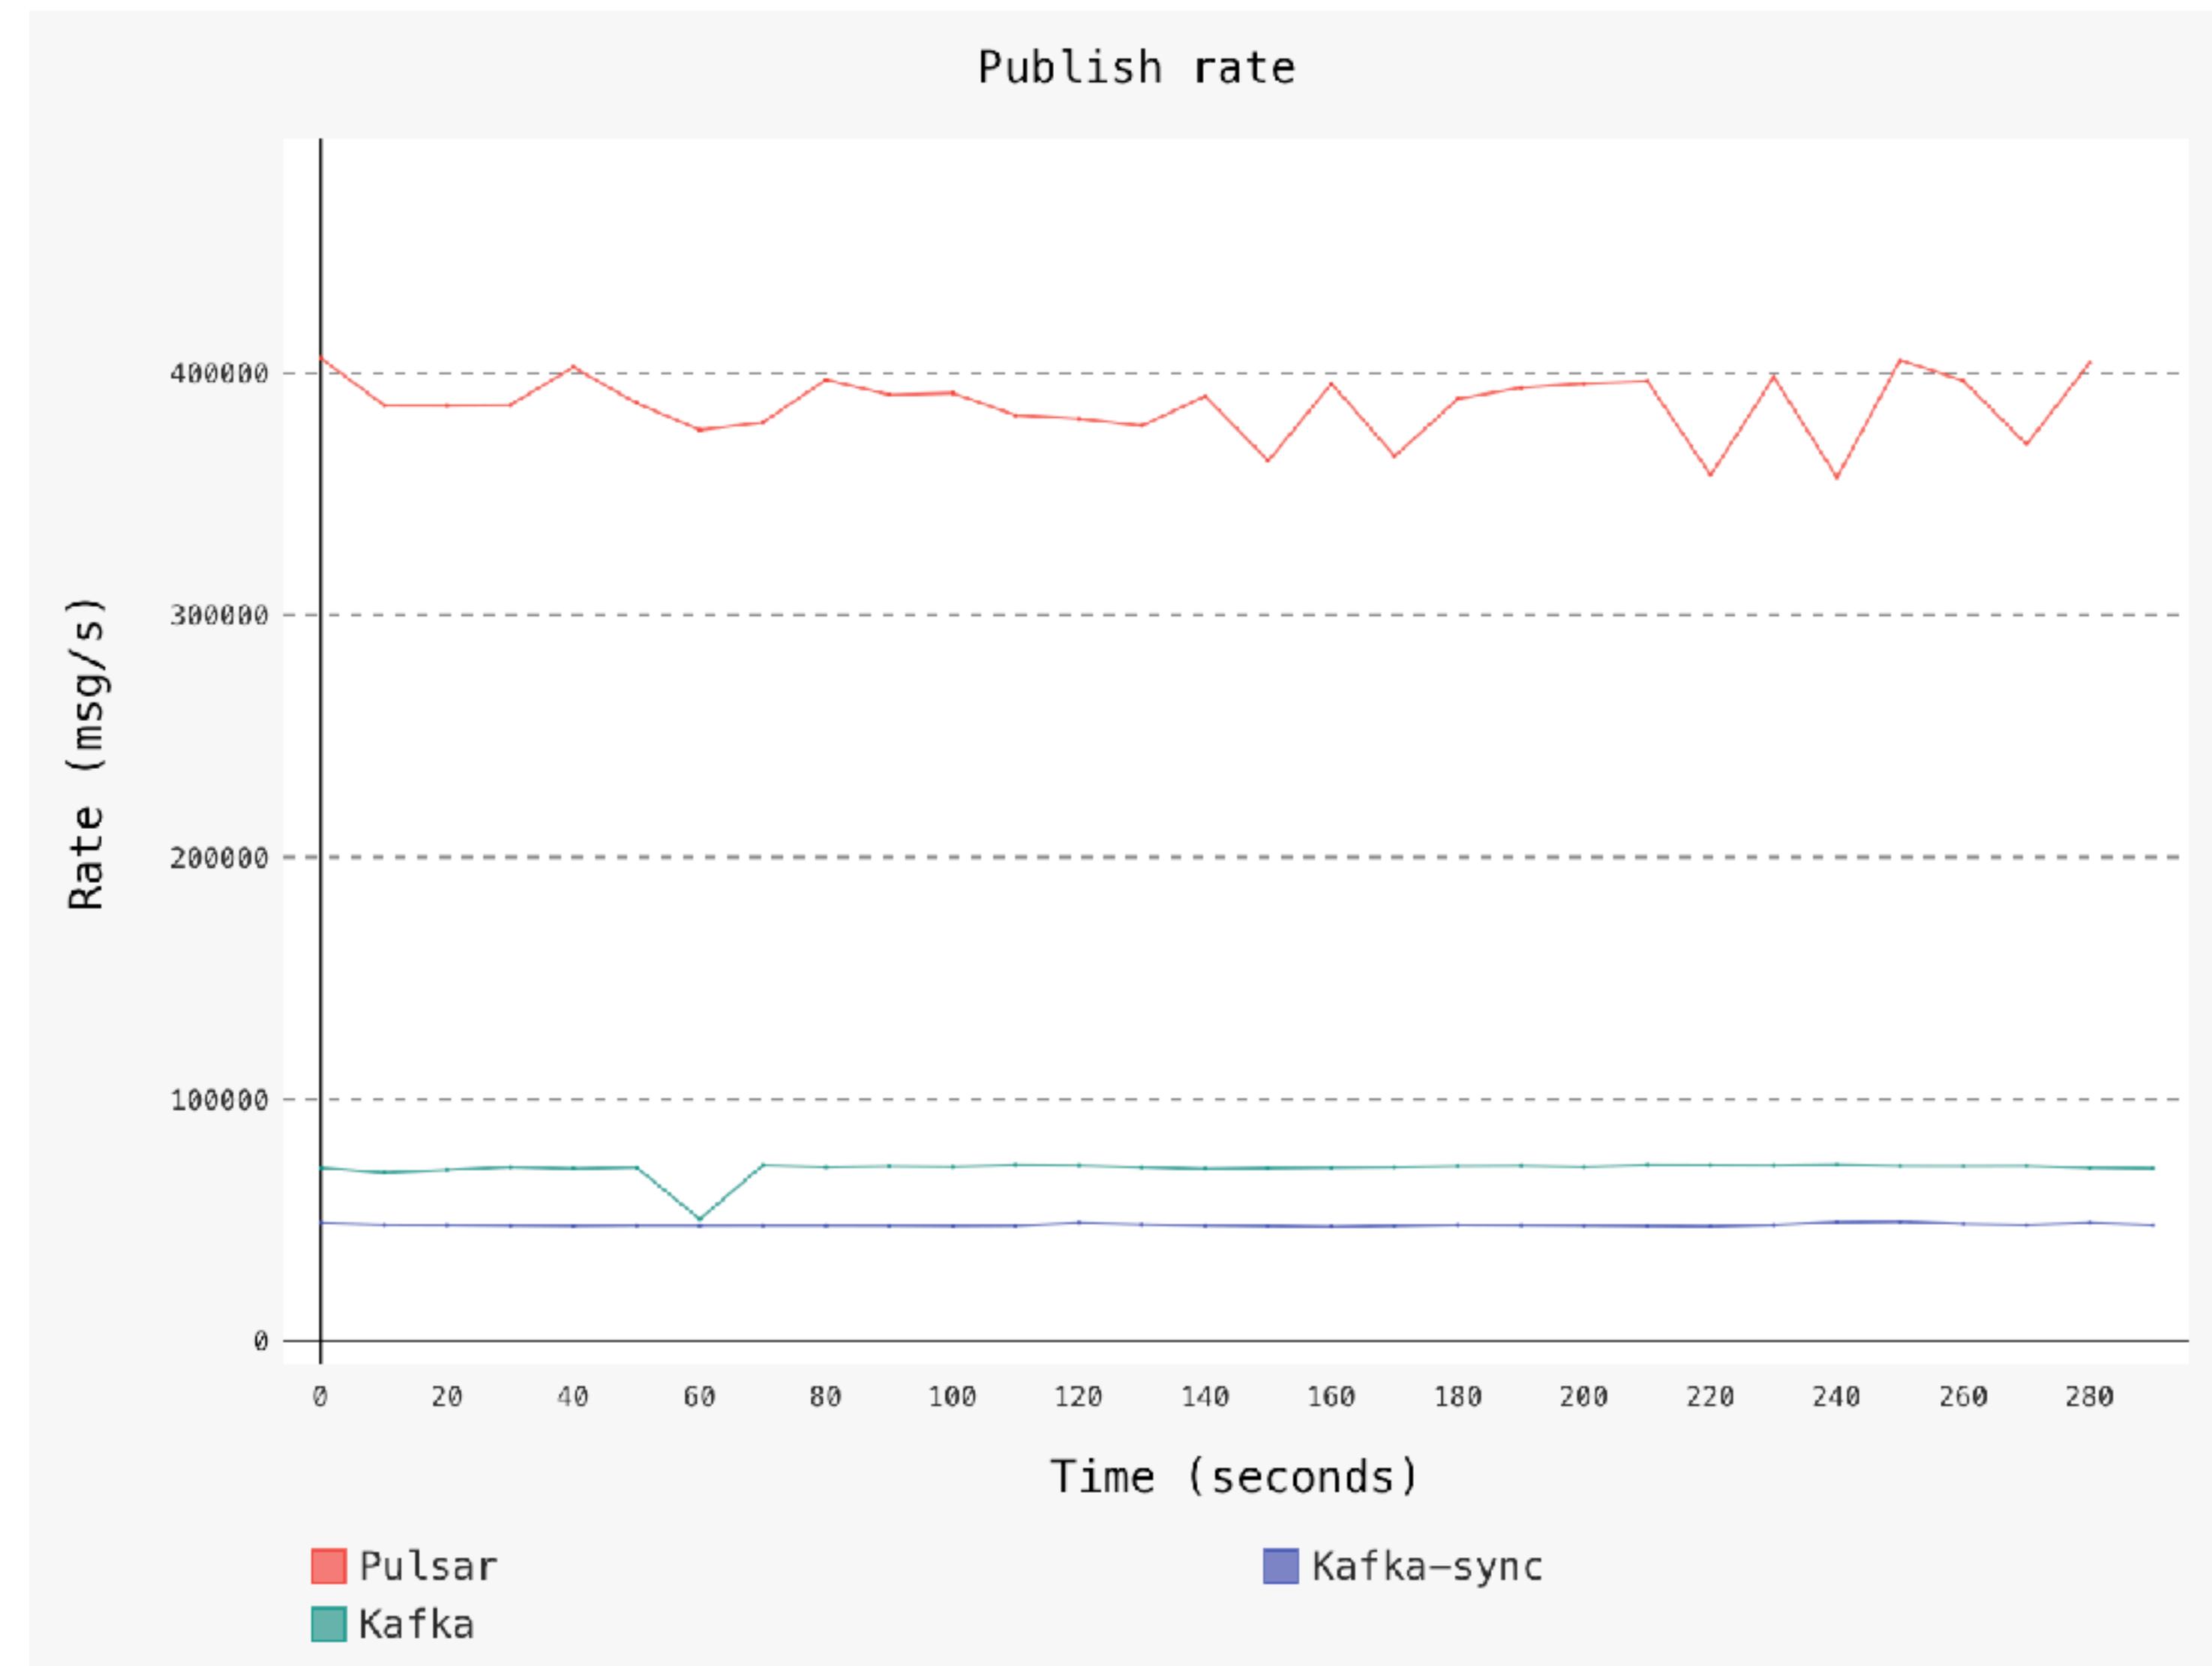
# PULSAR IO - EXECUTION



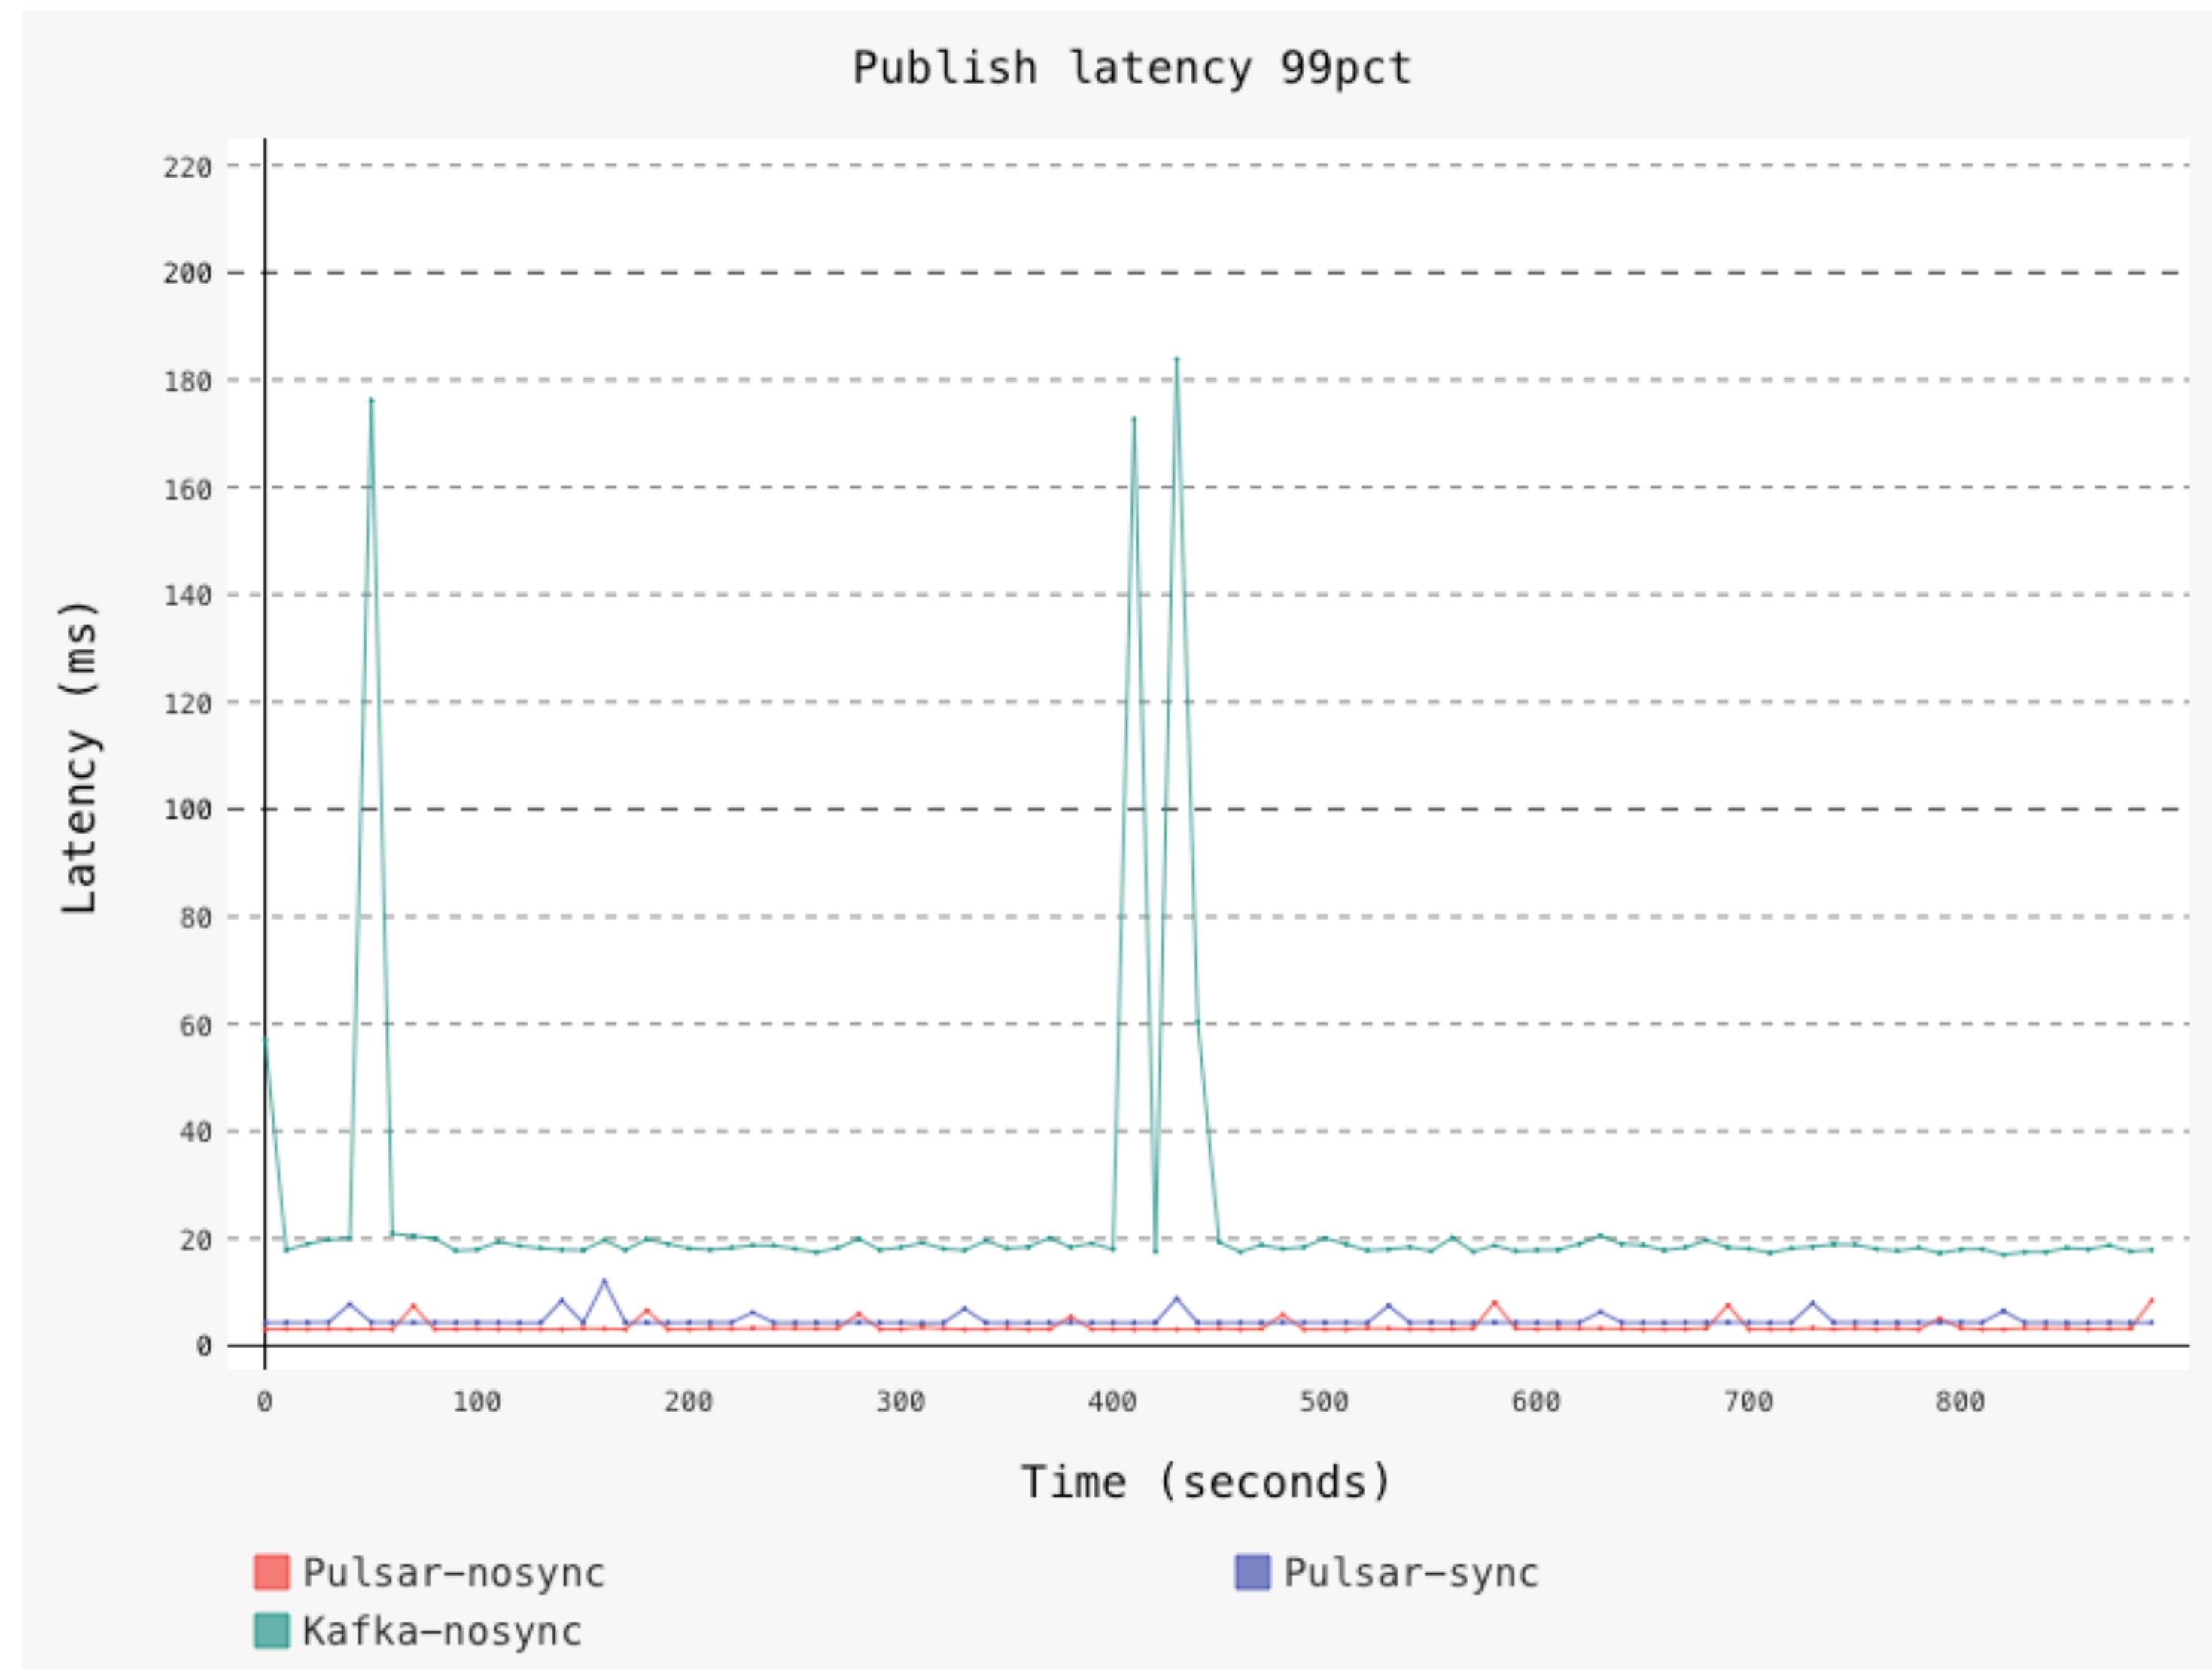
# INTERACTIVE QUERYING OF STREAMS - PULSAR SQL



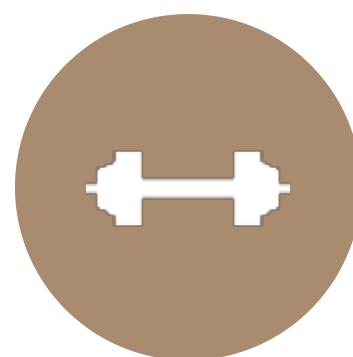
# PULSAR PERFORMANCE



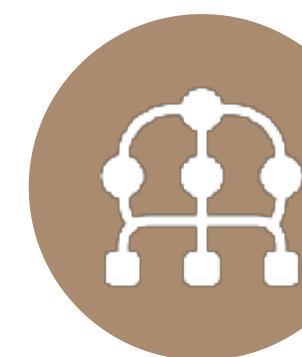
# PULSAR PERFORMANCE - LATENCY



# APACHE PULSAR vs. APACHE KAFKA

**Durability**

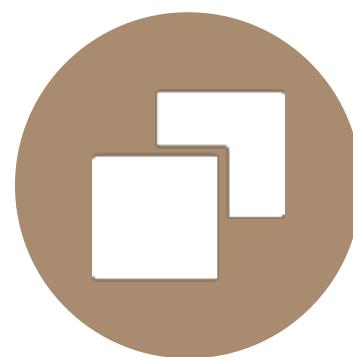
Data replicated and synced to disk

**Multi-tenancy**

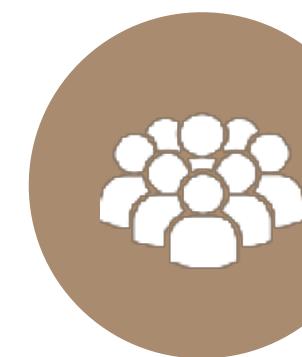
A single cluster can support many tenants and use cases

**Tiered Storage**

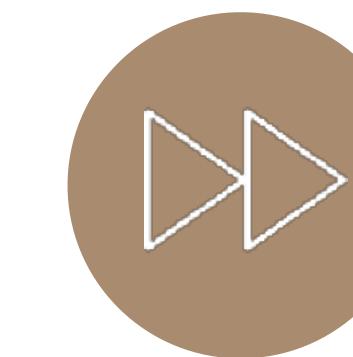
Hot/warm data for real time access and cold event data in cheaper storage

**Geo-replication**

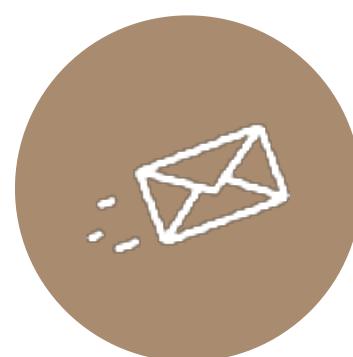
Out of box support for geographically distributed applications

**Seamless Cluster Expansion**

Expand the cluster without any down time

**Pulsar Functions**

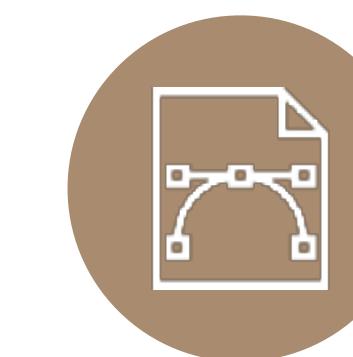
Flexible light weight compute

**Unified messaging model**

Support both Topic & Queue semantic in a single model

**High throughput & Low Latency**

Can reach 1.8 M messages/s in a single partition and publish latency of 5ms at 99pct

**Highly scalable**

Can support millions of topics, makes data modeling easier

# Examples of companies using Apache Pulsar

Streamlio  
outreach



**Growing funnel of validation and leads from outbound, inbound and open source**

Open source  
adopters



Open source  
evaluators



# Yahoo!

## Scenario

Need to collect and distribute user and data events to distributed global applications at Internet scale

## Challenges

- Multiple technologies to handle messaging needs
- Multiple, siloed messaging clusters
- Hard to meet scale and performance
- Complex, fragile environment



## Solution

- Central event data bus using Apache Pulsar
- Consolidated multiple technologies and clusters into a single solution
- Fully-replicated across 8 global datacenter
- Processing >100B messages / day, 2.3M topics

# APACHE PULSAR IN PRODUCTION @SCALE

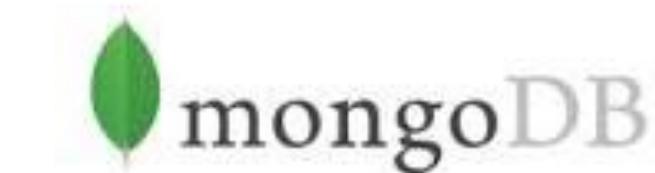
- 4+ years
- Serves 2.3 million topics
- 700 billion messages/day
- 500+ bookie nodes
- 200+ broker nodes
- Average latency < 5 ms
- 99.9% 15 ms (strong durability guarantees)
- Zero data loss
- 150+ applications
- Self served provisioning
- Full-mesh cross-datacenter replication - 8+ data centers

# Growing ecosystem

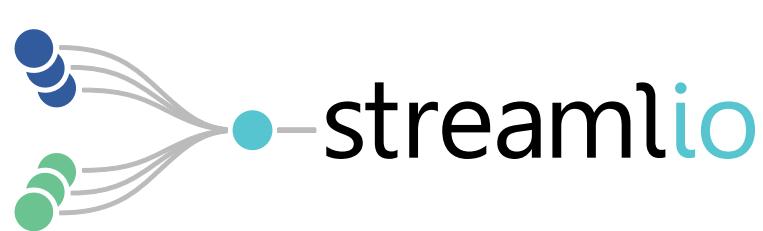
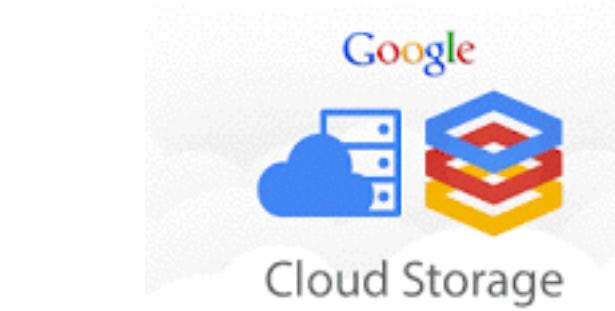
AEROSPIKE



ORACLE



Spark  
Streaming

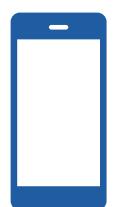


# Use Cases

# Example use cases



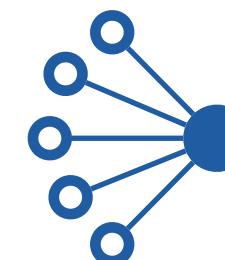
Real-time monitoring  
and notifications



Interactive  
applications



Log processing  
and analytics



IoT analytics



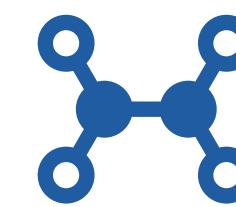
Streaming data  
transformation



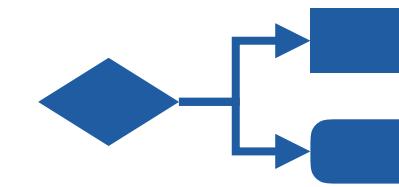
Real-time  
analytics



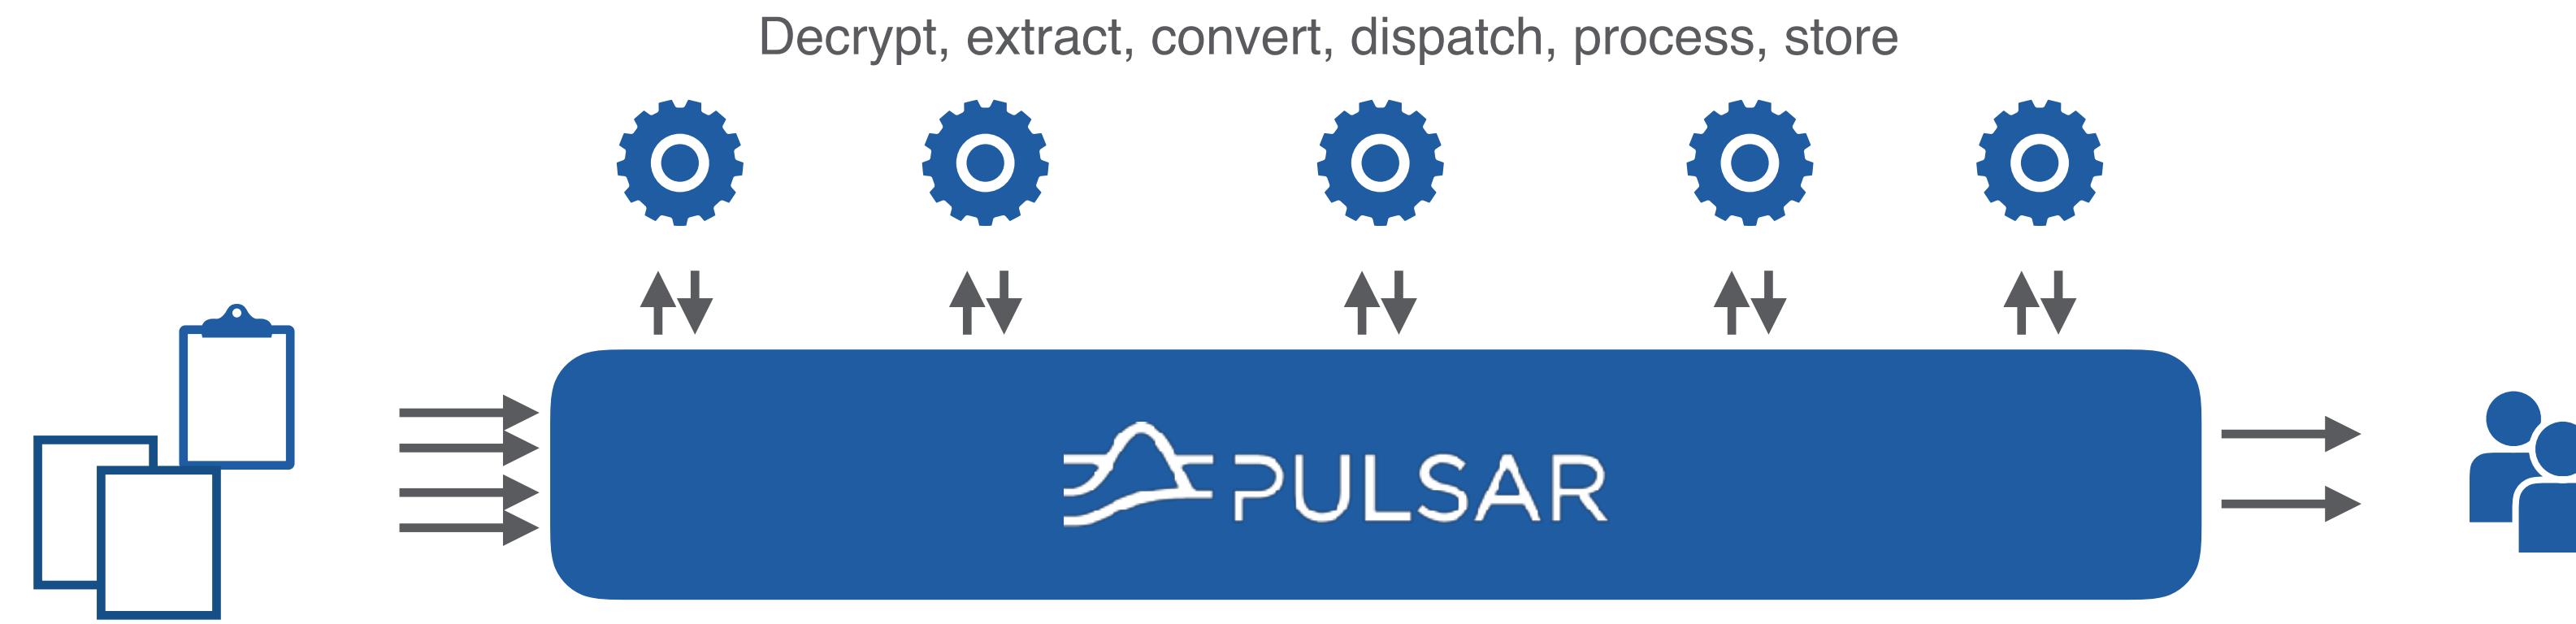
Data  
distribution



Event-driven  
workflows



# Data-driven workflows



## Scenario

Application processes incoming events and documents that generate processing workflows

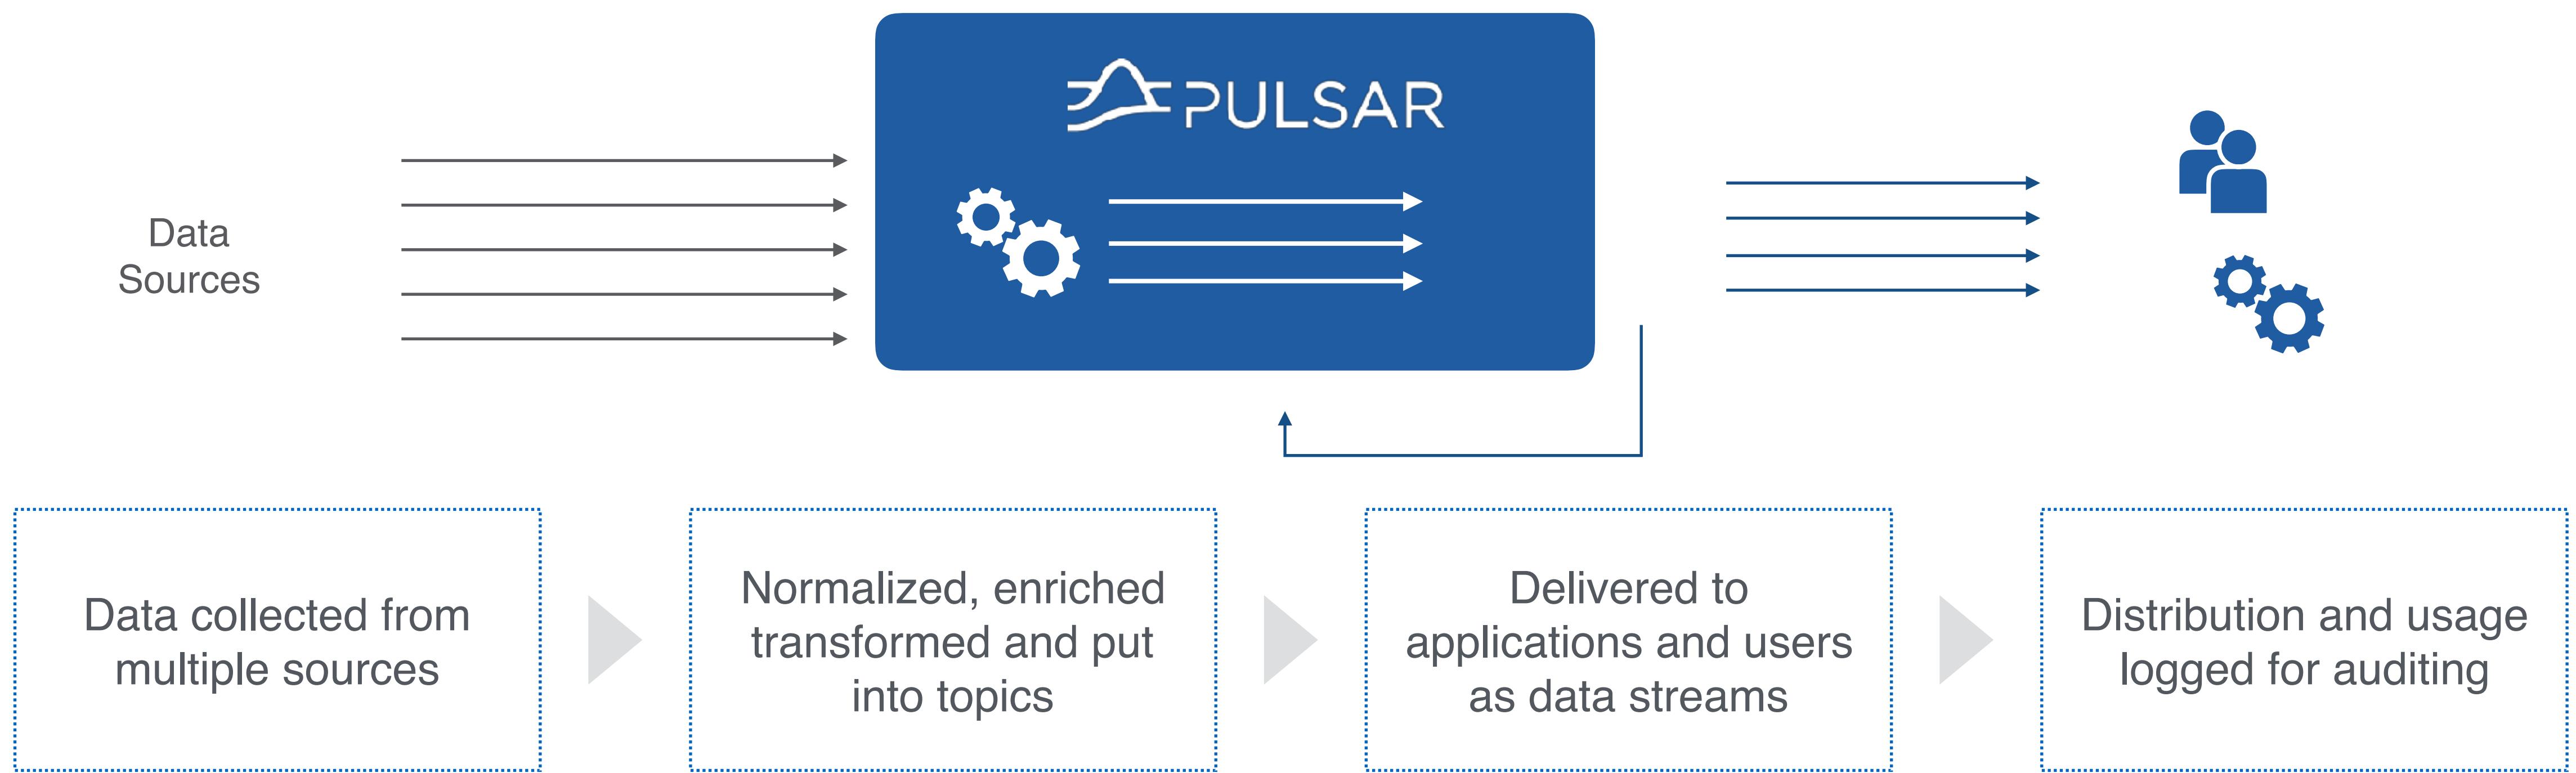
## Challenges

Operational burdens and scalability challenges of existing technologies growing as data grows

## Solution

Process incoming events and data and create work queues in same system

# Data distribution



# Simplifying the data pipeline

## Scenario

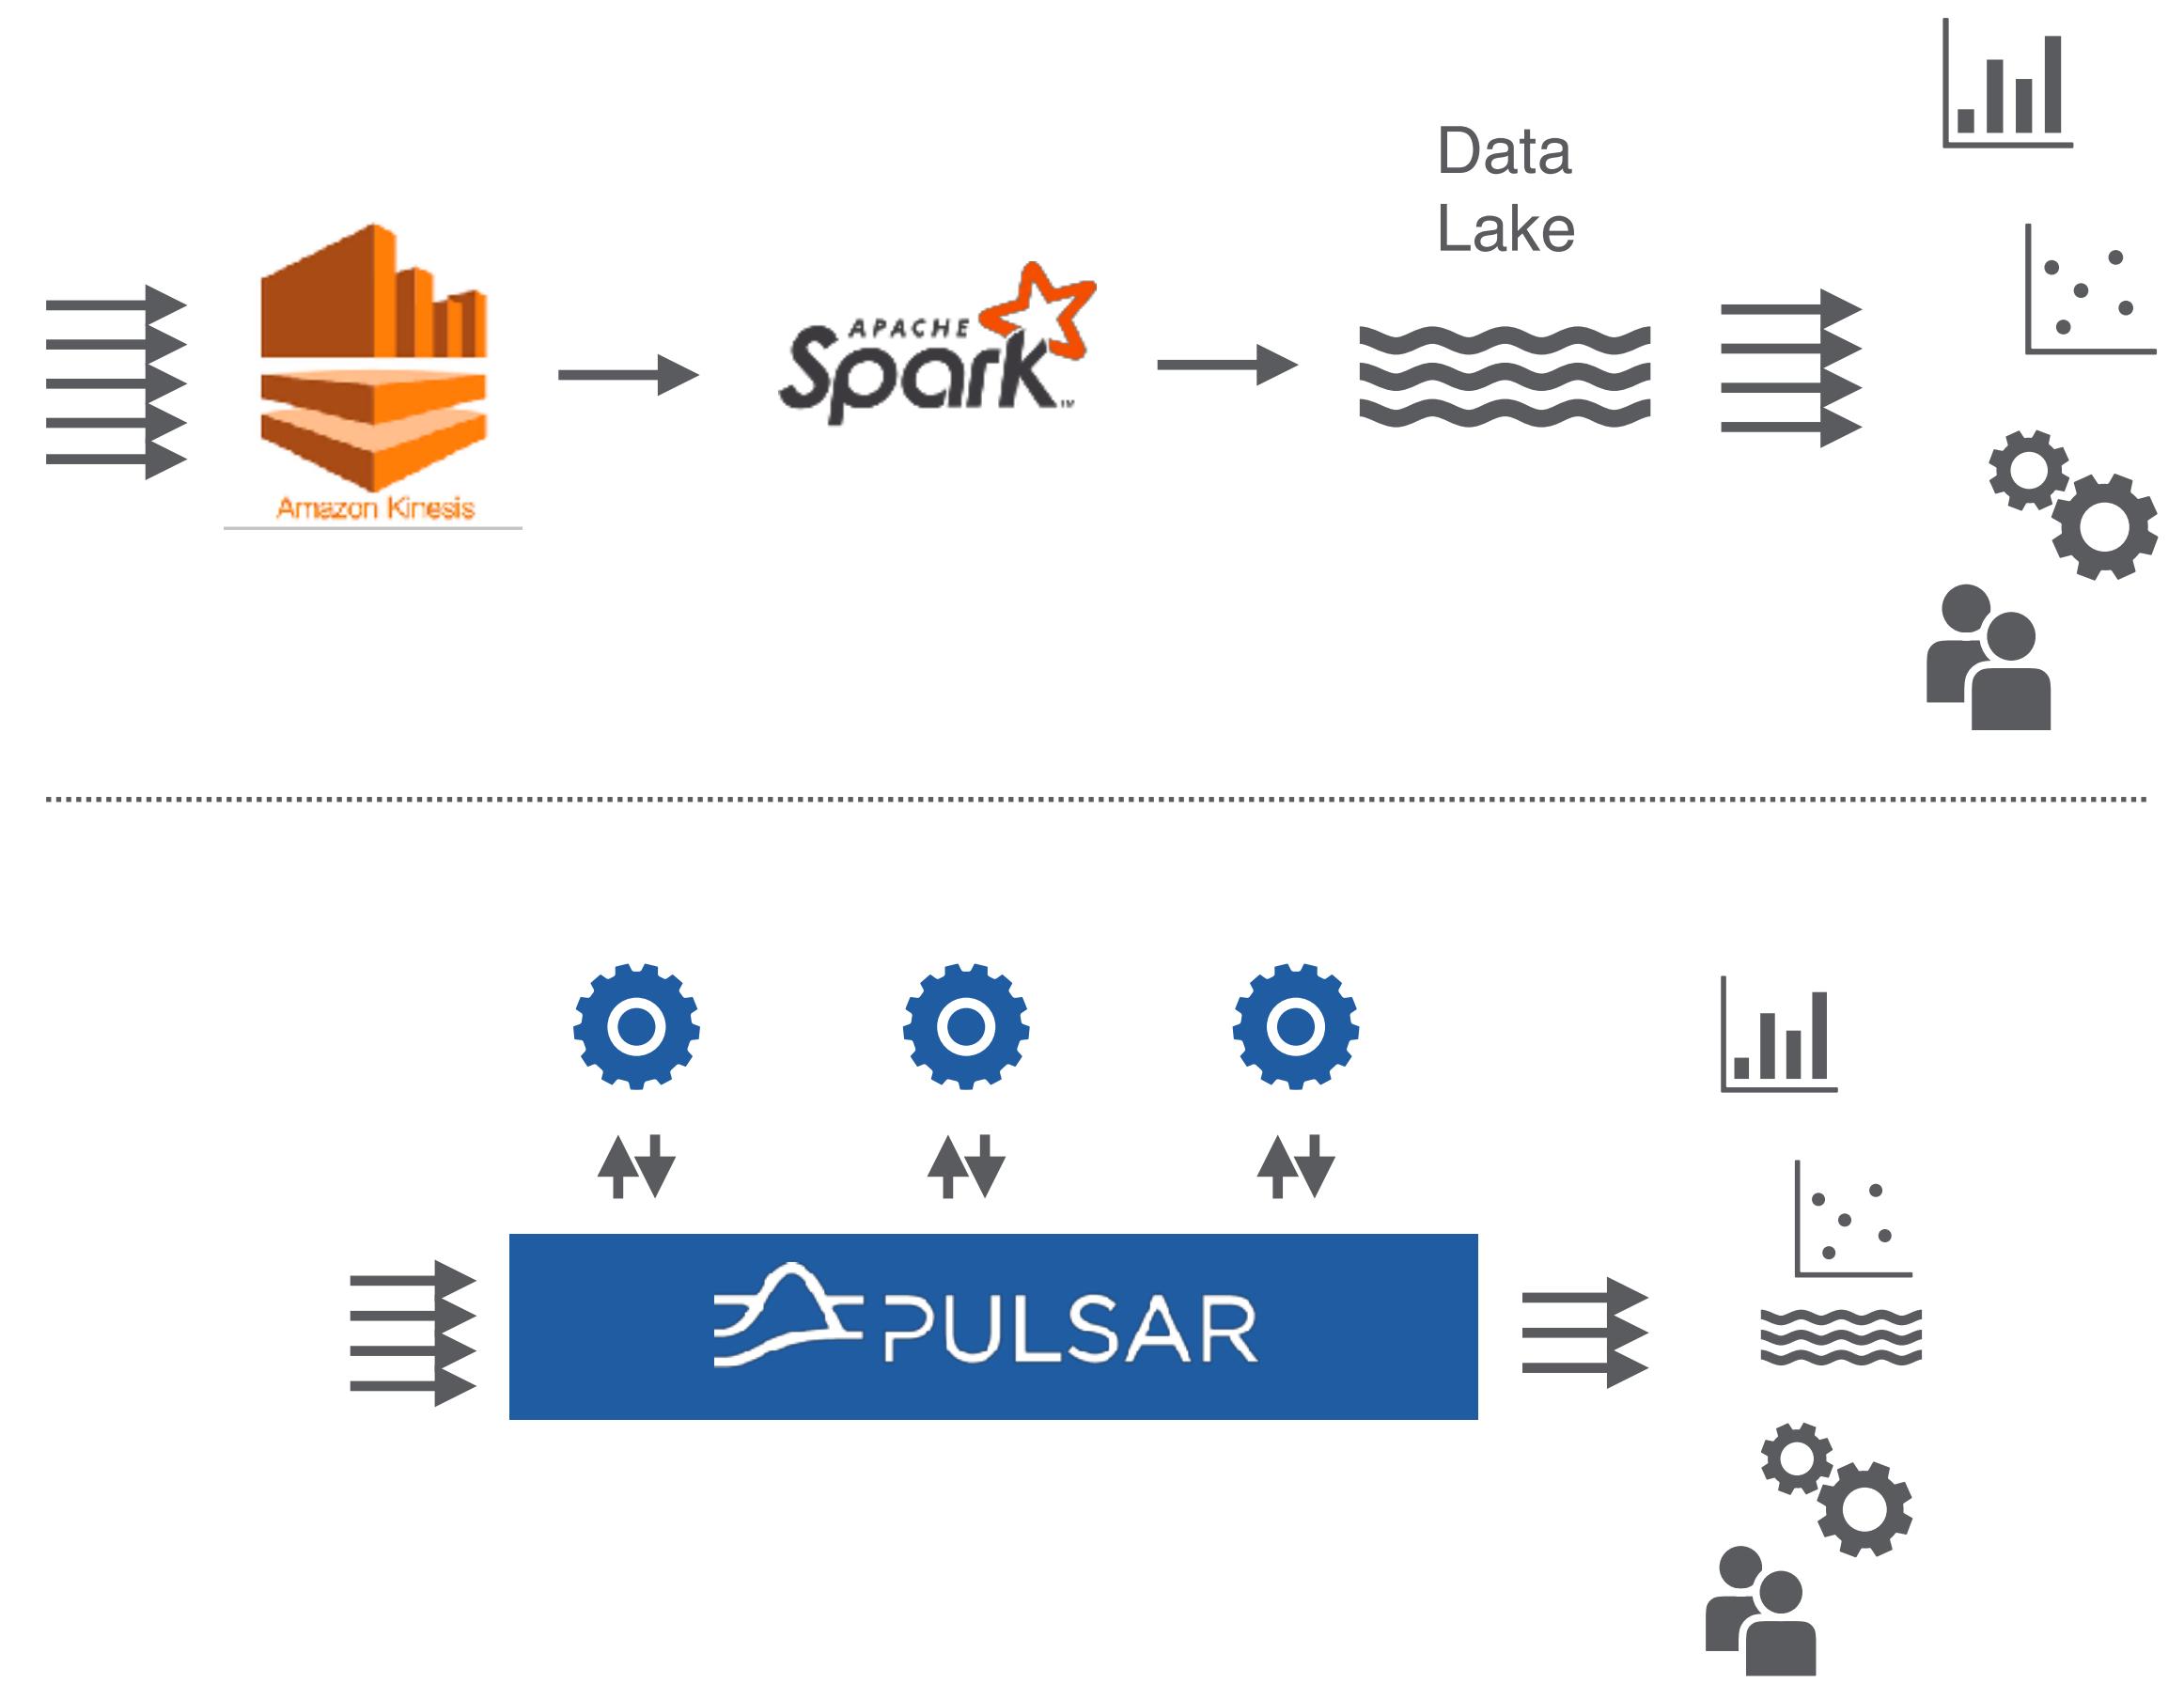
Retail analytics software provider brings together operational and market research data for insights.

## Challenges

Existing Kinesis + Spark + data lake infrastructure was unnecessarily complex and burdensome to operate and maintain.

## Solution

- Replaced Kinesis + Spark with Apache Pulsar
- Simplified data transformation pipeline
- Reduced operations burdens



# Event sourcing

## Problem

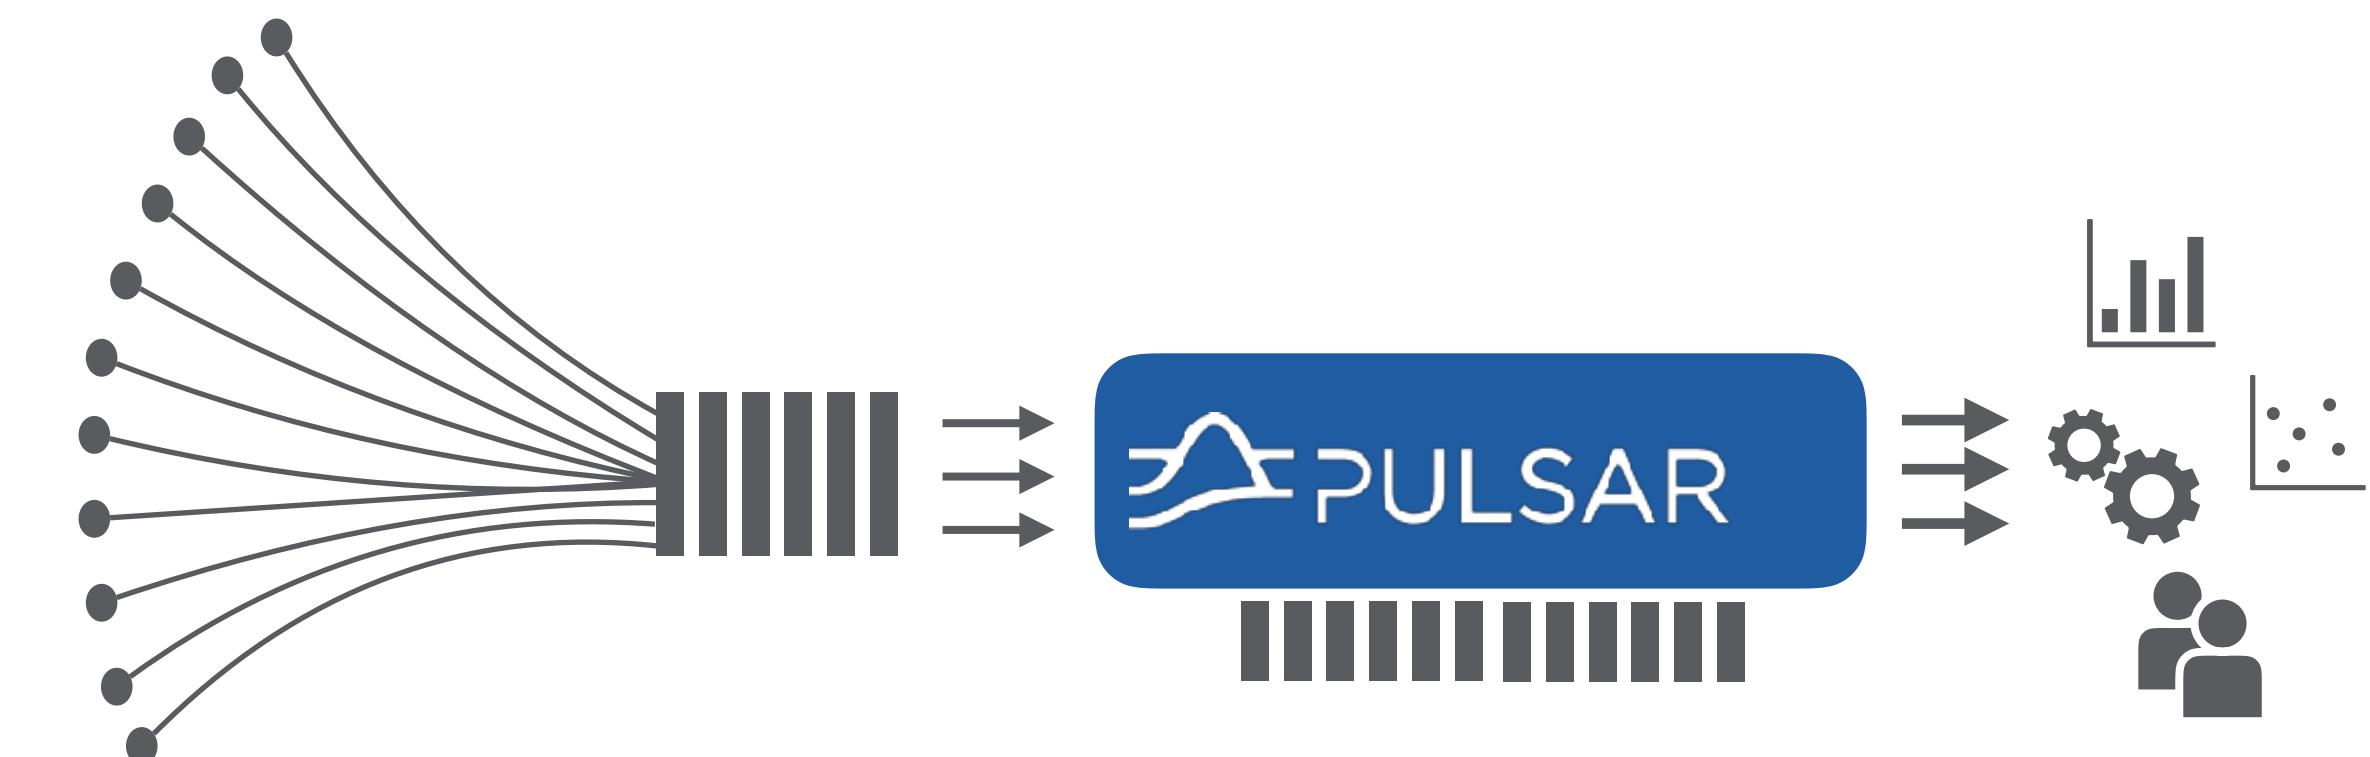
Event-driven applications require long-term retention of data streams, but current technologies are cumbersome and expensive to use for data retention and cannot efficiently replay data.

## Solution

Deploy Apache Pulsar for long-term retention and scalable processing and distribution of event data.

## Why Streamlio

- Architected for scalable and efficient long-term storage
- High performance, scalable processing and distribution of data due to unique architecture



# IOT ENVIRONMENT

## Light Device



- ◆ Typically sensors
- ◆ Only one functionality
- ◆ Simple to configure
- ◆ Light weight protocols to communicate



## Smart Device



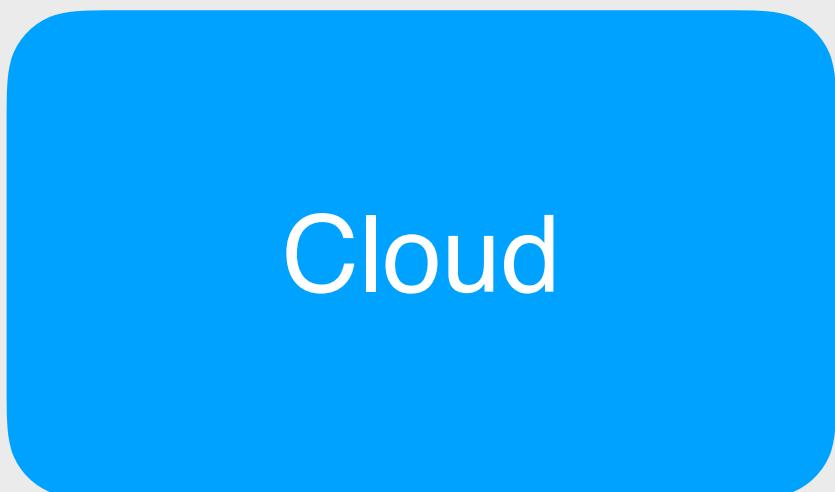
- ◆ Typically ARM based
- ◆ Multiple functionality
- ◆ Basic but generic computational logic, limited storage
- ◆ Light weight and proprietary protocols to communicate

## Edge Node



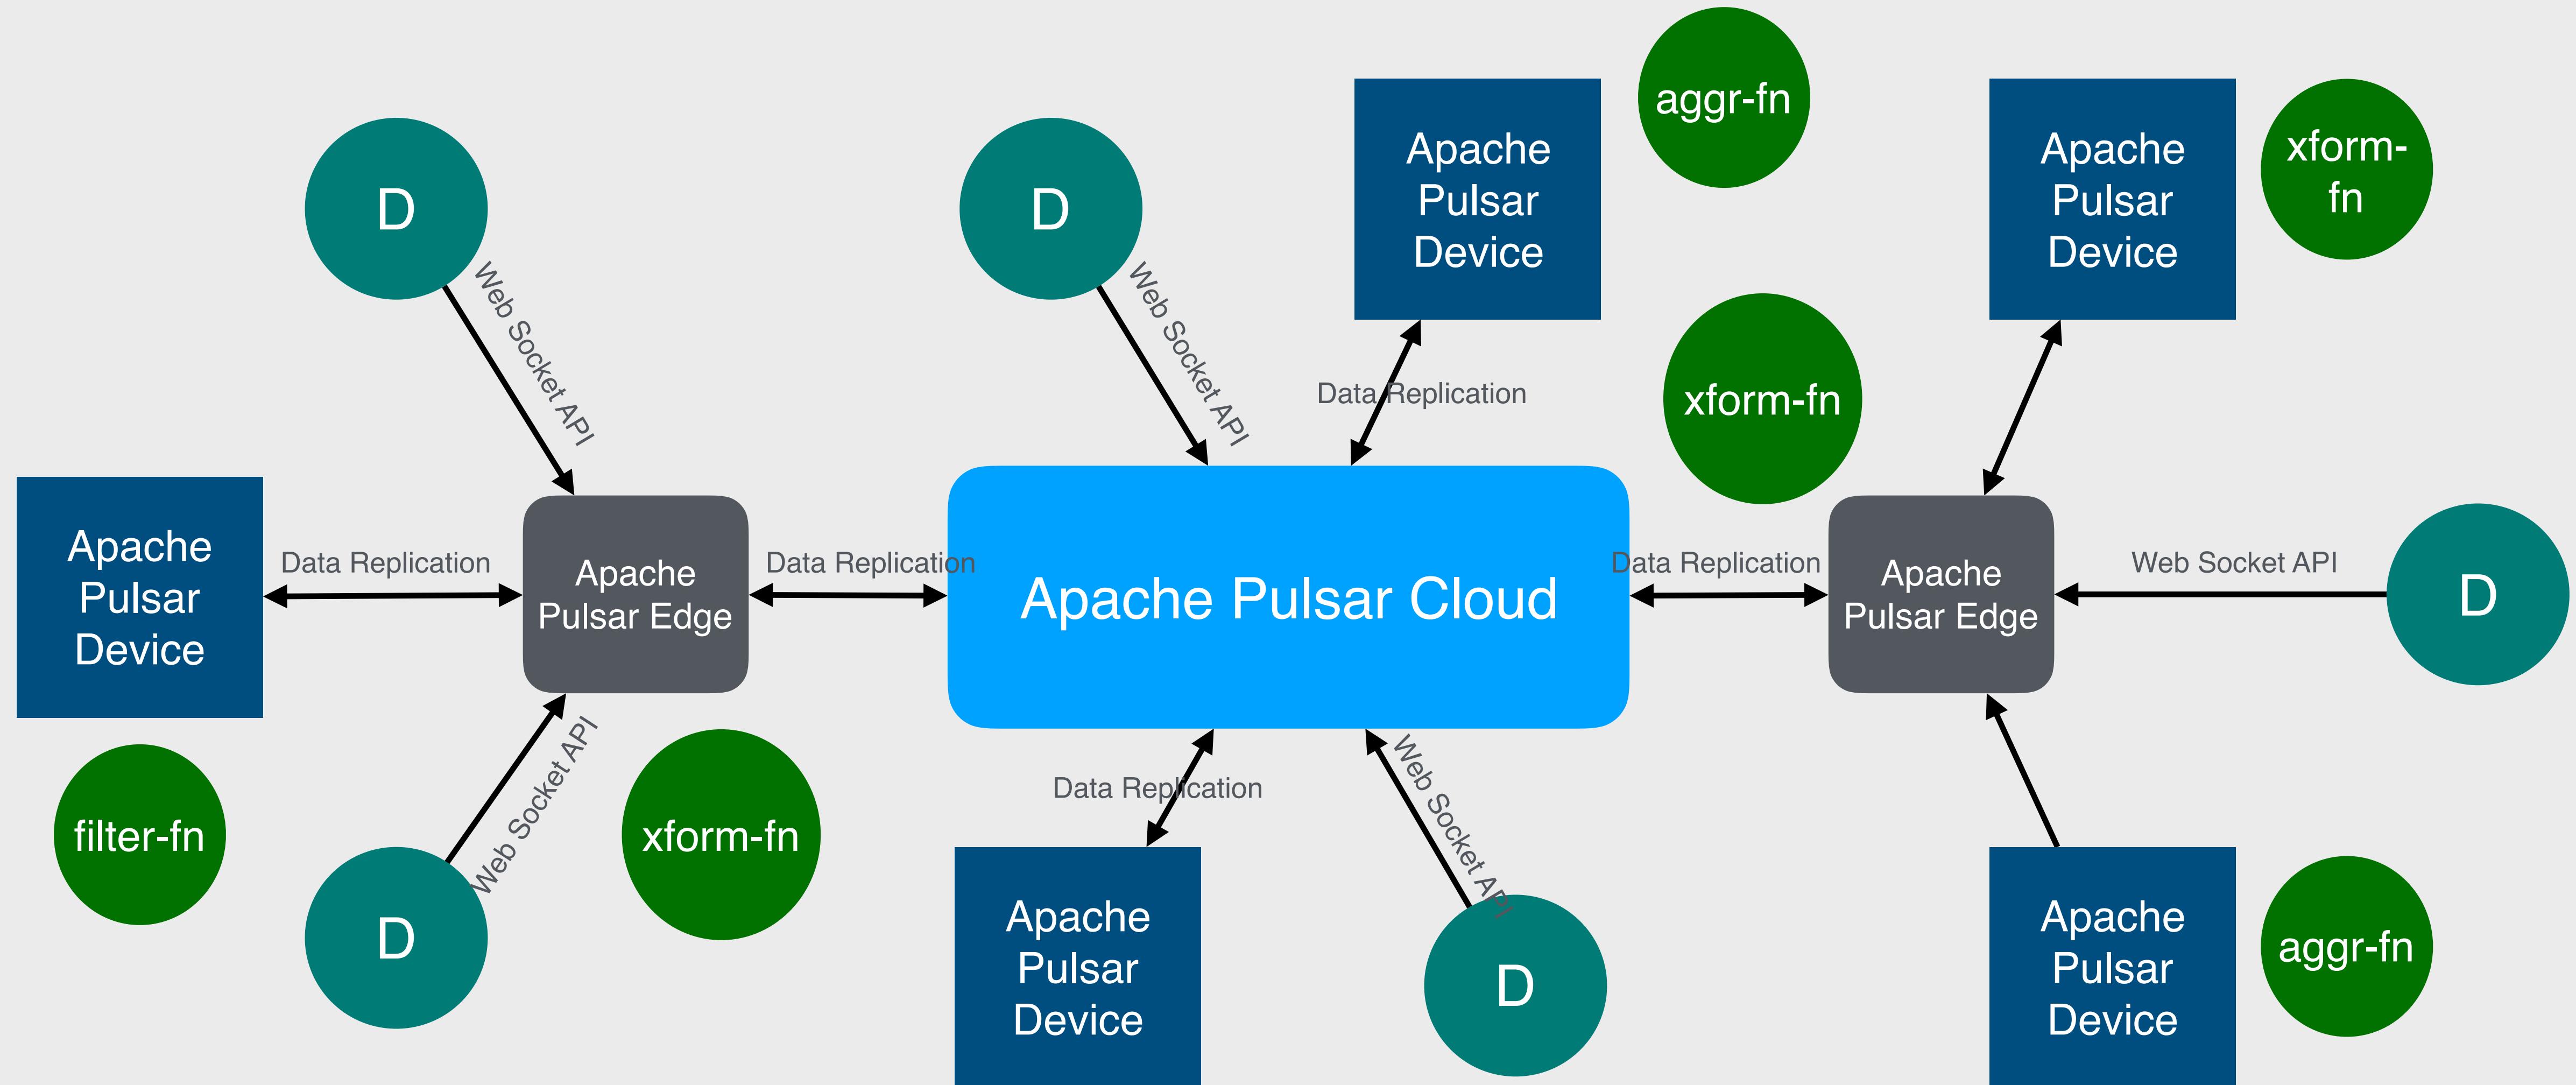
- ◆ Multicore based
- ◆ Versatile functionality
- ◆ Complex and generic computational logic, decent amount of storage
- ◆ Light weight and proprietary protocols to communicate

## Cloud



- ◆ Multiple machines
- ◆ Versatile functionality
- ◆ Complex and generic computational logic
- ◆ Lots of storage

# IoT DATA FABRIC WITH APACHE PULSAR



# Large Car Manufacturer: Connected vehicle



## Scenario

Continuously-arriving data generated by connected cars needs to be quickly collected, processed and distributed to applications and partners

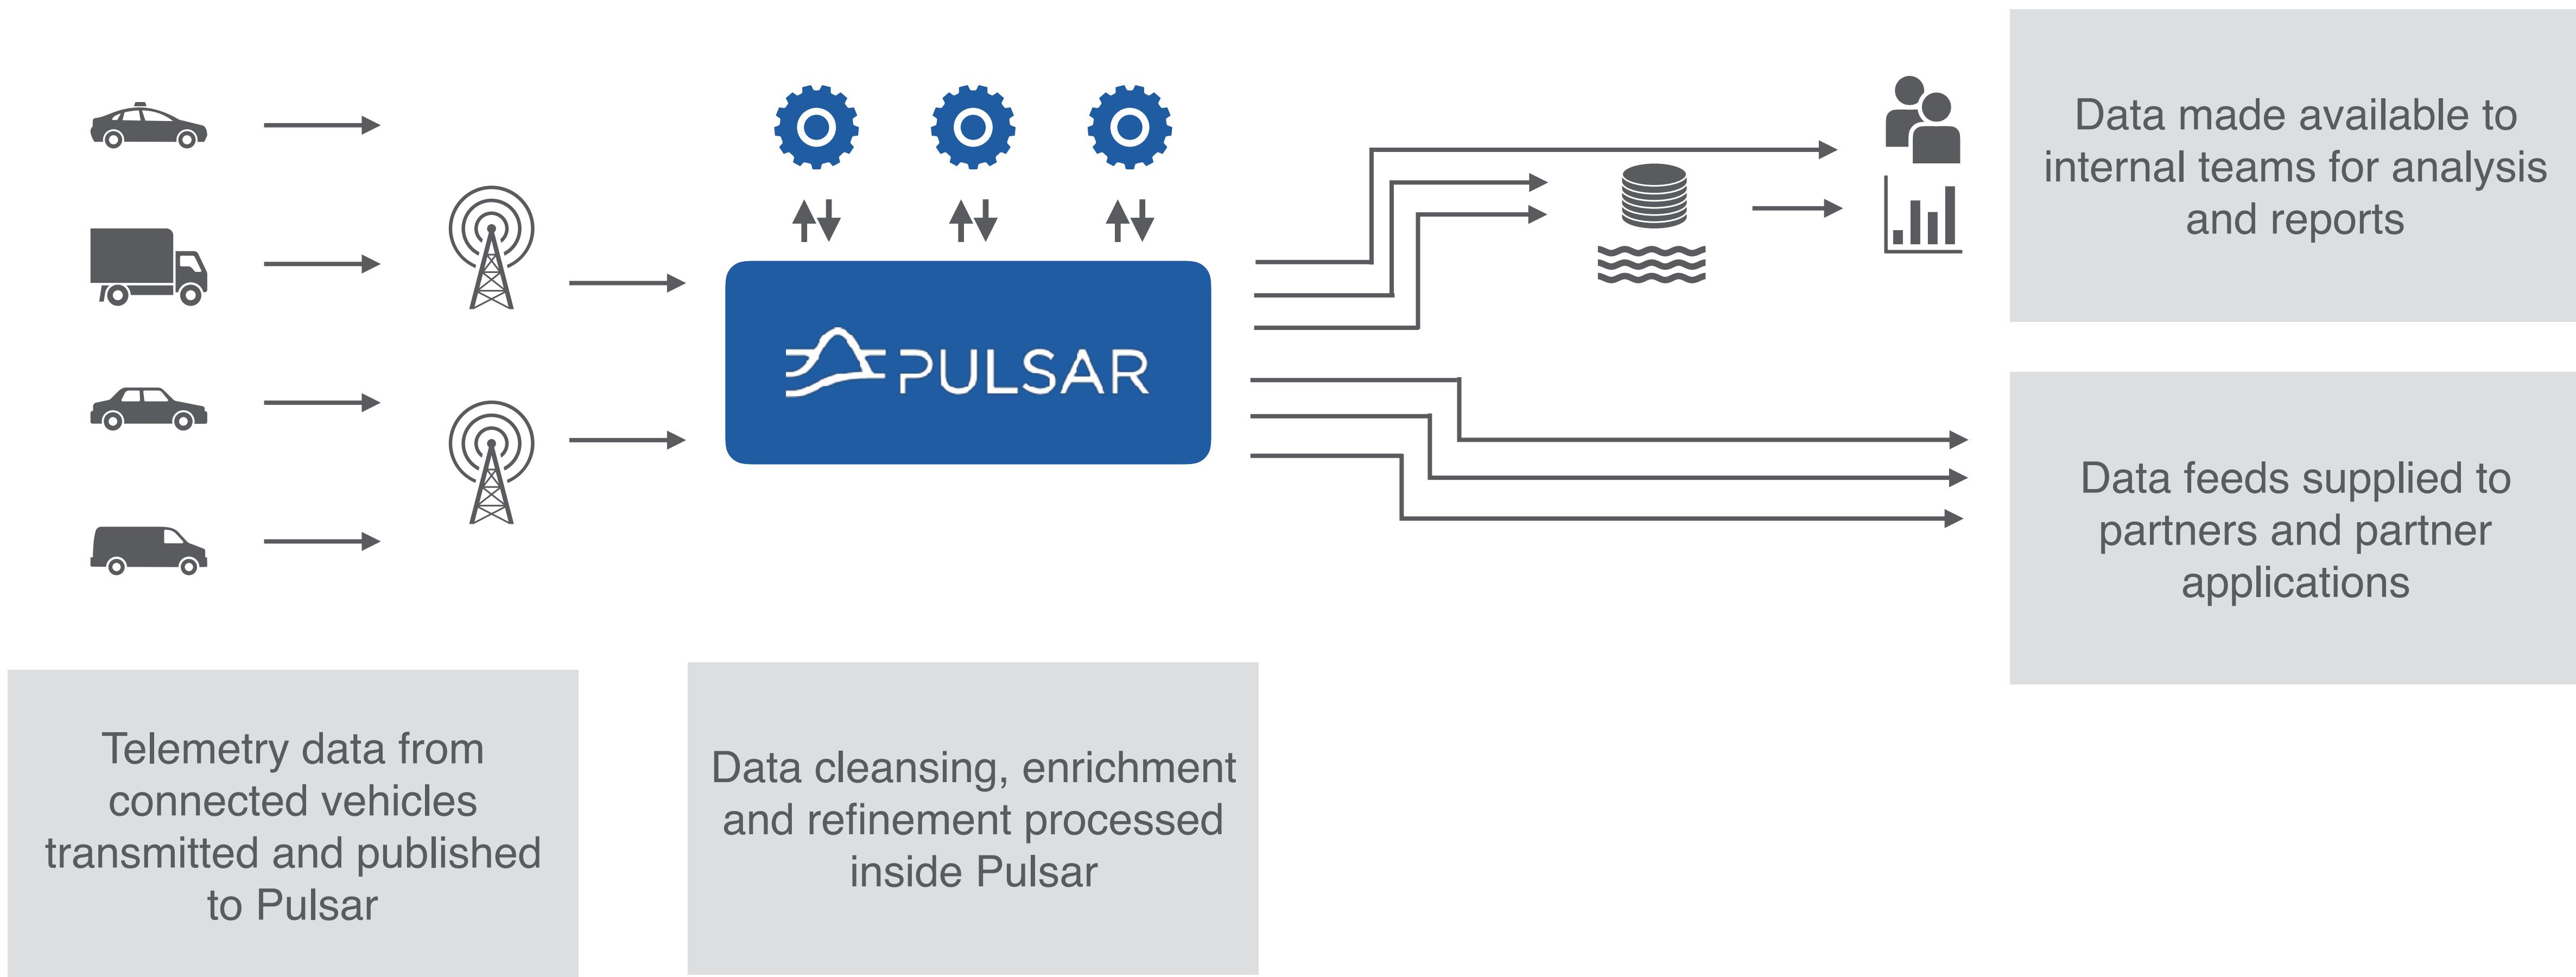
## Challenges

Require scalability to handle growing data sources and volumes without complex mix of technologies

## Solution

Leverage Streamlio solution to provide data backbone that can receive, transform, and distribute data at scale

# Large Car Manufacturer: Connected vehicle



# Large Car Manufacturer: Big Data Logging System



## Scenario

Continuously ingest logs from big data system for distributed to appropriate teams with appropriate log transformations and enrichment

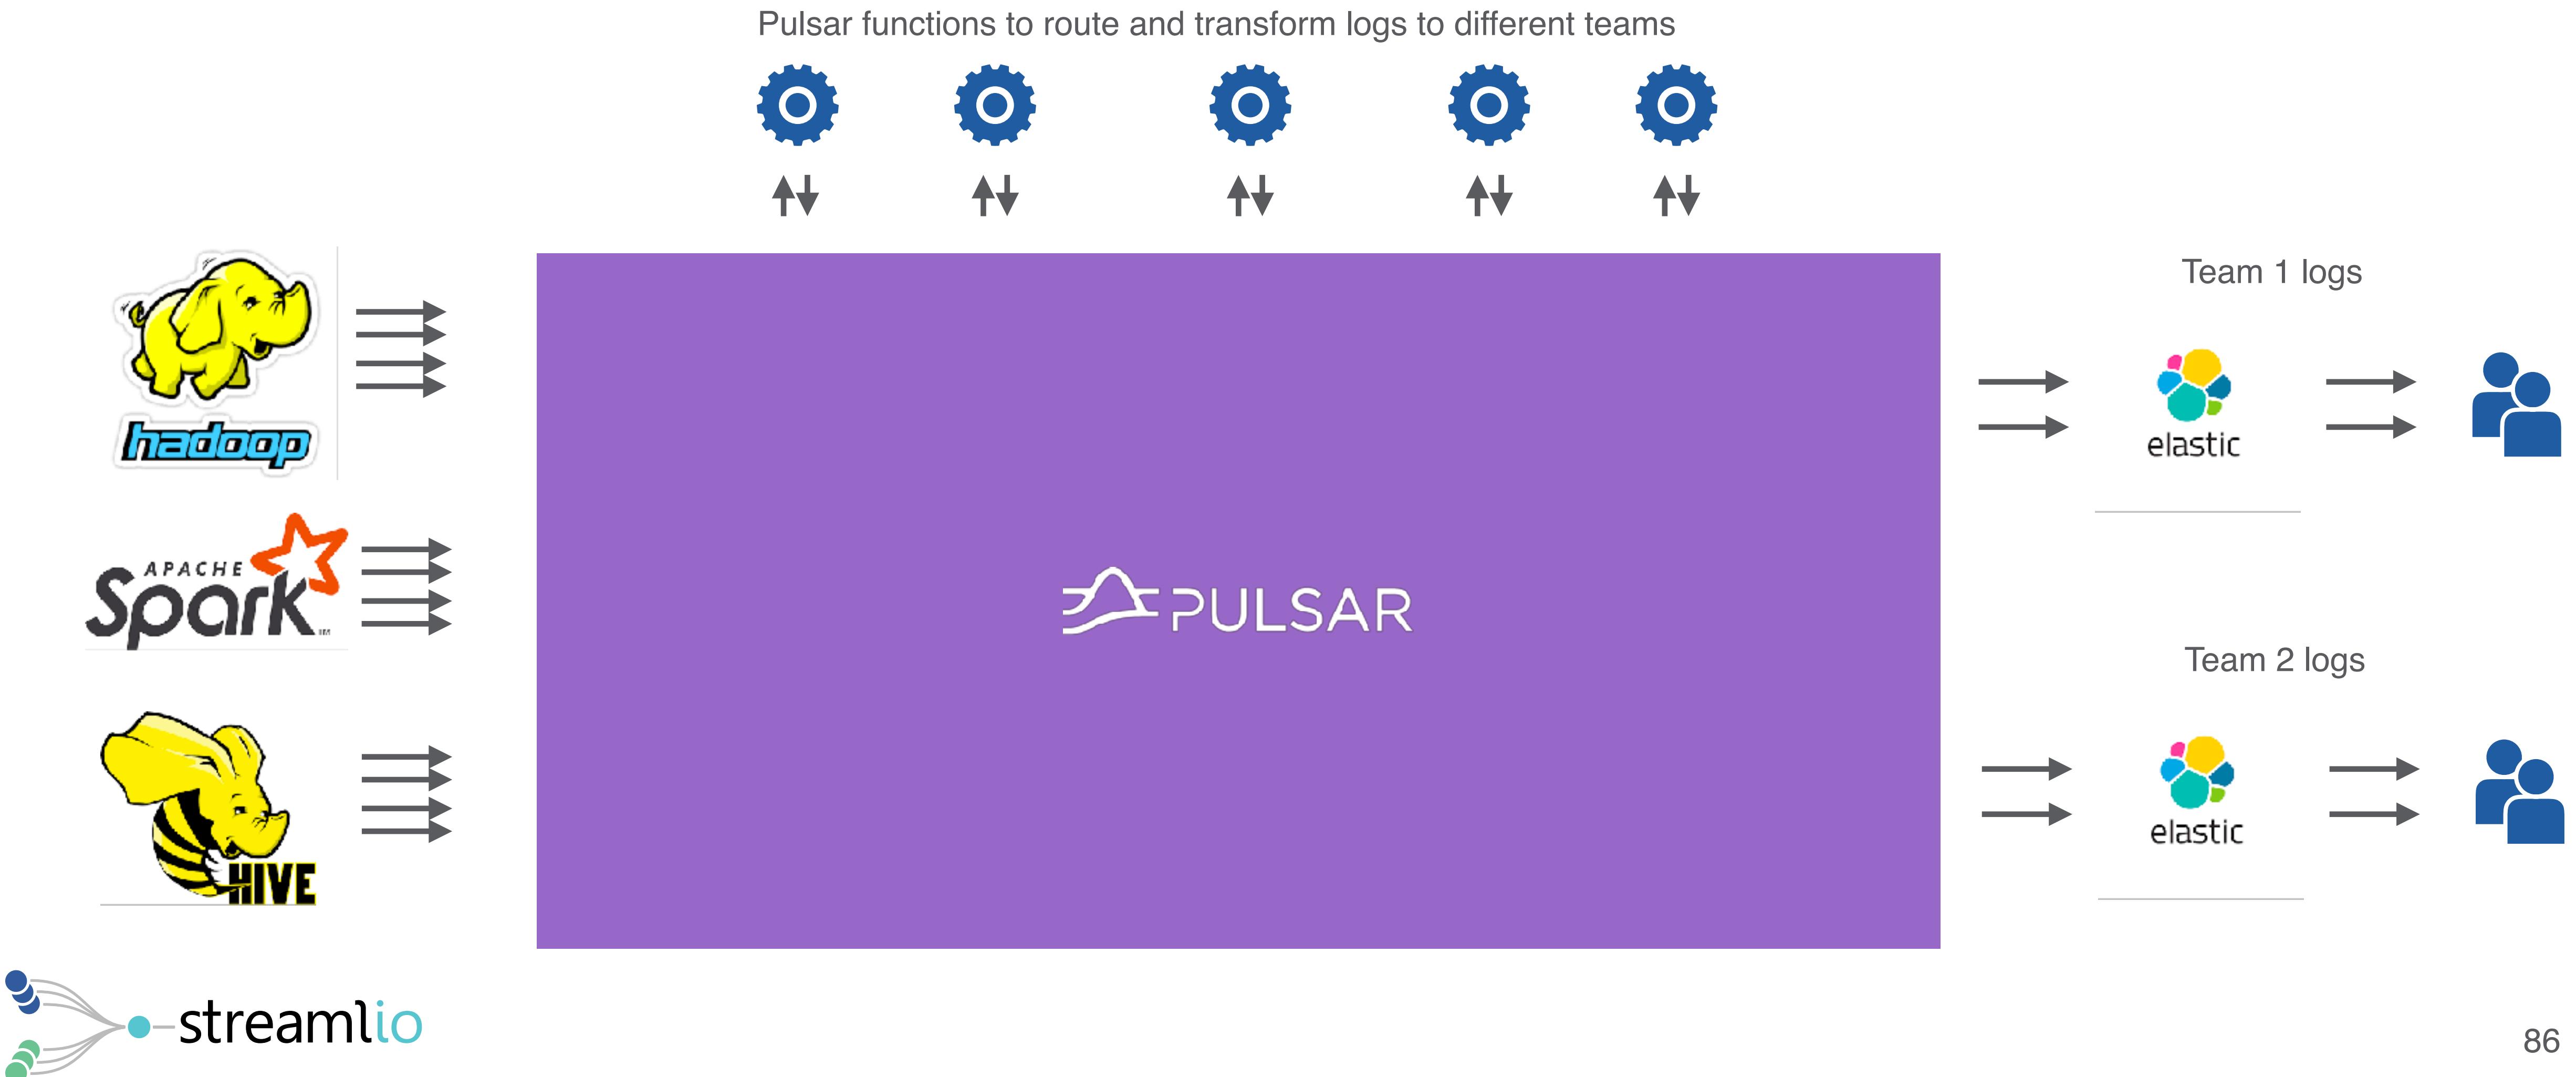
## Challenges

Require scalability to handle growing set of big data systems and larger log volumes

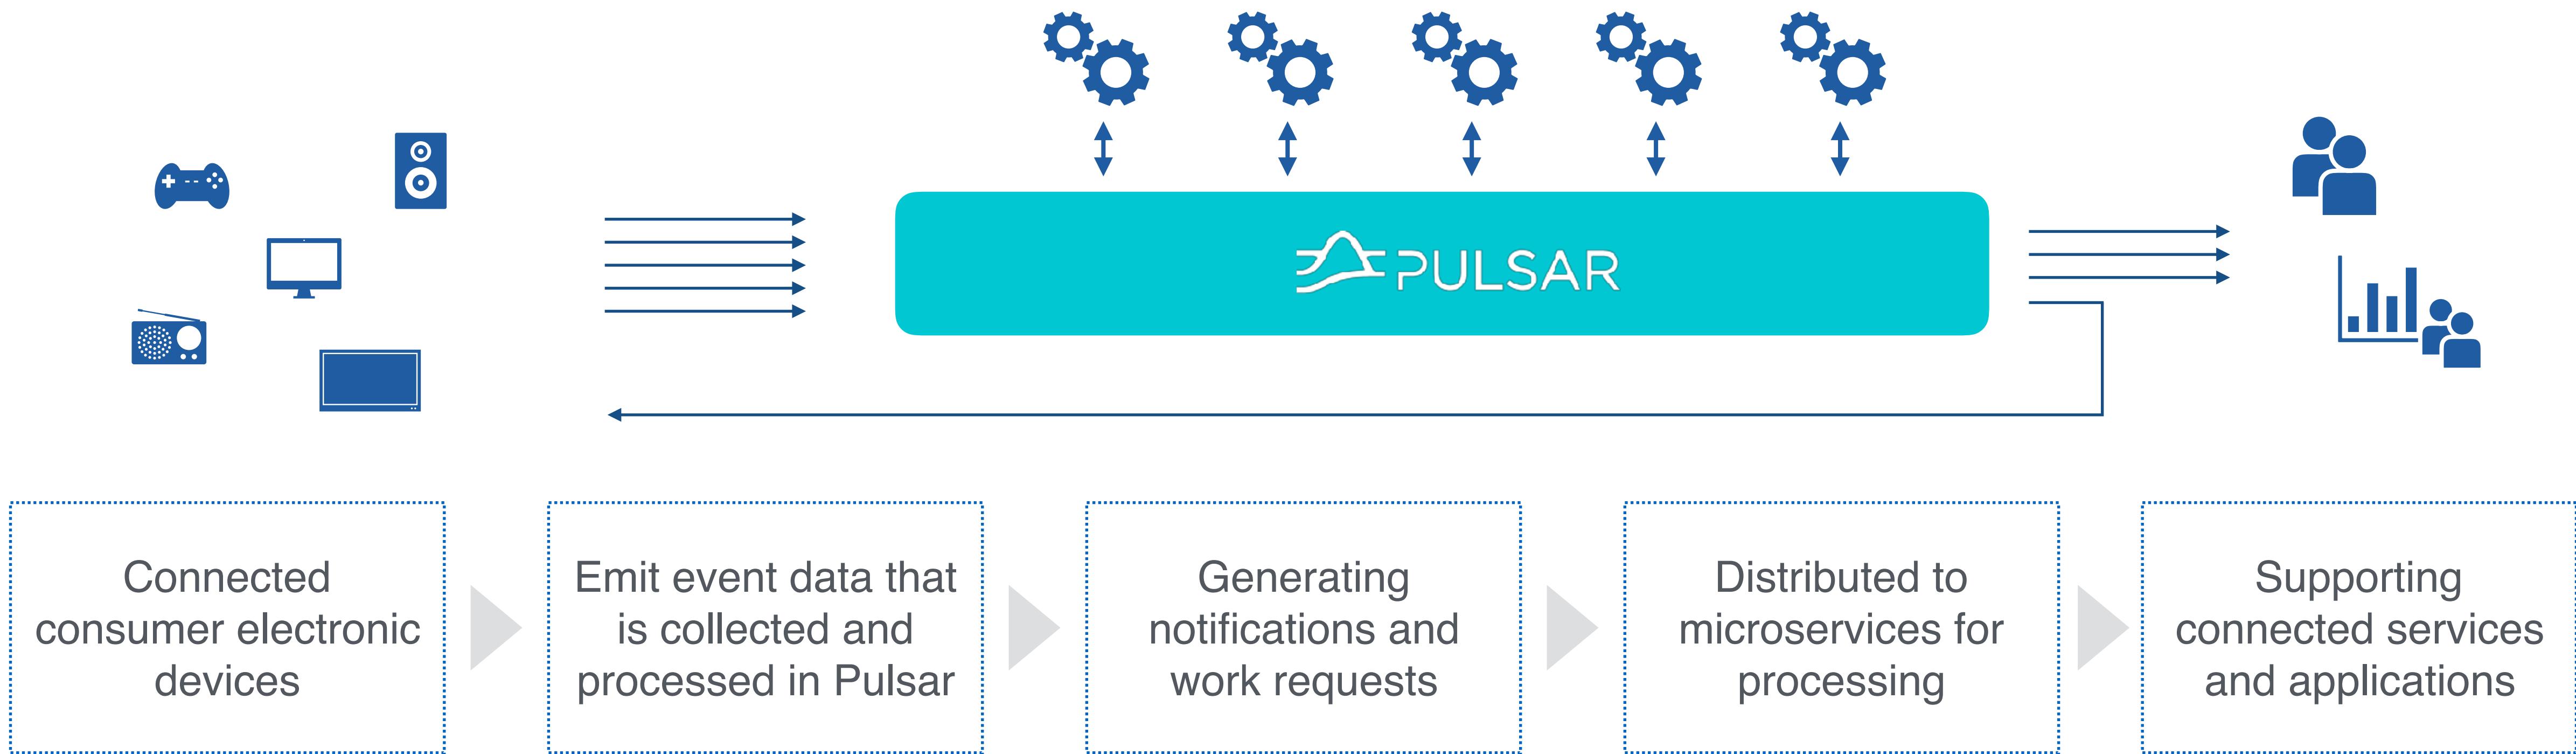
## Solution

Leverage Streamlio Pulsar solution to provide logging backbone that can ingest, transform, and distribute logs at scale

# Large Car Manufacturer: Big Data Logging System



# Connected consumer



## MORE READINGS

✓ **Understanding How Pulsar Works**

[https://jack-vanightly.com/blog/2018/10/2/understanding-how-apache-pulsar-works](https://jack-vanlightly.com/blog/2018/10/2/understanding-how-apache-pulsar-works)

✓ **How To (Not) Lose Messages on Apache Pulsar Cluster**

<https://jack-vanightly.com/blog/2018/10/21/how-to-not-lose-messages-on-an-apache-pulsar-cluster>

## MORE READINGS

- ✓ **Unified queuing and streaming**

<https://streamlio.io/blog/pulsar-streaming-queuing>

- ✓ **Segment centric storage**

<https://streamlio.io/blog/pulsar-segment-based-architecture>

- ✓ **Messaging, Storage or Both**

<https://streamlio.io/blog/messaging-storage-or-both>

- ✓ **Access patterns and tiered storage**

<https://streamlio.io/blog/access-patterns-and-tiered-storage-in-apache-pulsar>

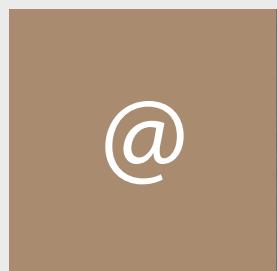
- ✓ **Tiered Storage in Apache Pulsar**

<https://streamlio.io/blog/tiered-storage-in-apache-pulsar>

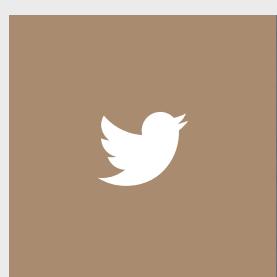
# QUESTIONS



# STAY IN TOUCH



EMAIL  
[karthik@streamli.io](mailto:karthik@streamli.io)



TWITTER  
[@karthikz](https://twitter.com/karthikz)



" My son is a corporate communications director.  
He never calls and he never writes."



 @karthikz